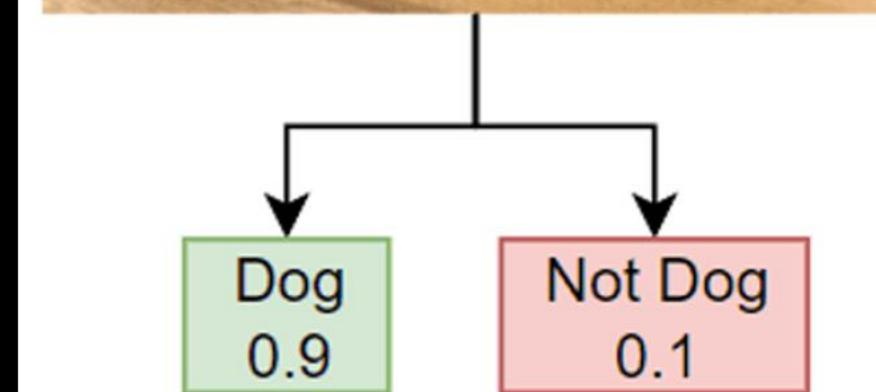
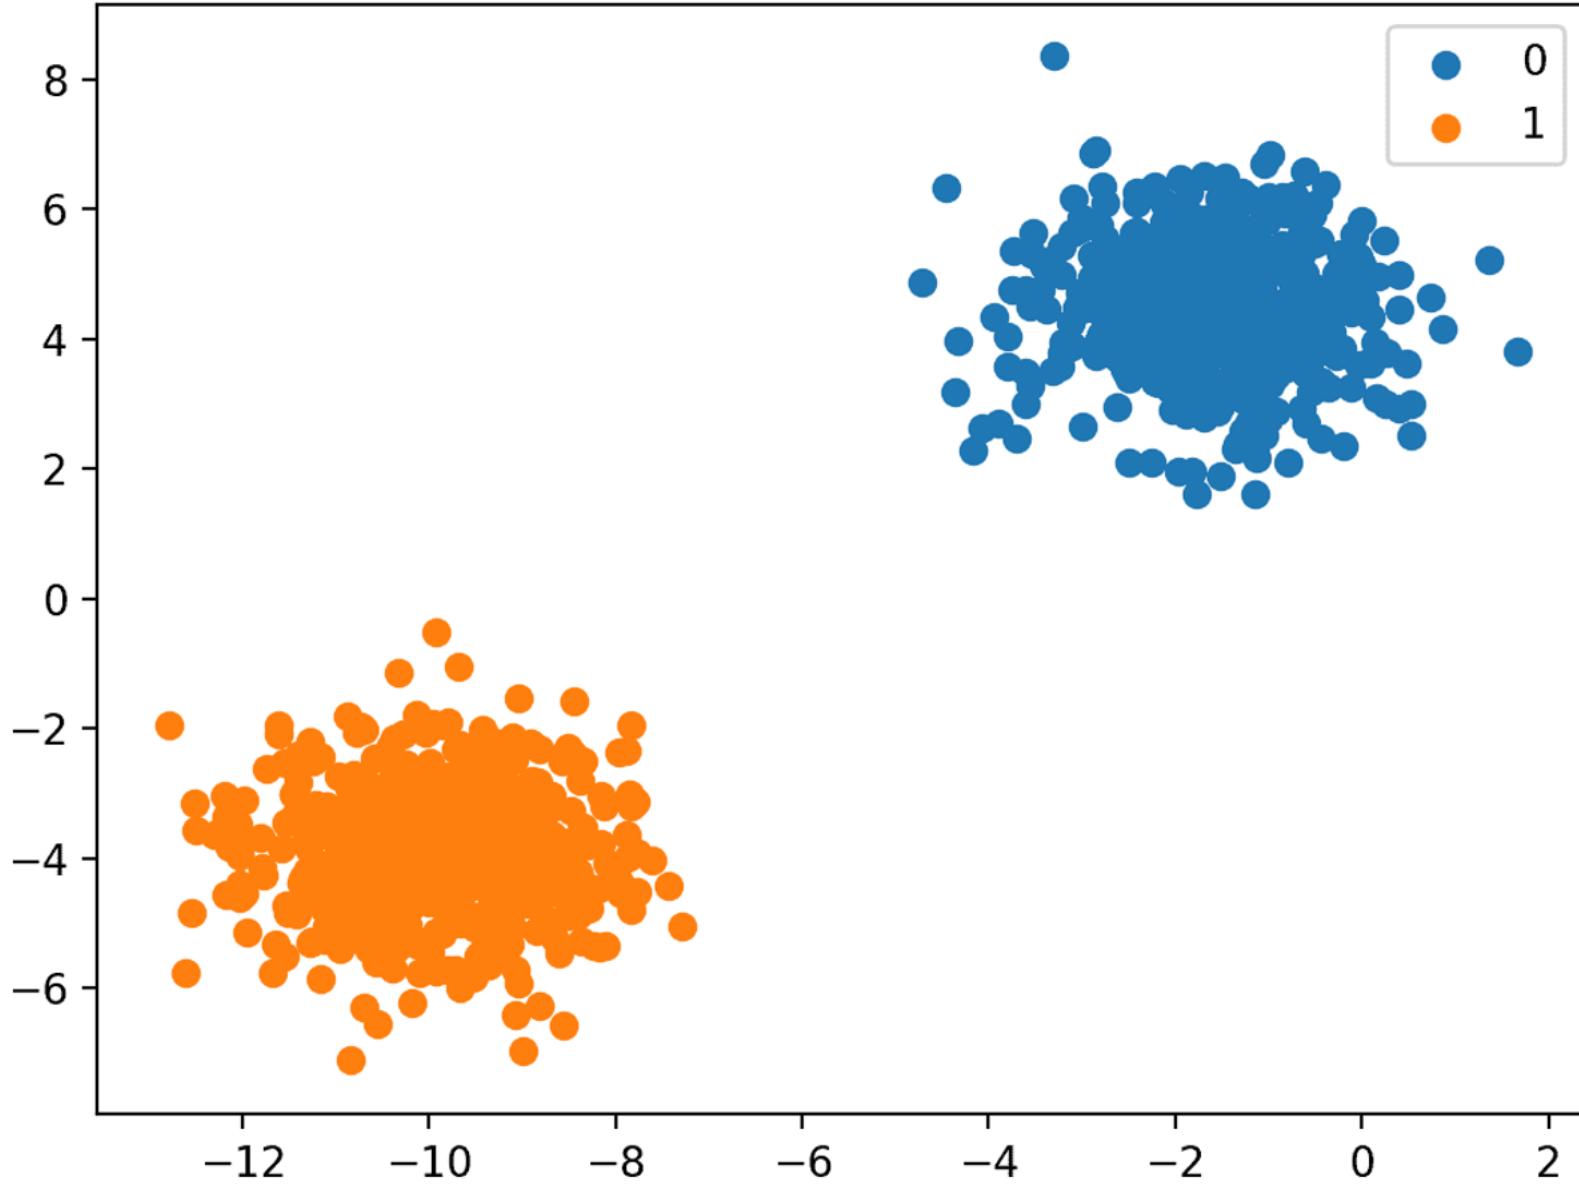


# classification

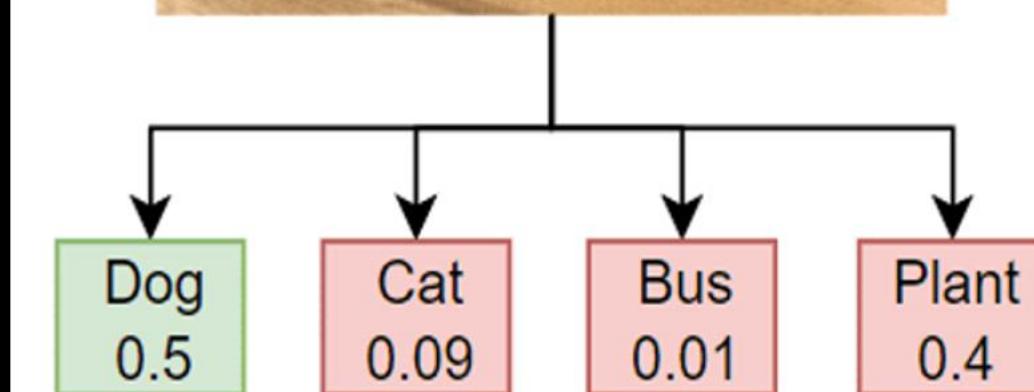
supervised learning task of learning a function  
mapping an input point to a discrete category

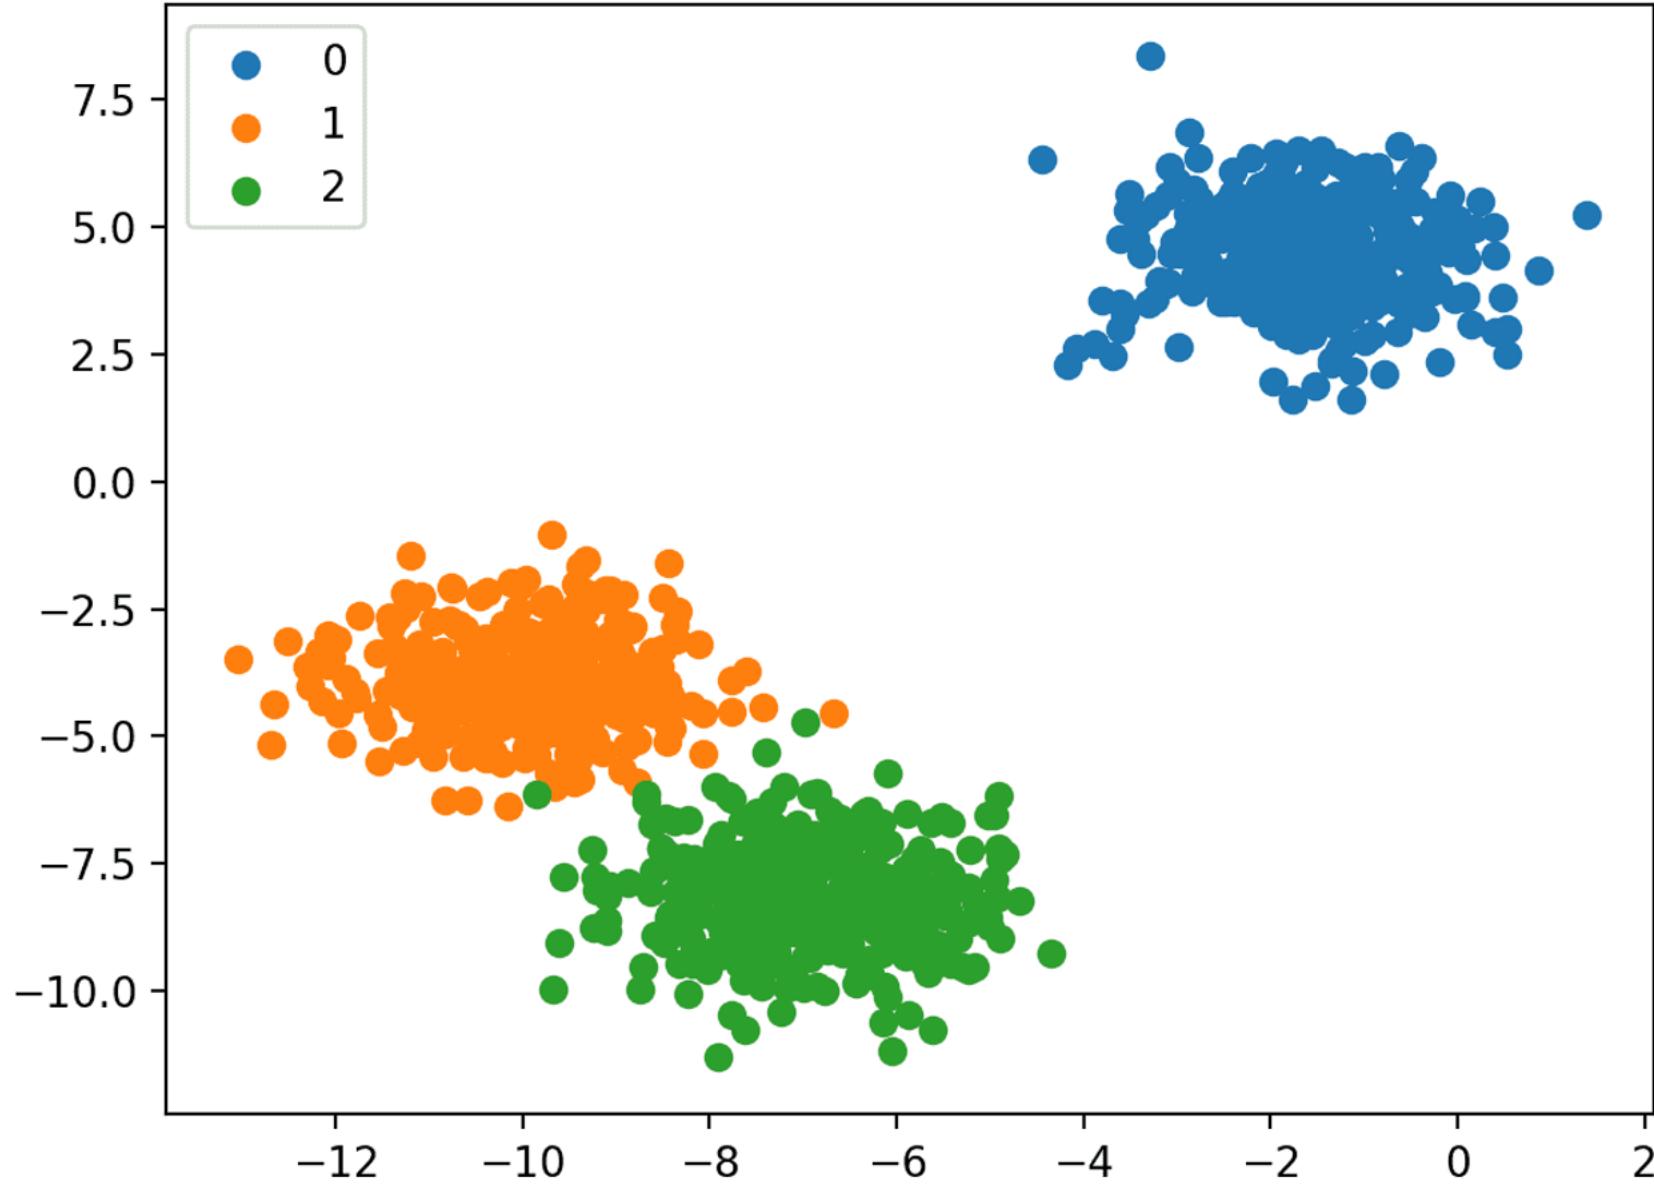
## Binary Classification



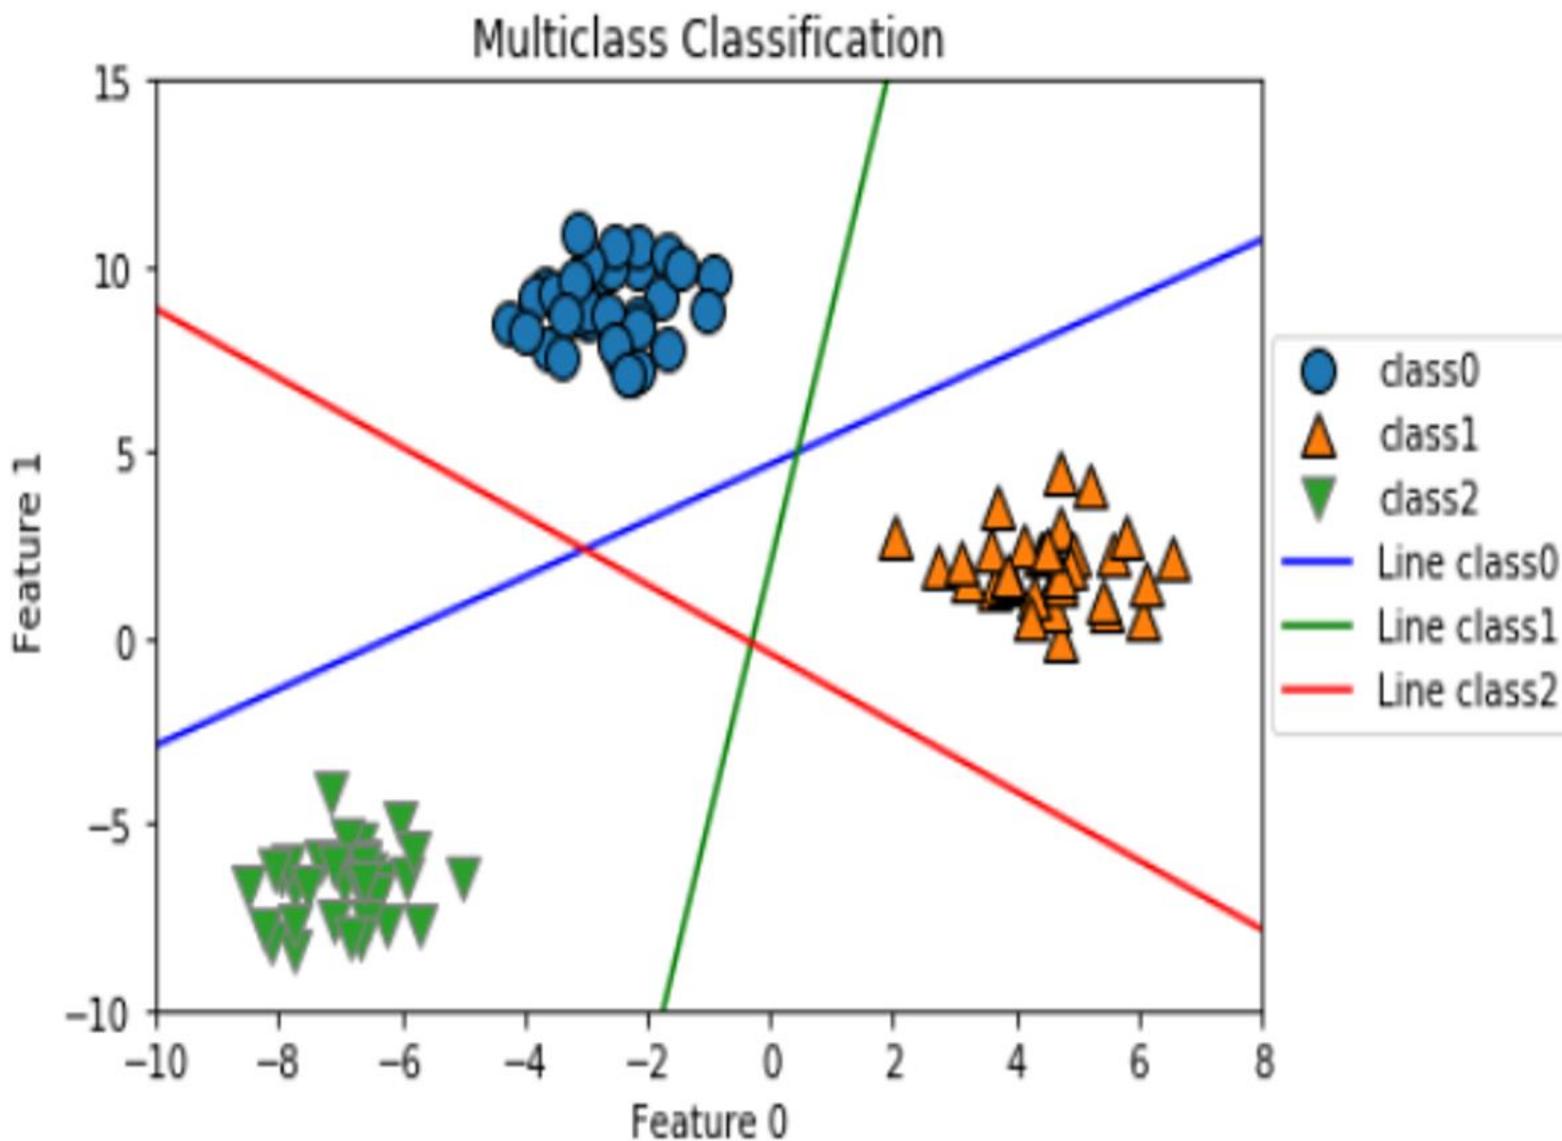


## Multiclass Classification

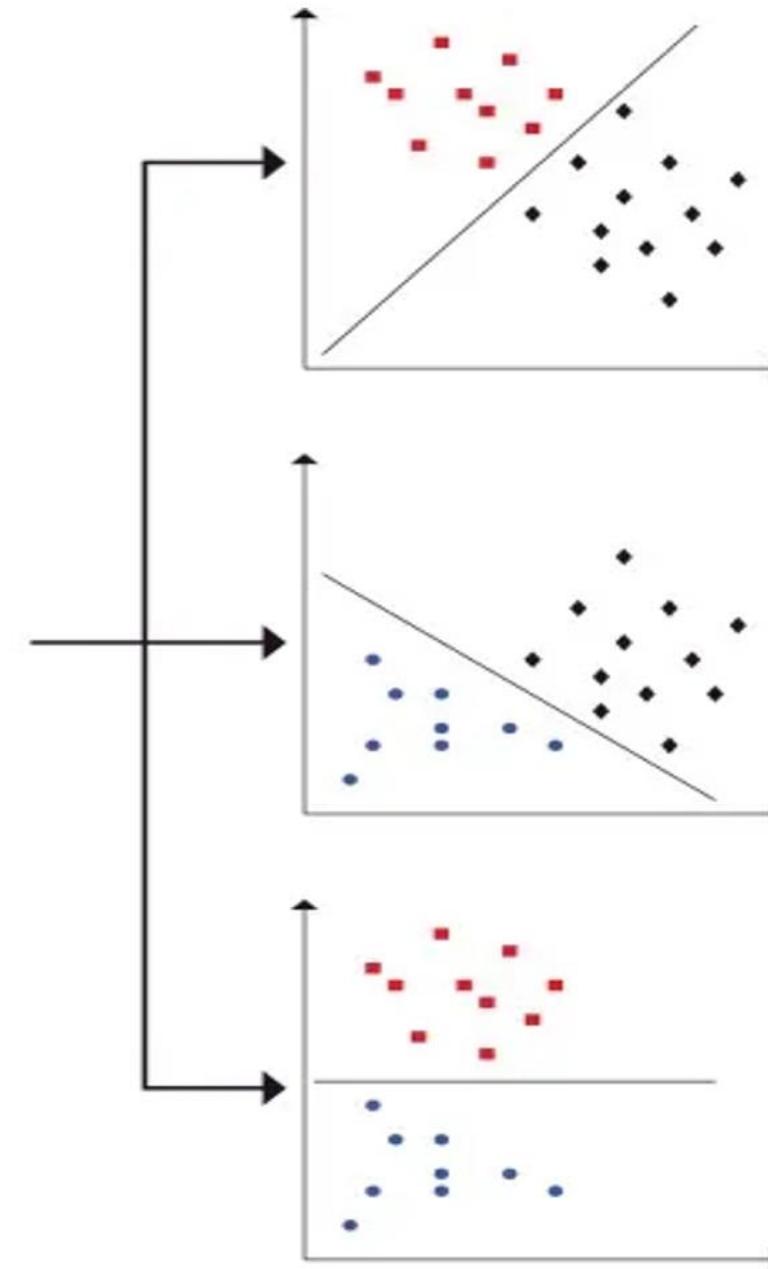
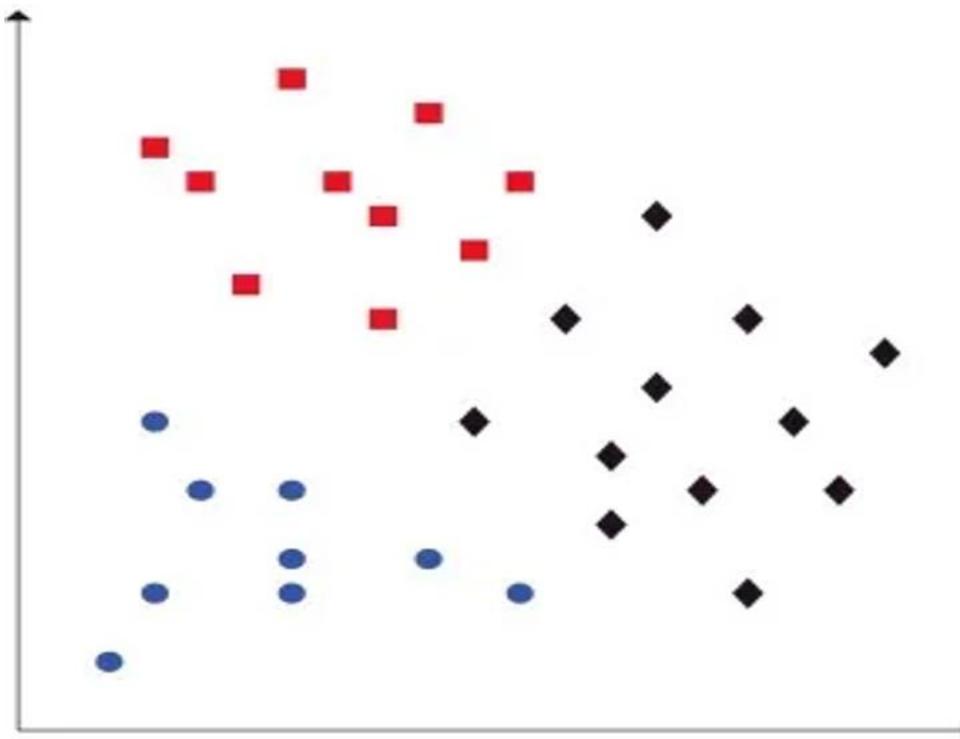




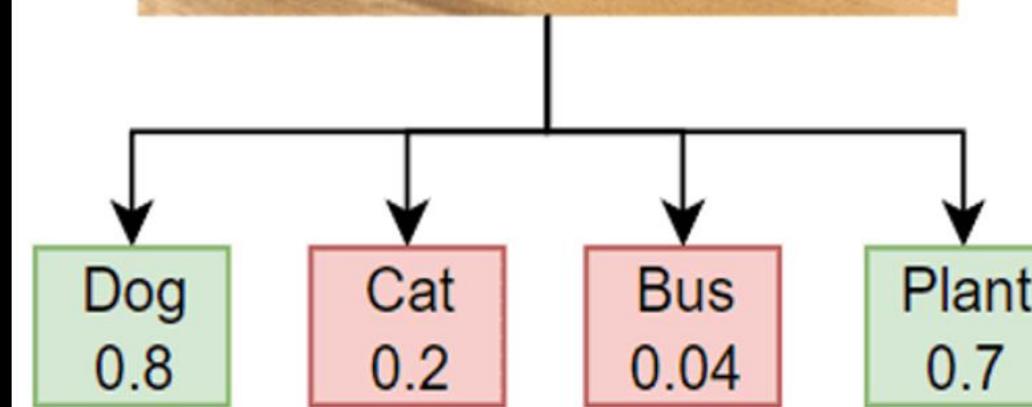
One-vs-All(Rest): Fit one binary classification model for each class vs. all other classes.



One-vs-One: Fit one binary classification model for each pair of classes.



## Multilabel Classification







Date	Humidity (relative humidity)	Pressure (sea level, mb)	Rain
January 1	93%	999.7	Rain
January 2	49%	1015.5	No Rain
January 3	79%	1031.1	No Rain
January 4	65%	984.9	Rain
January 5	90%	975.2	Rain

$f(\text{humidity}, \text{pressure})$

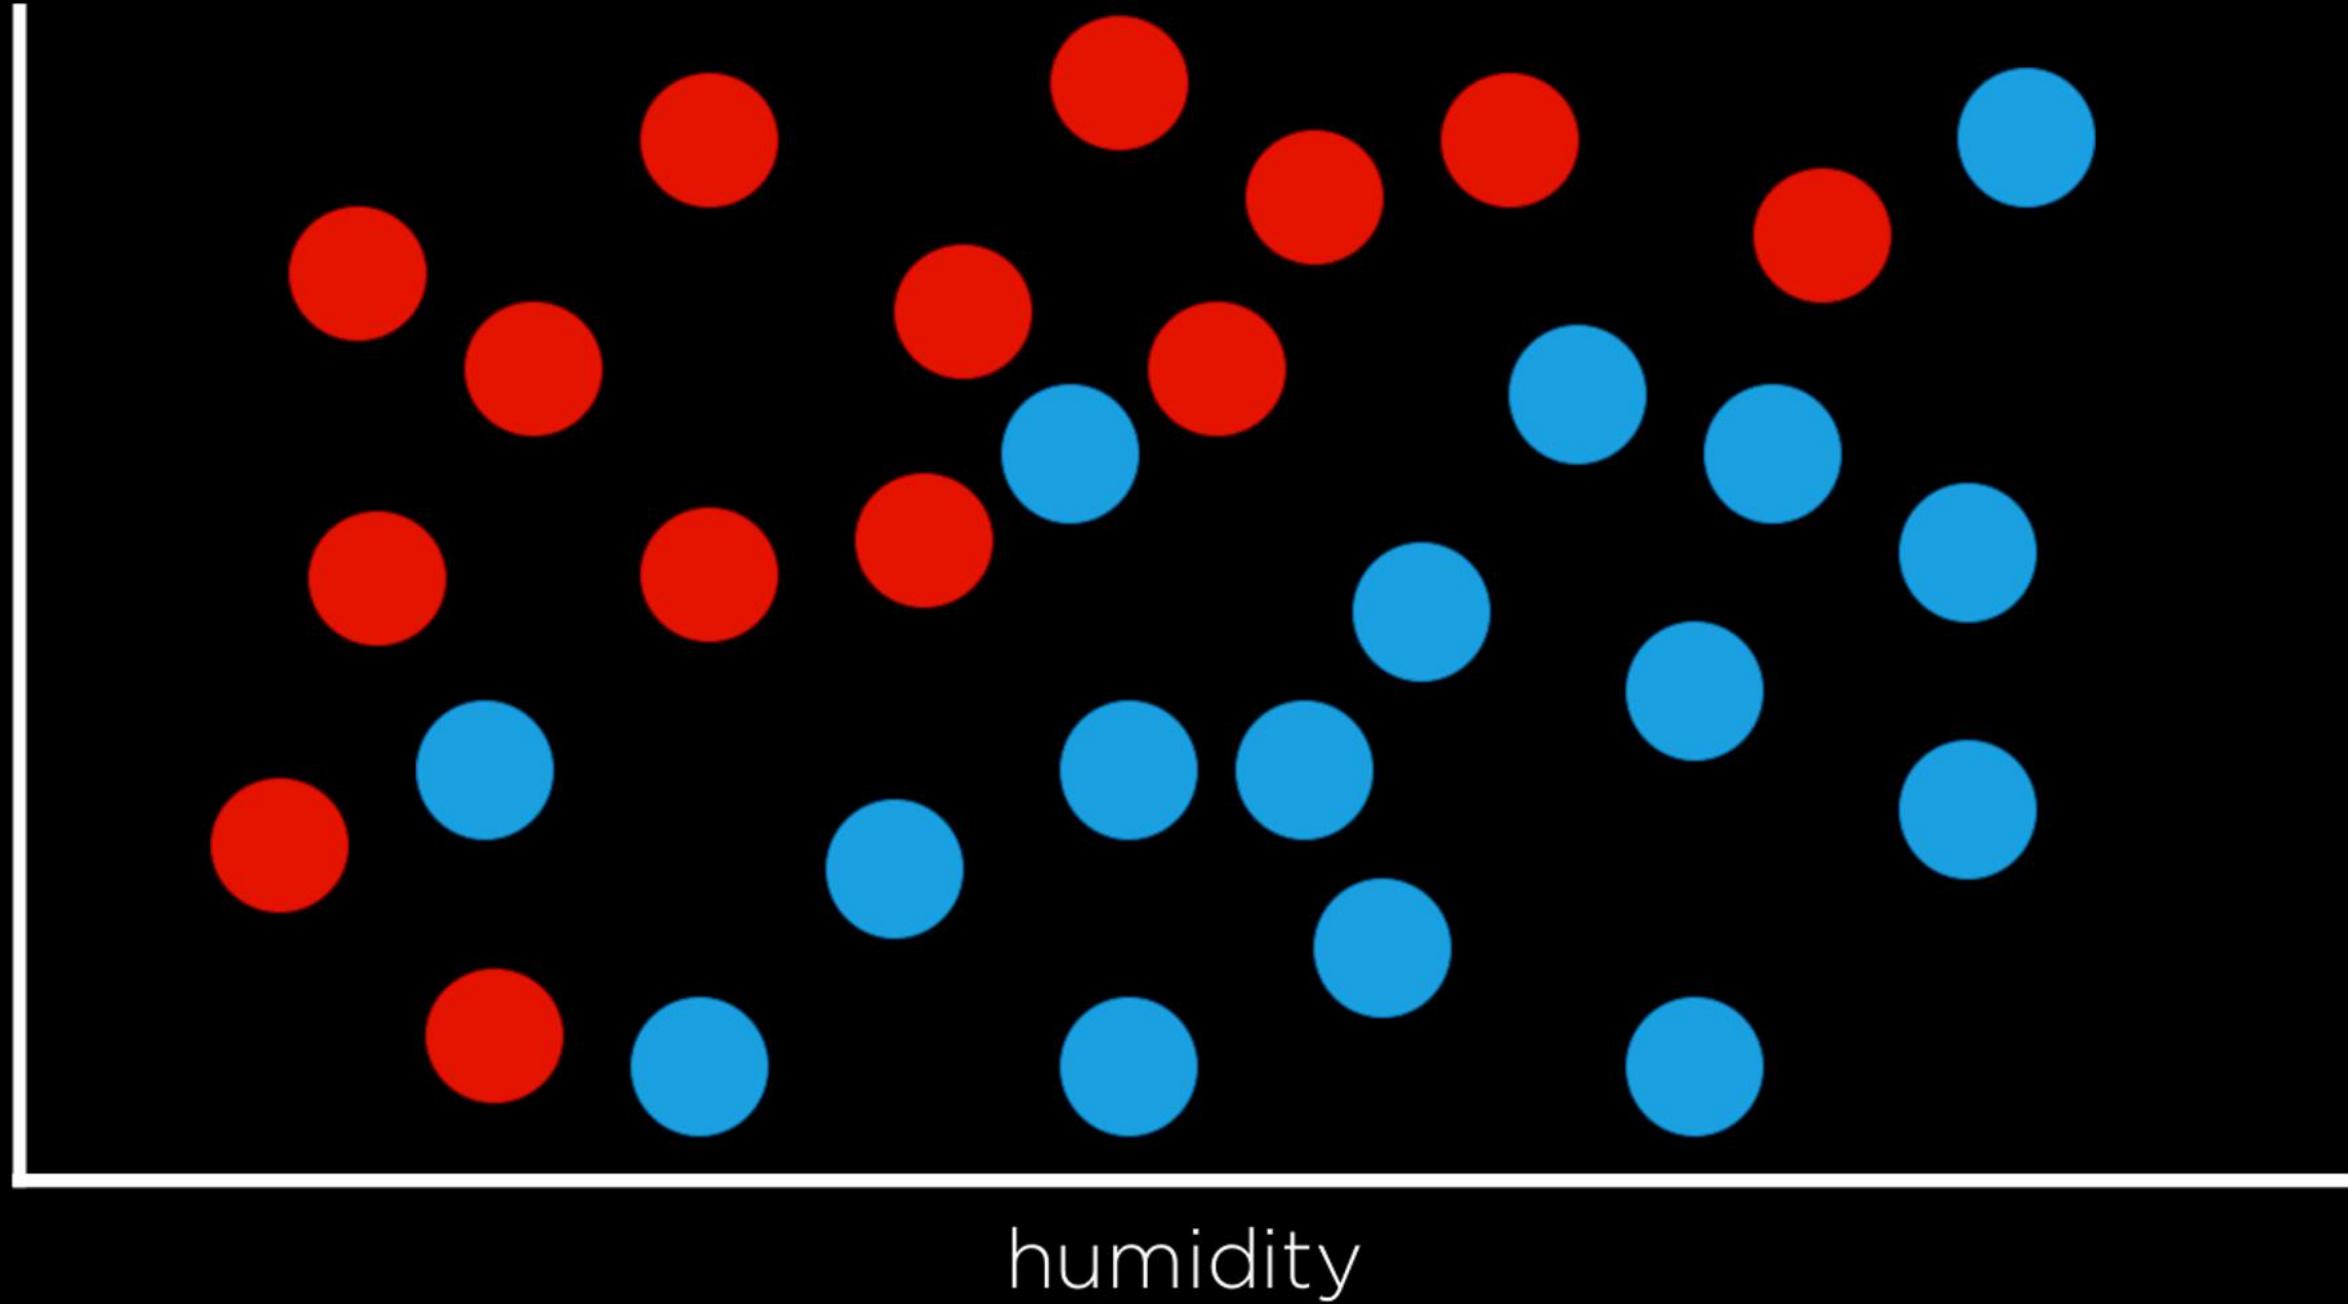
$f(93, 999.7) = \text{Rain}$

$f(49, 1015.5) = \text{No Rain}$

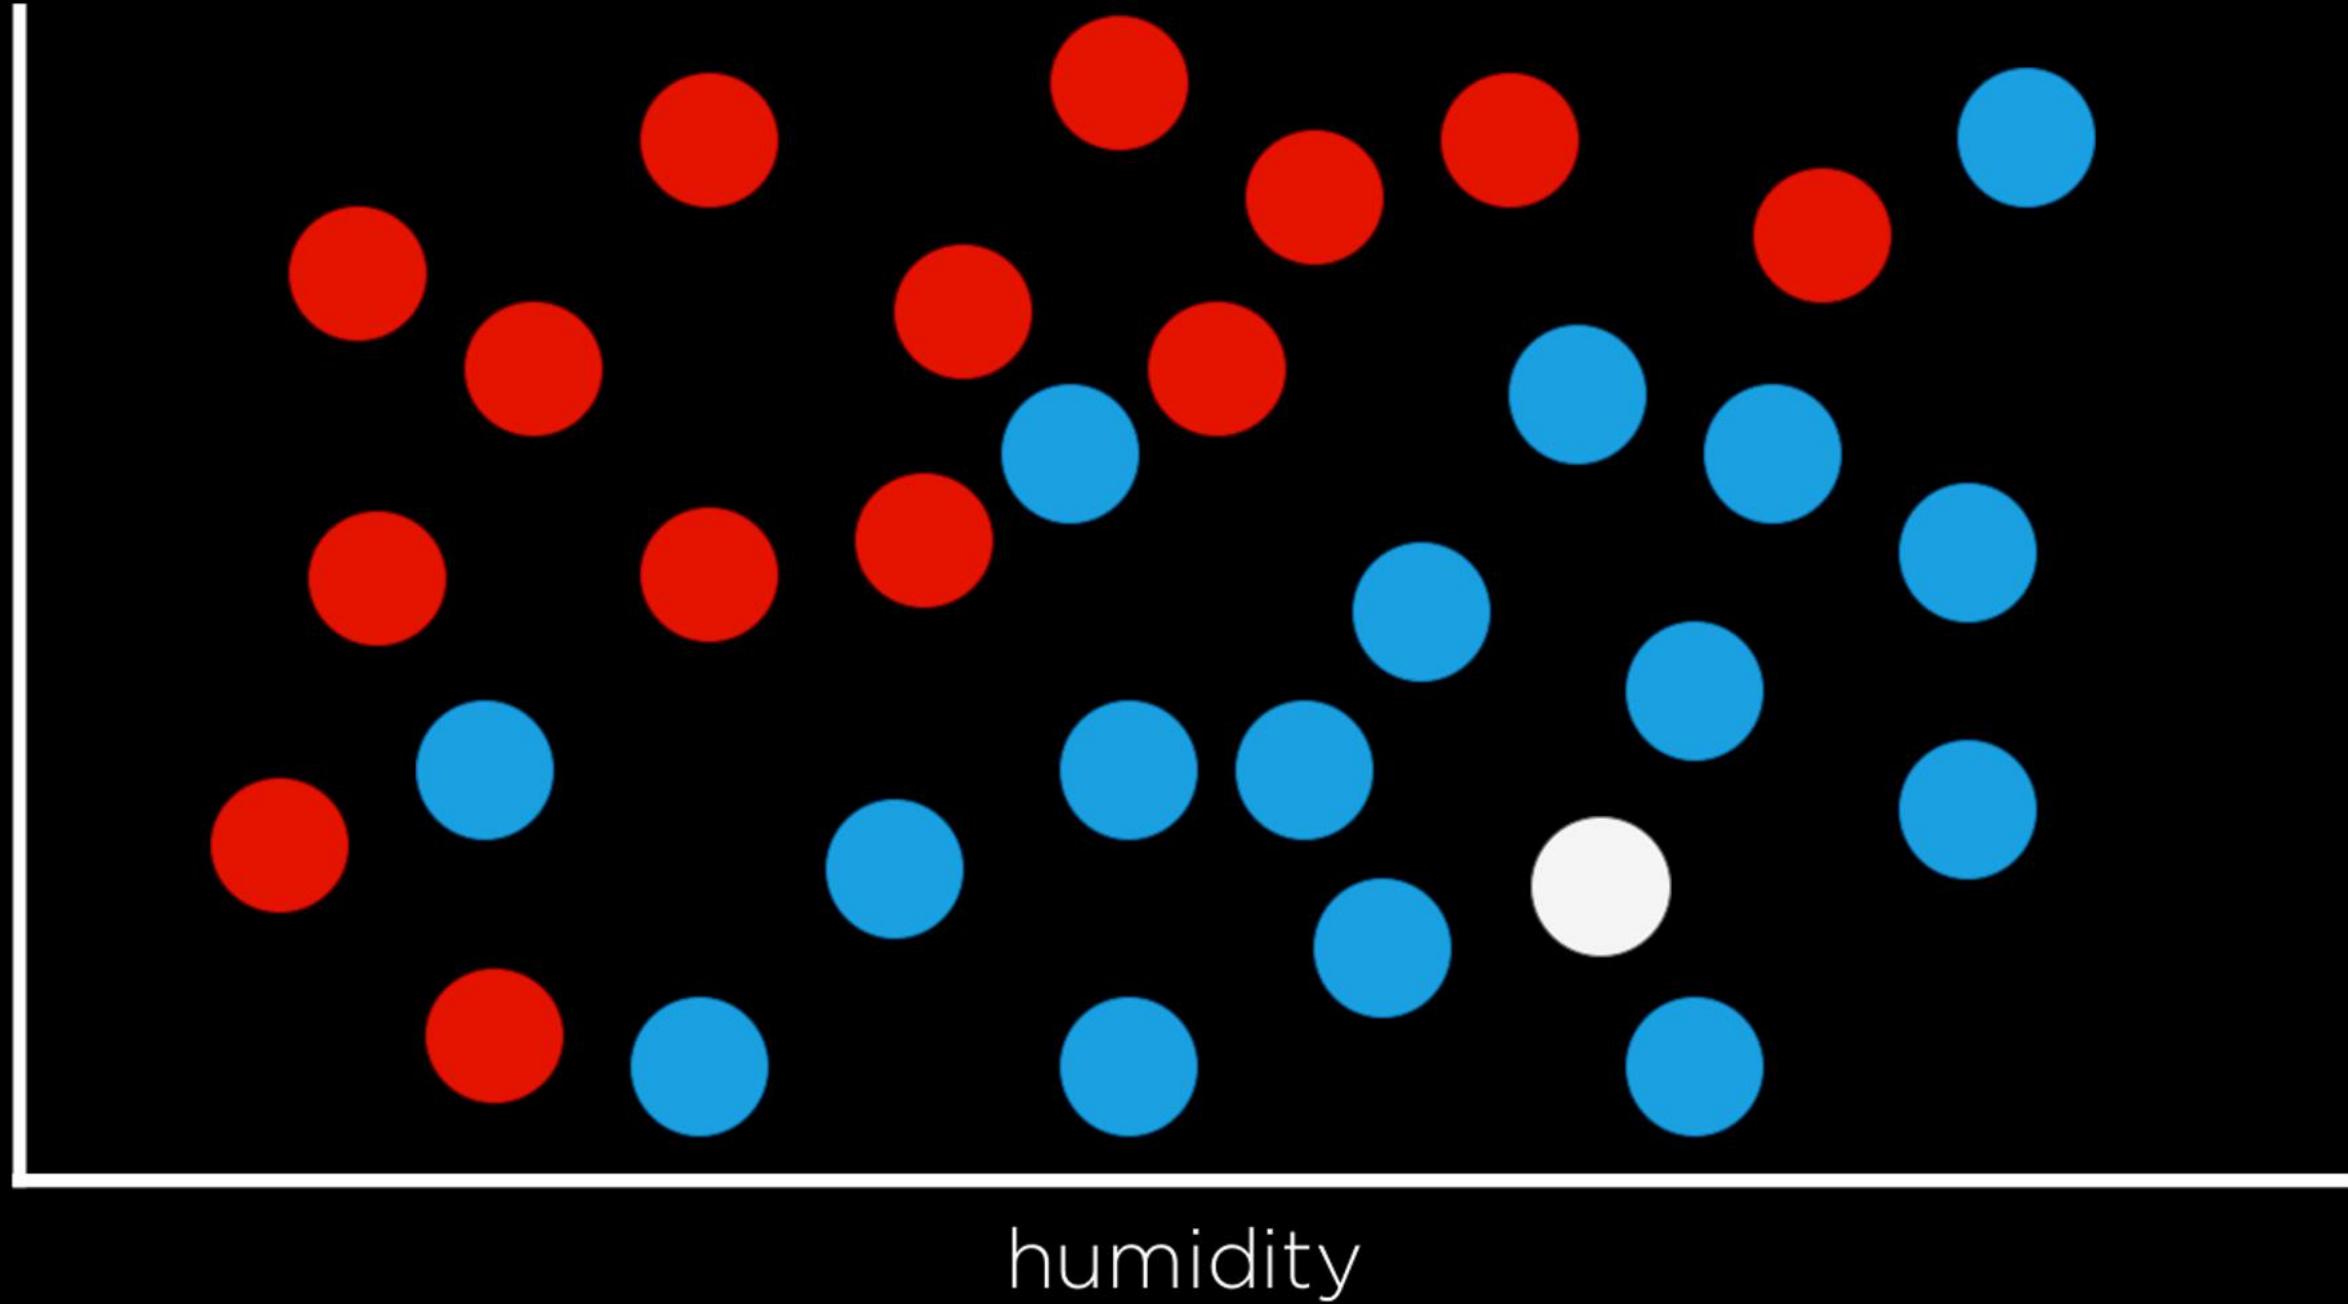
$f(79, 1031.1) = \text{No Rain}$

$h(\text{humidity}, \text{pressure})$

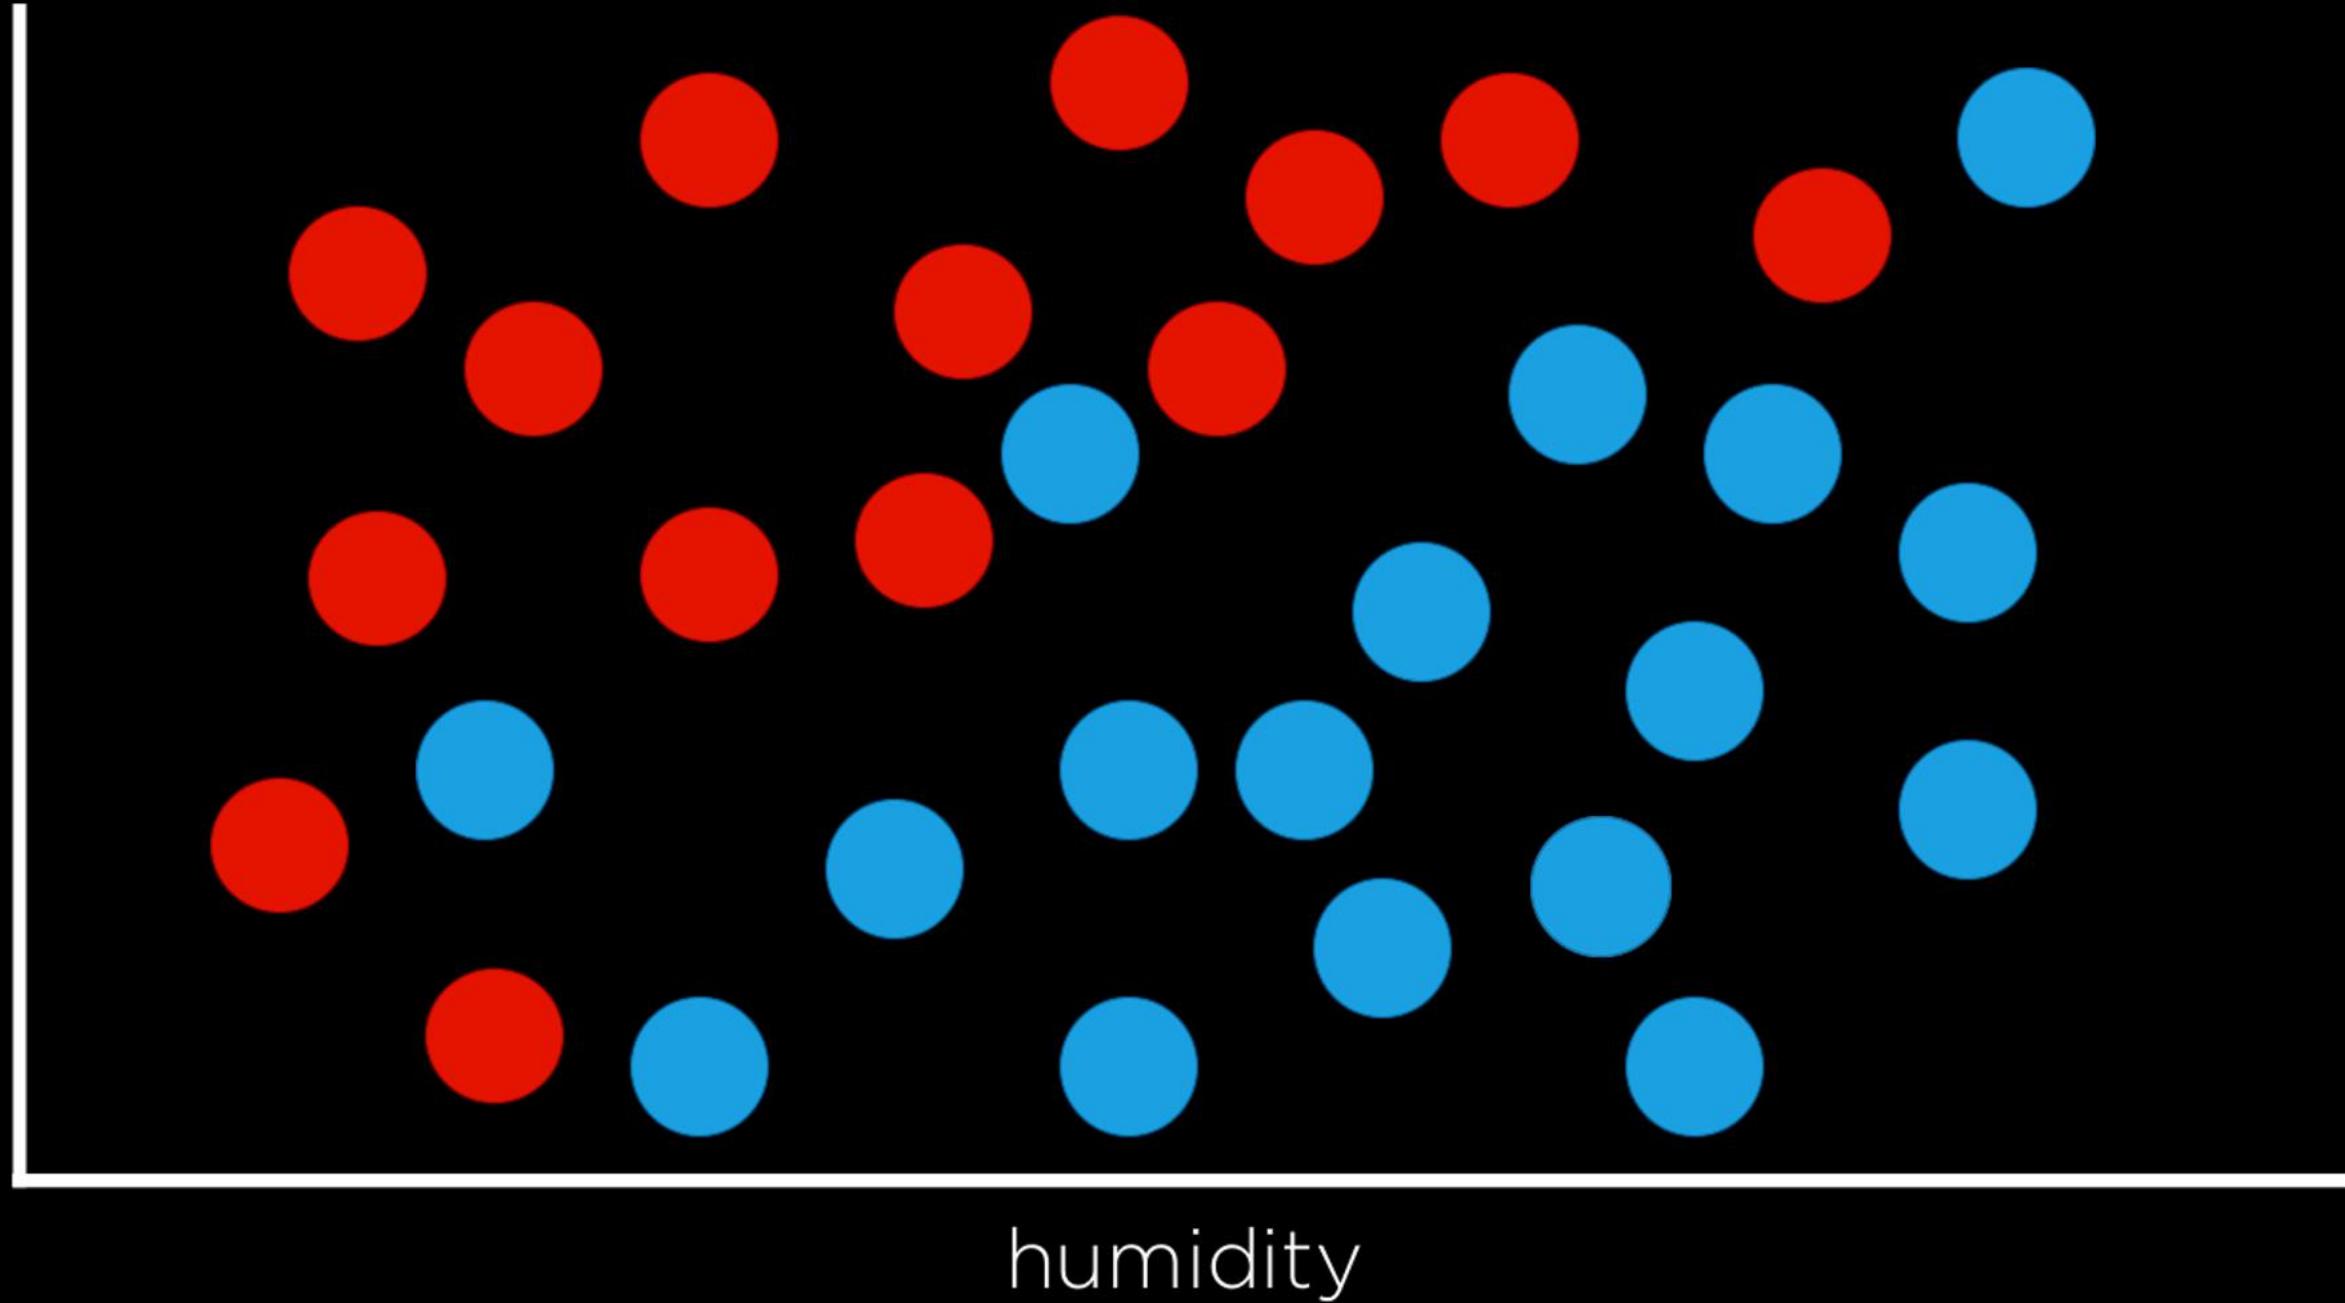
pressure



pressure



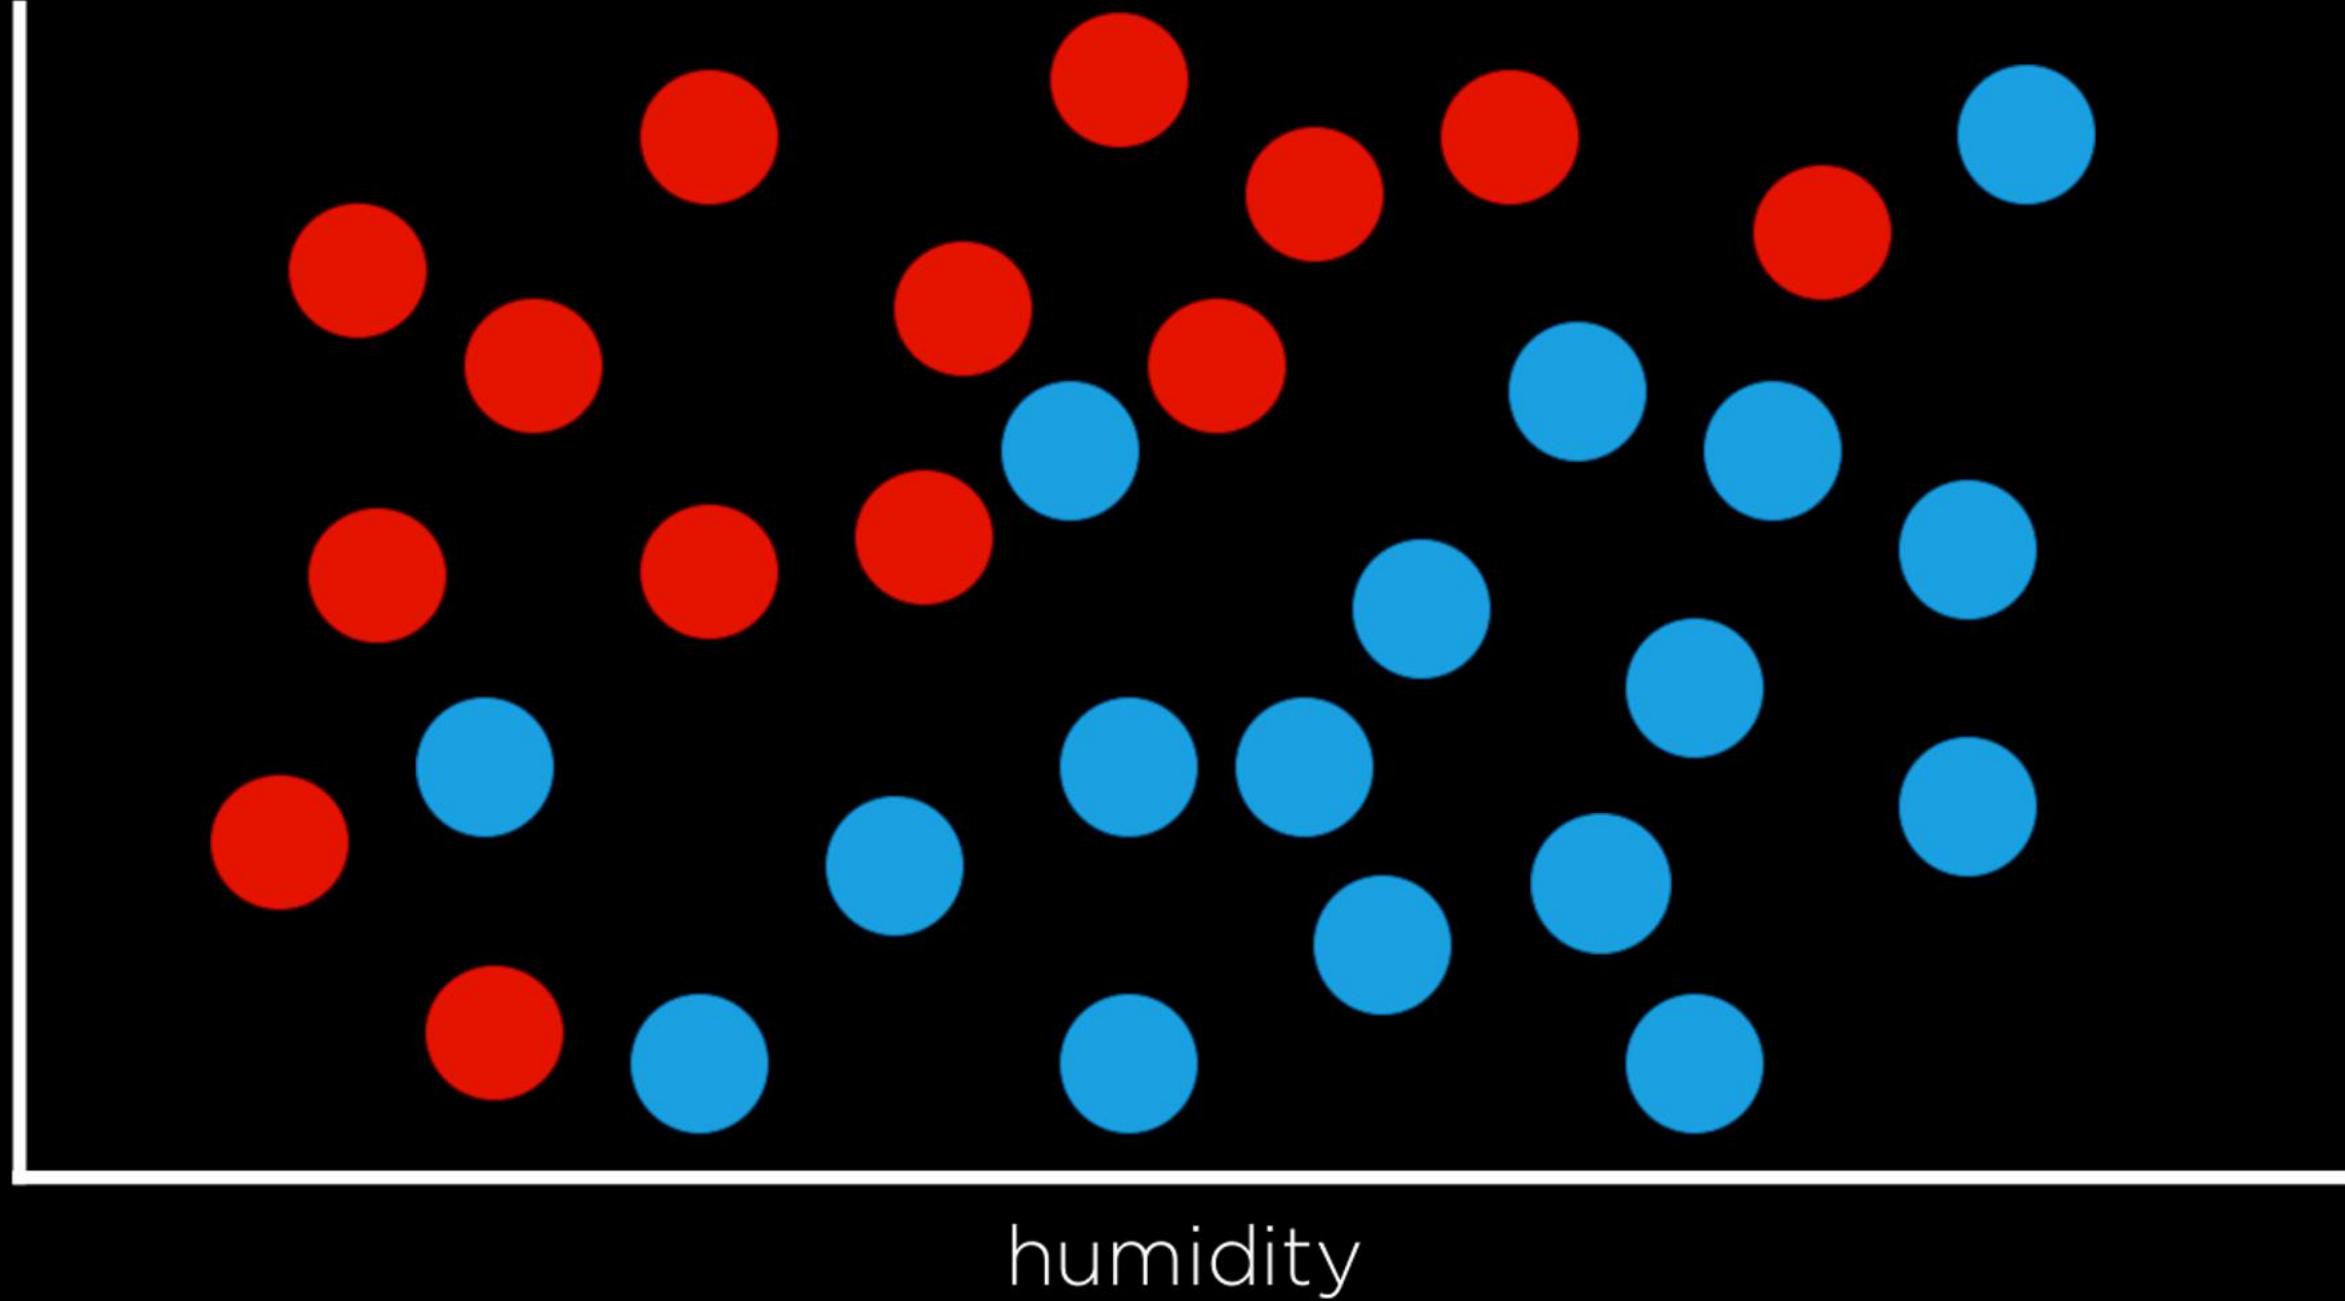
pressure



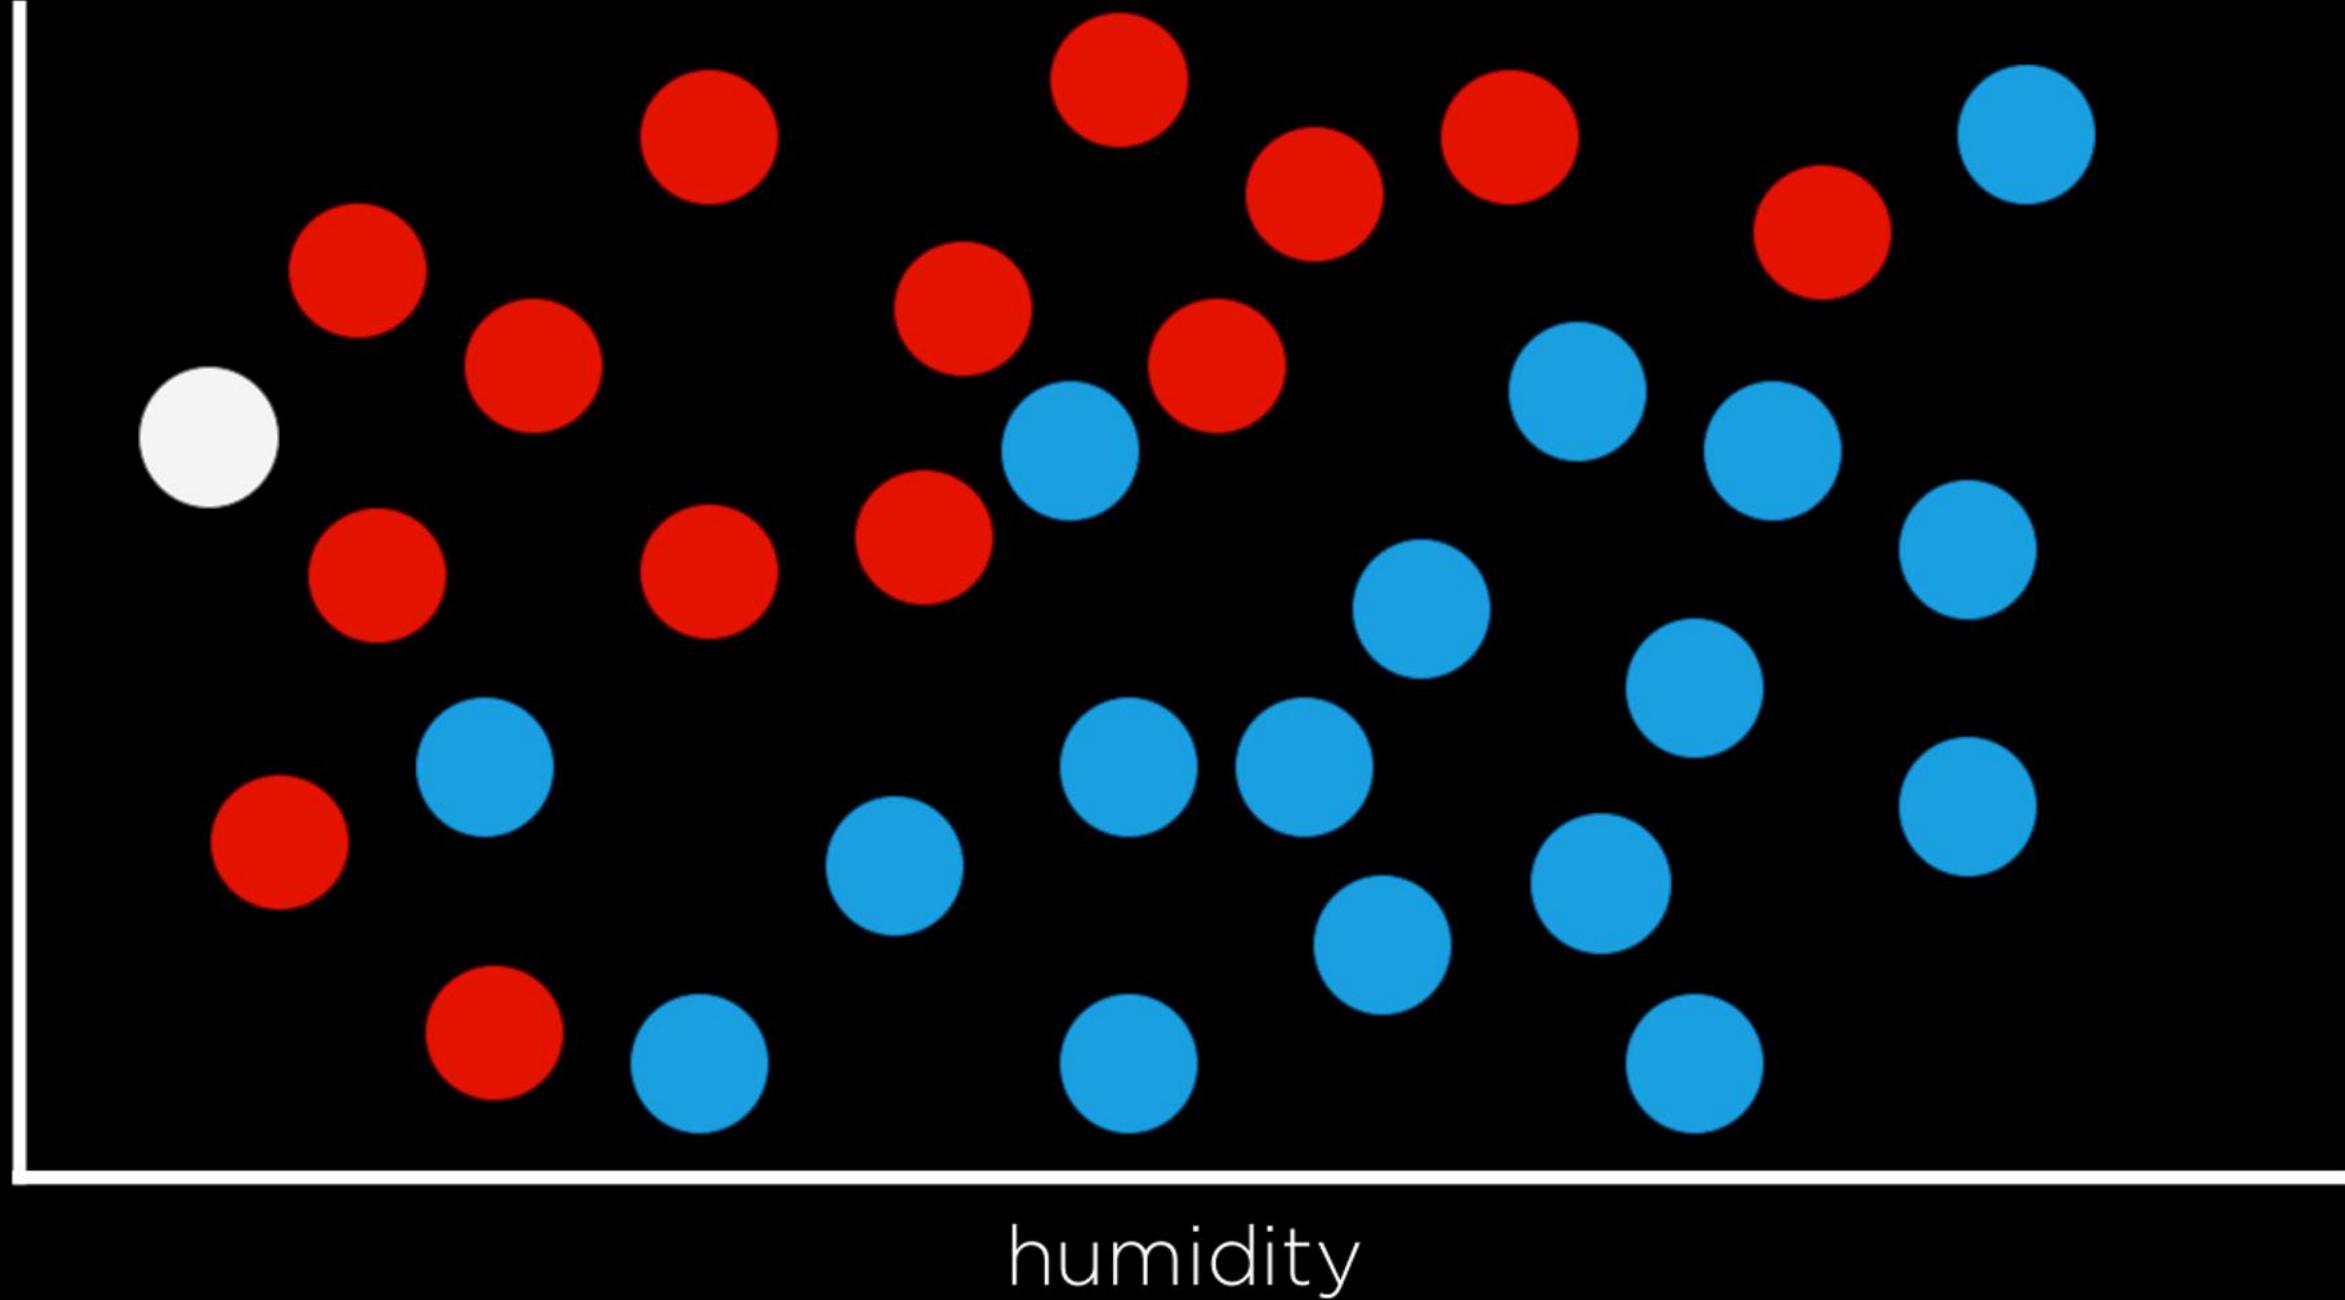
# **nearest-neighbor classification**

algorithm that, given an input, chooses the class of the nearest data point to that input

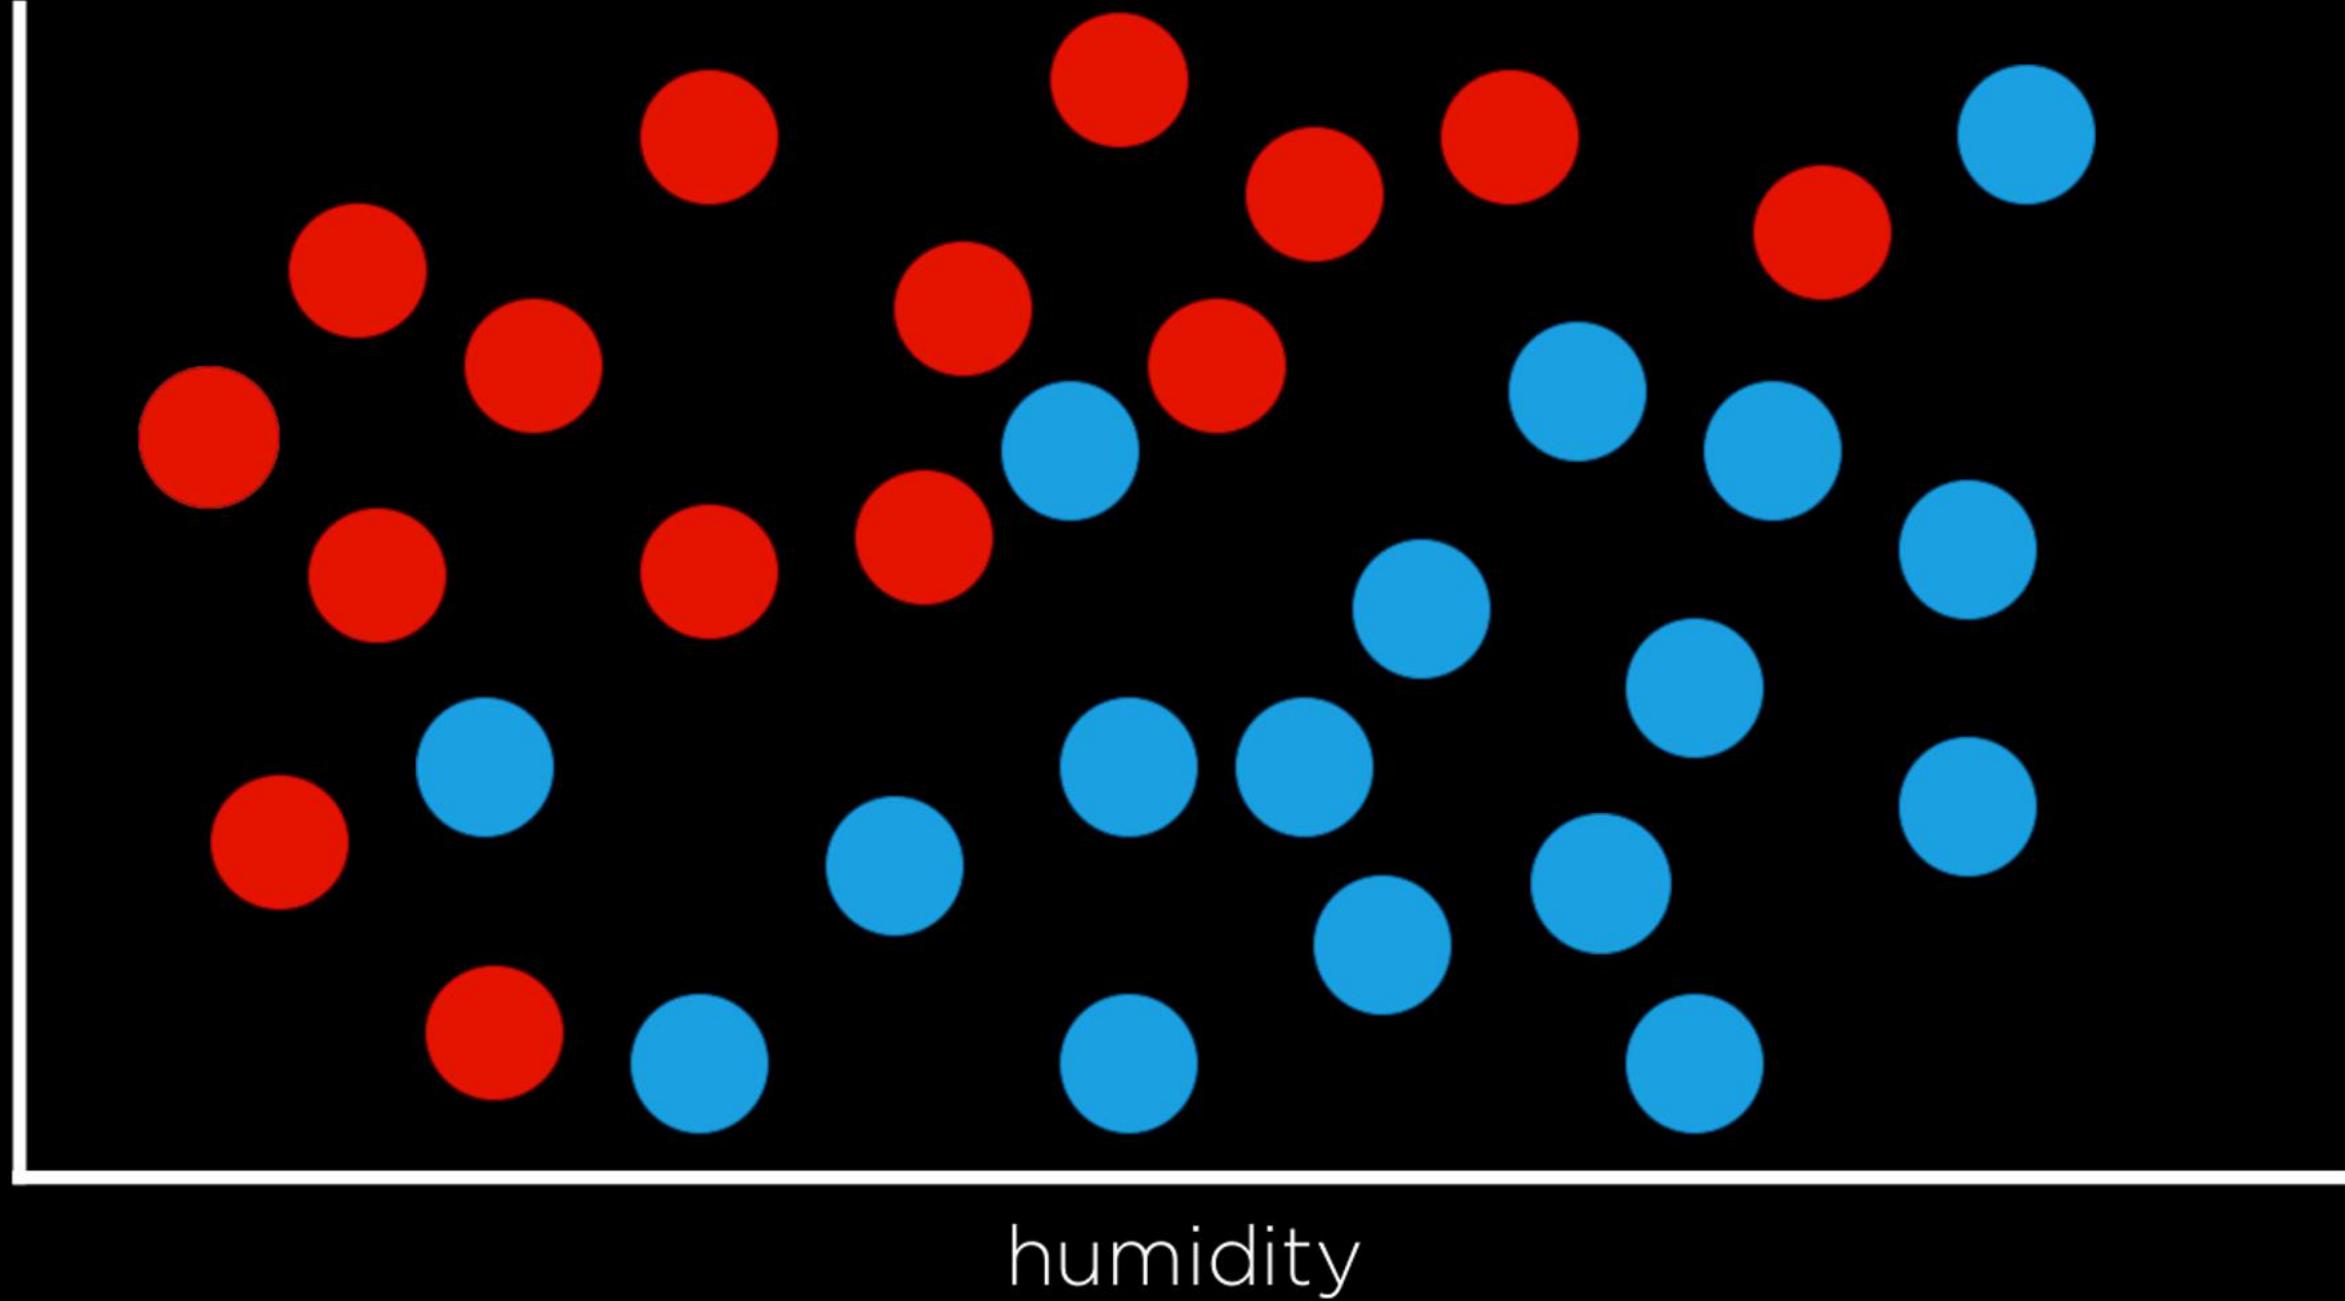
pressure

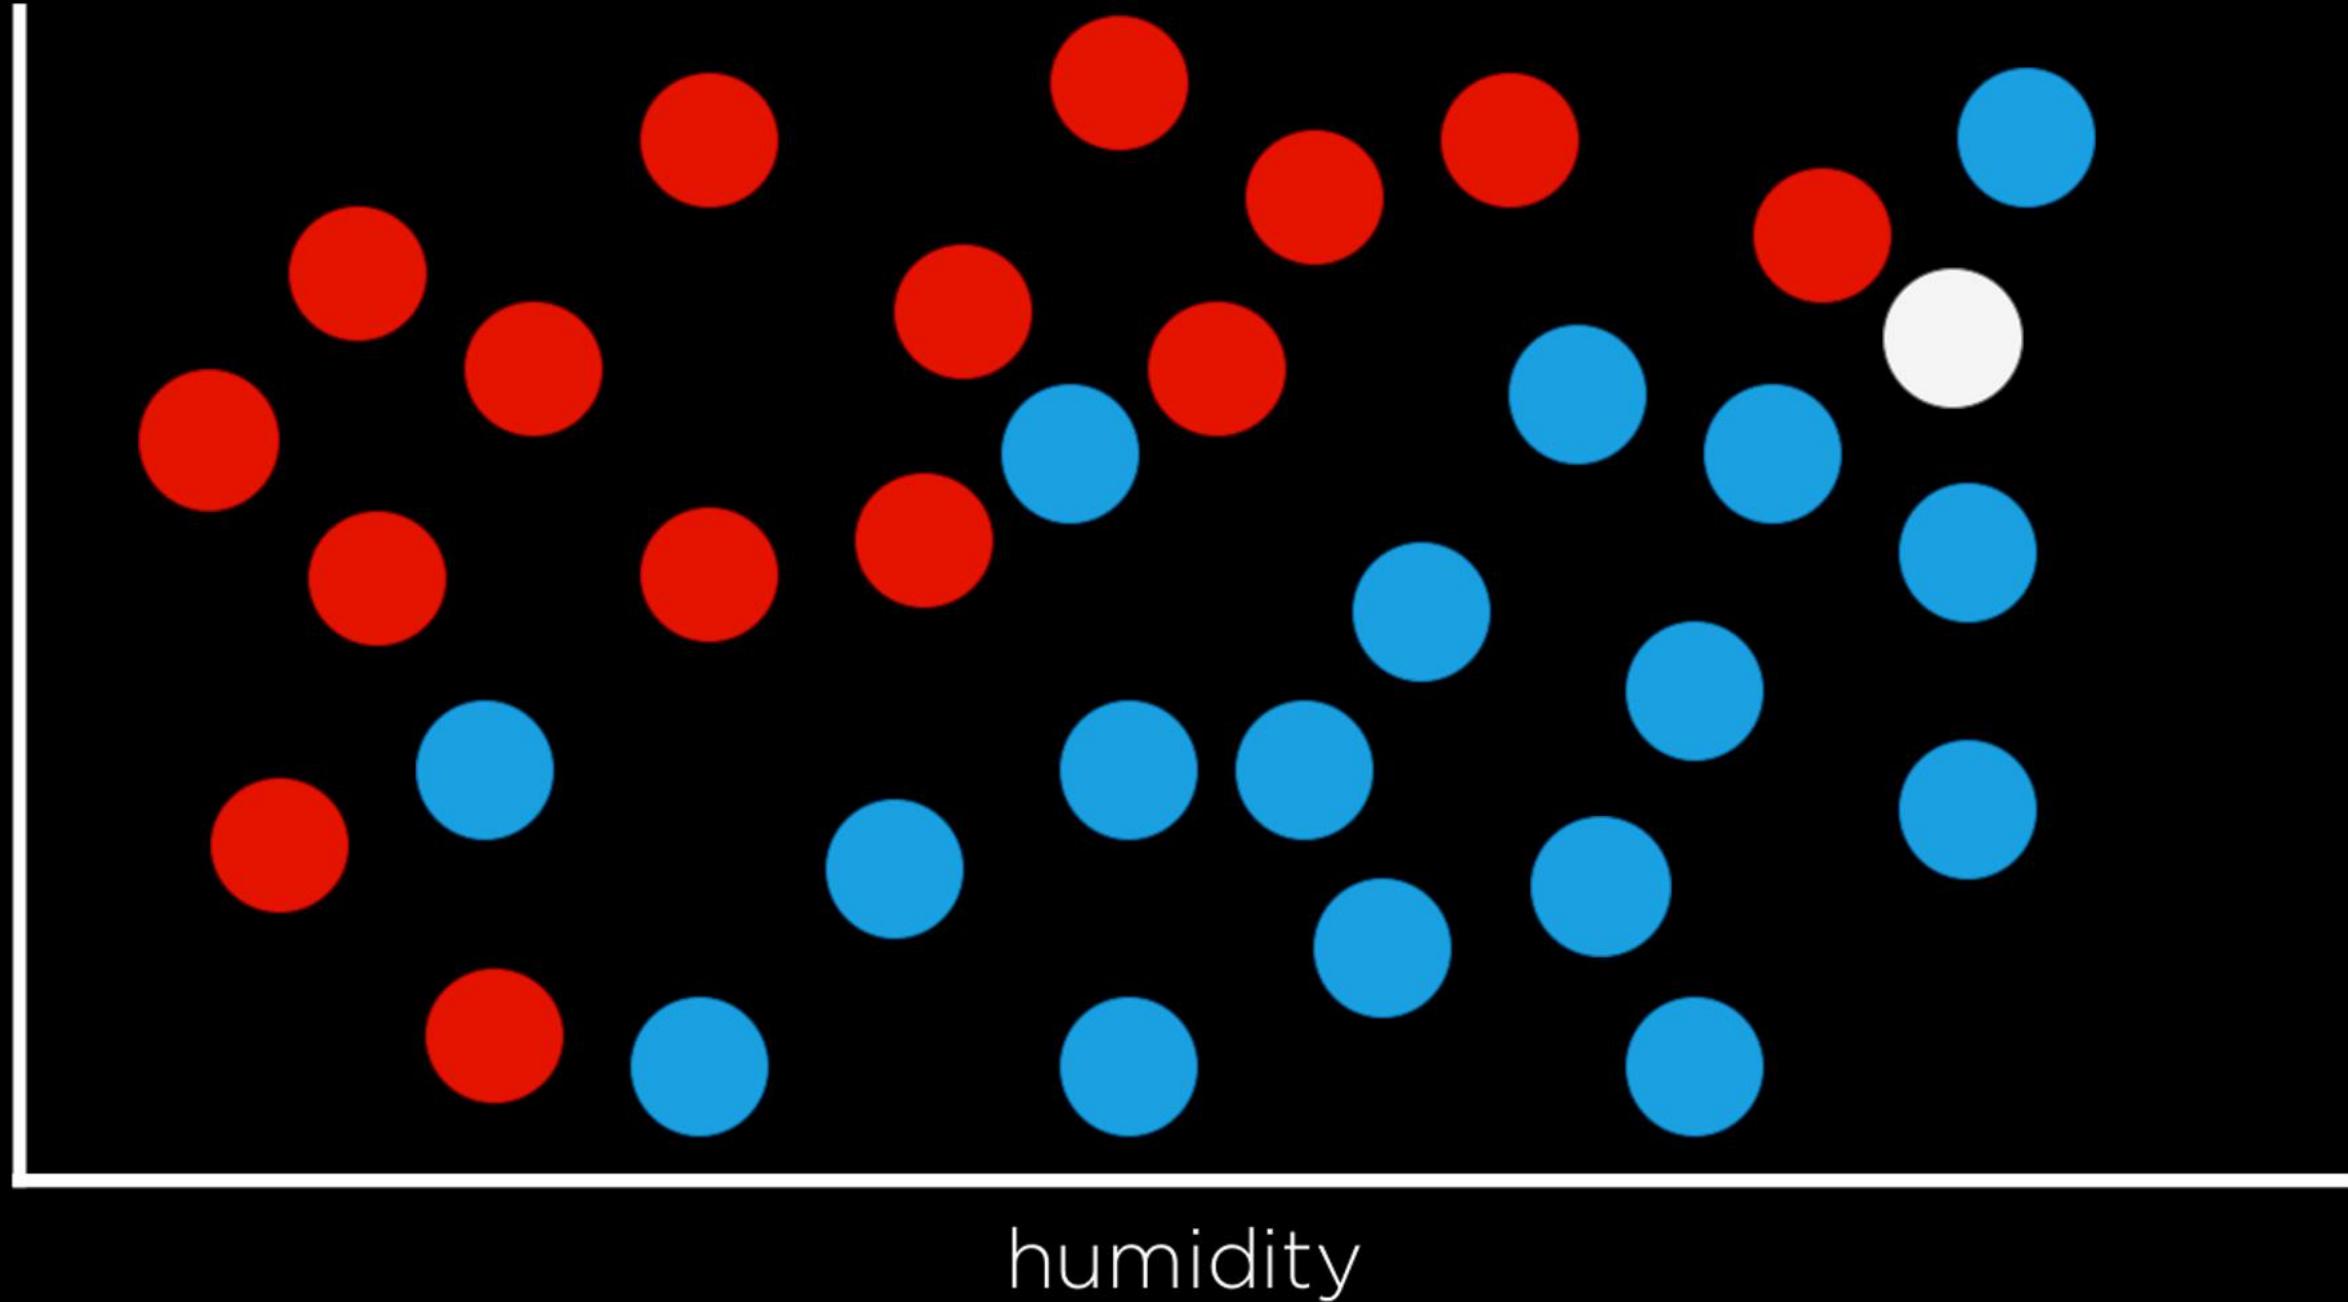


pressure

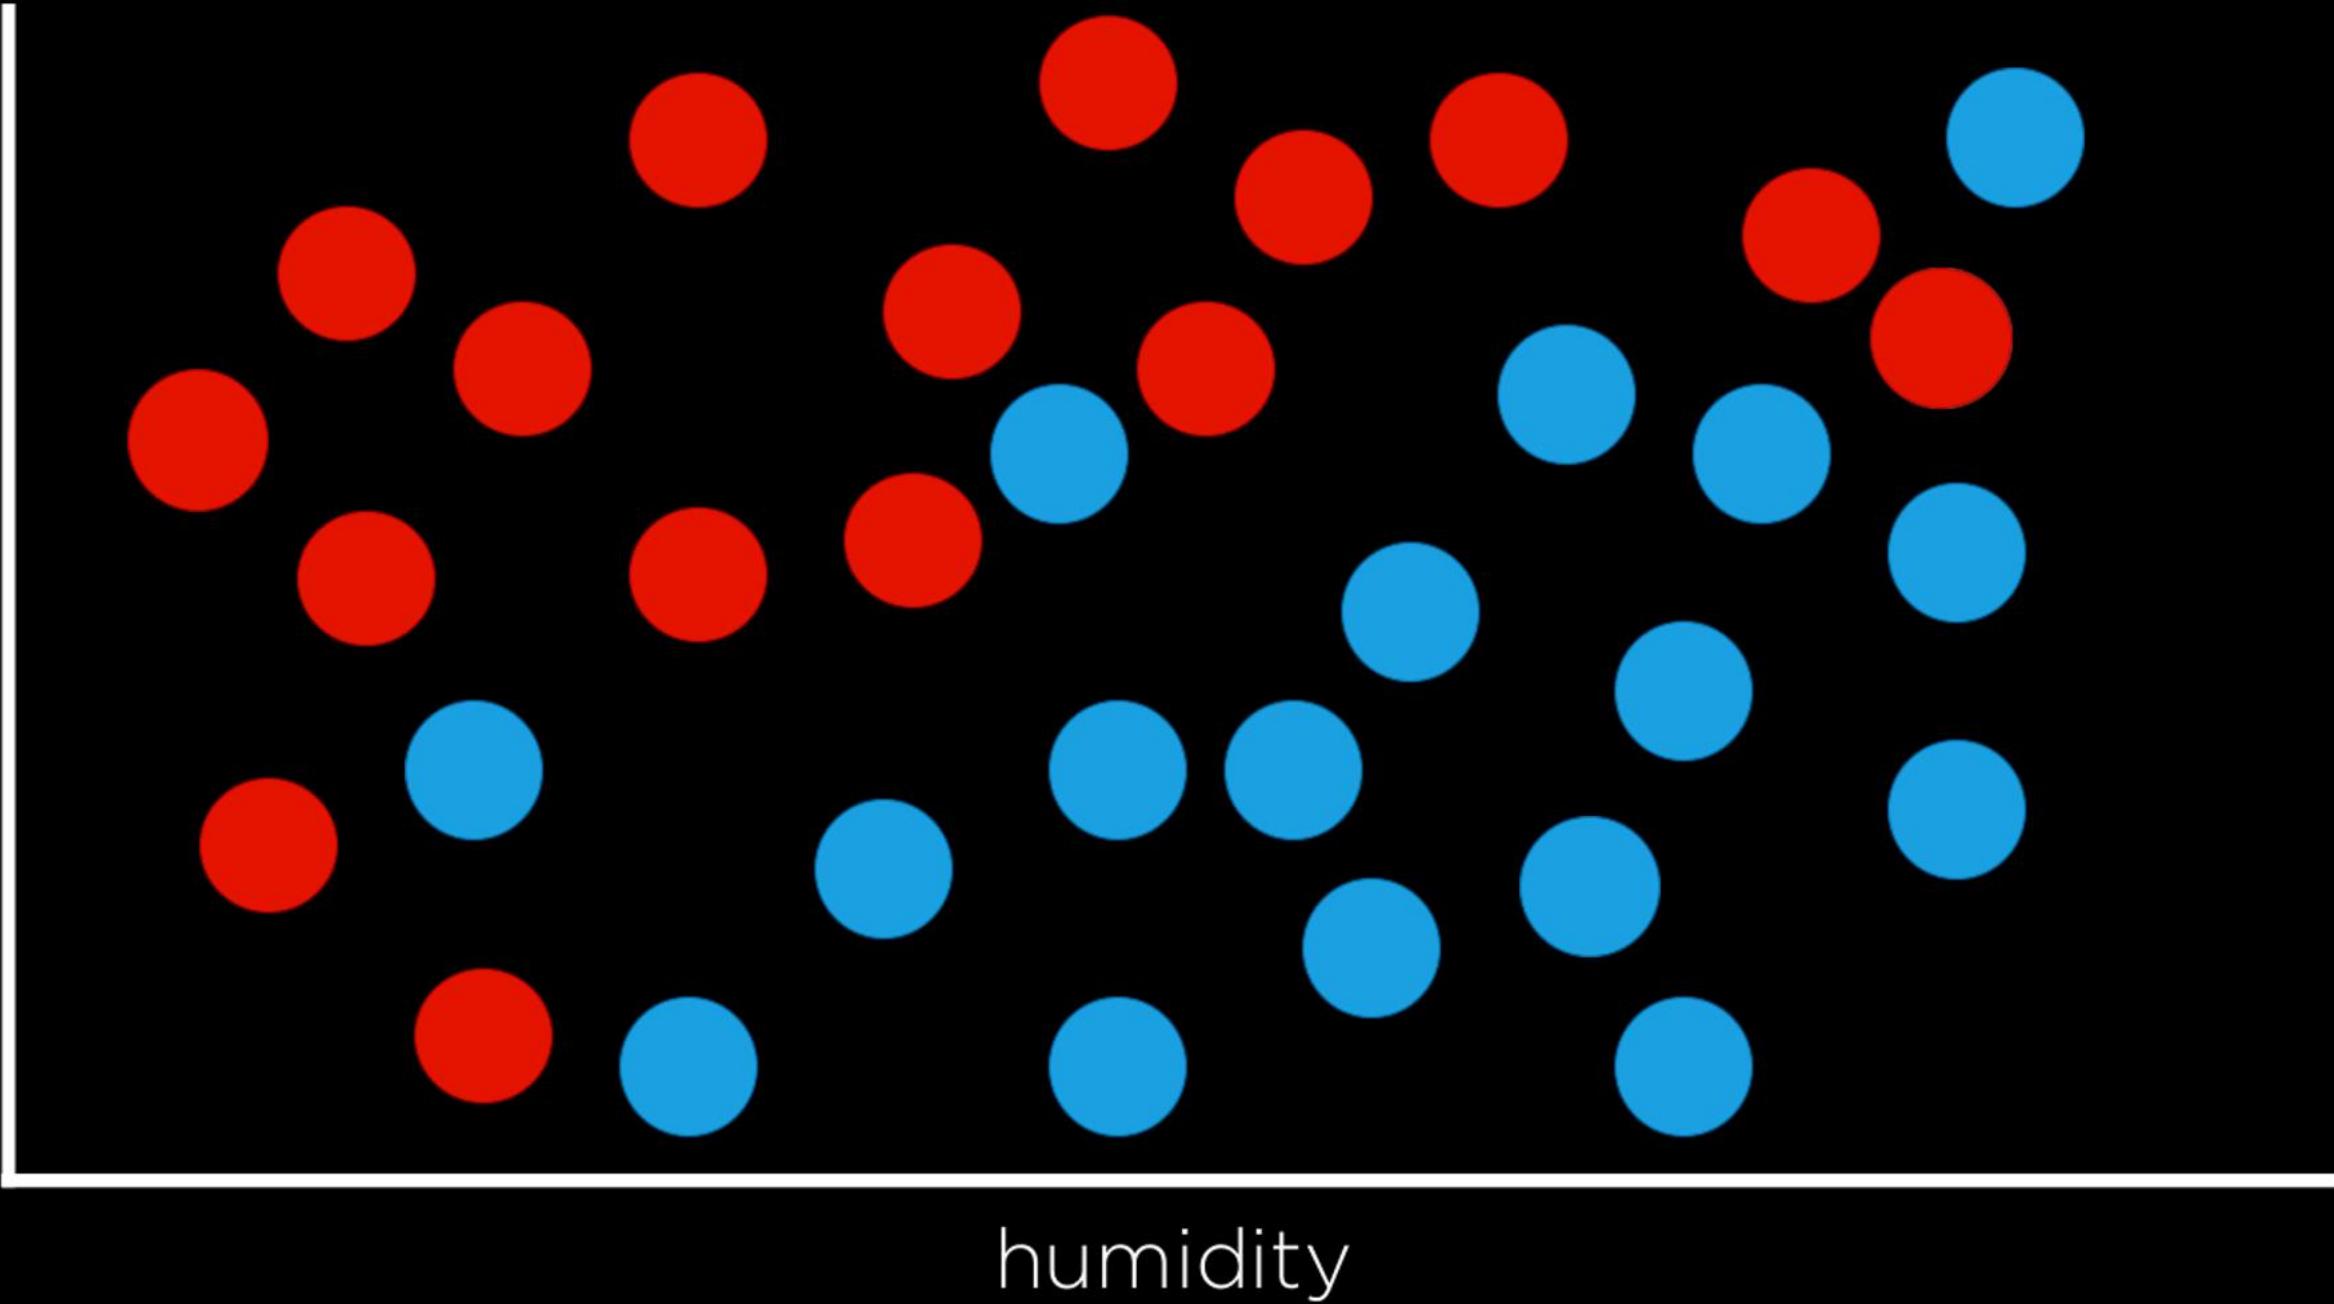


pressure





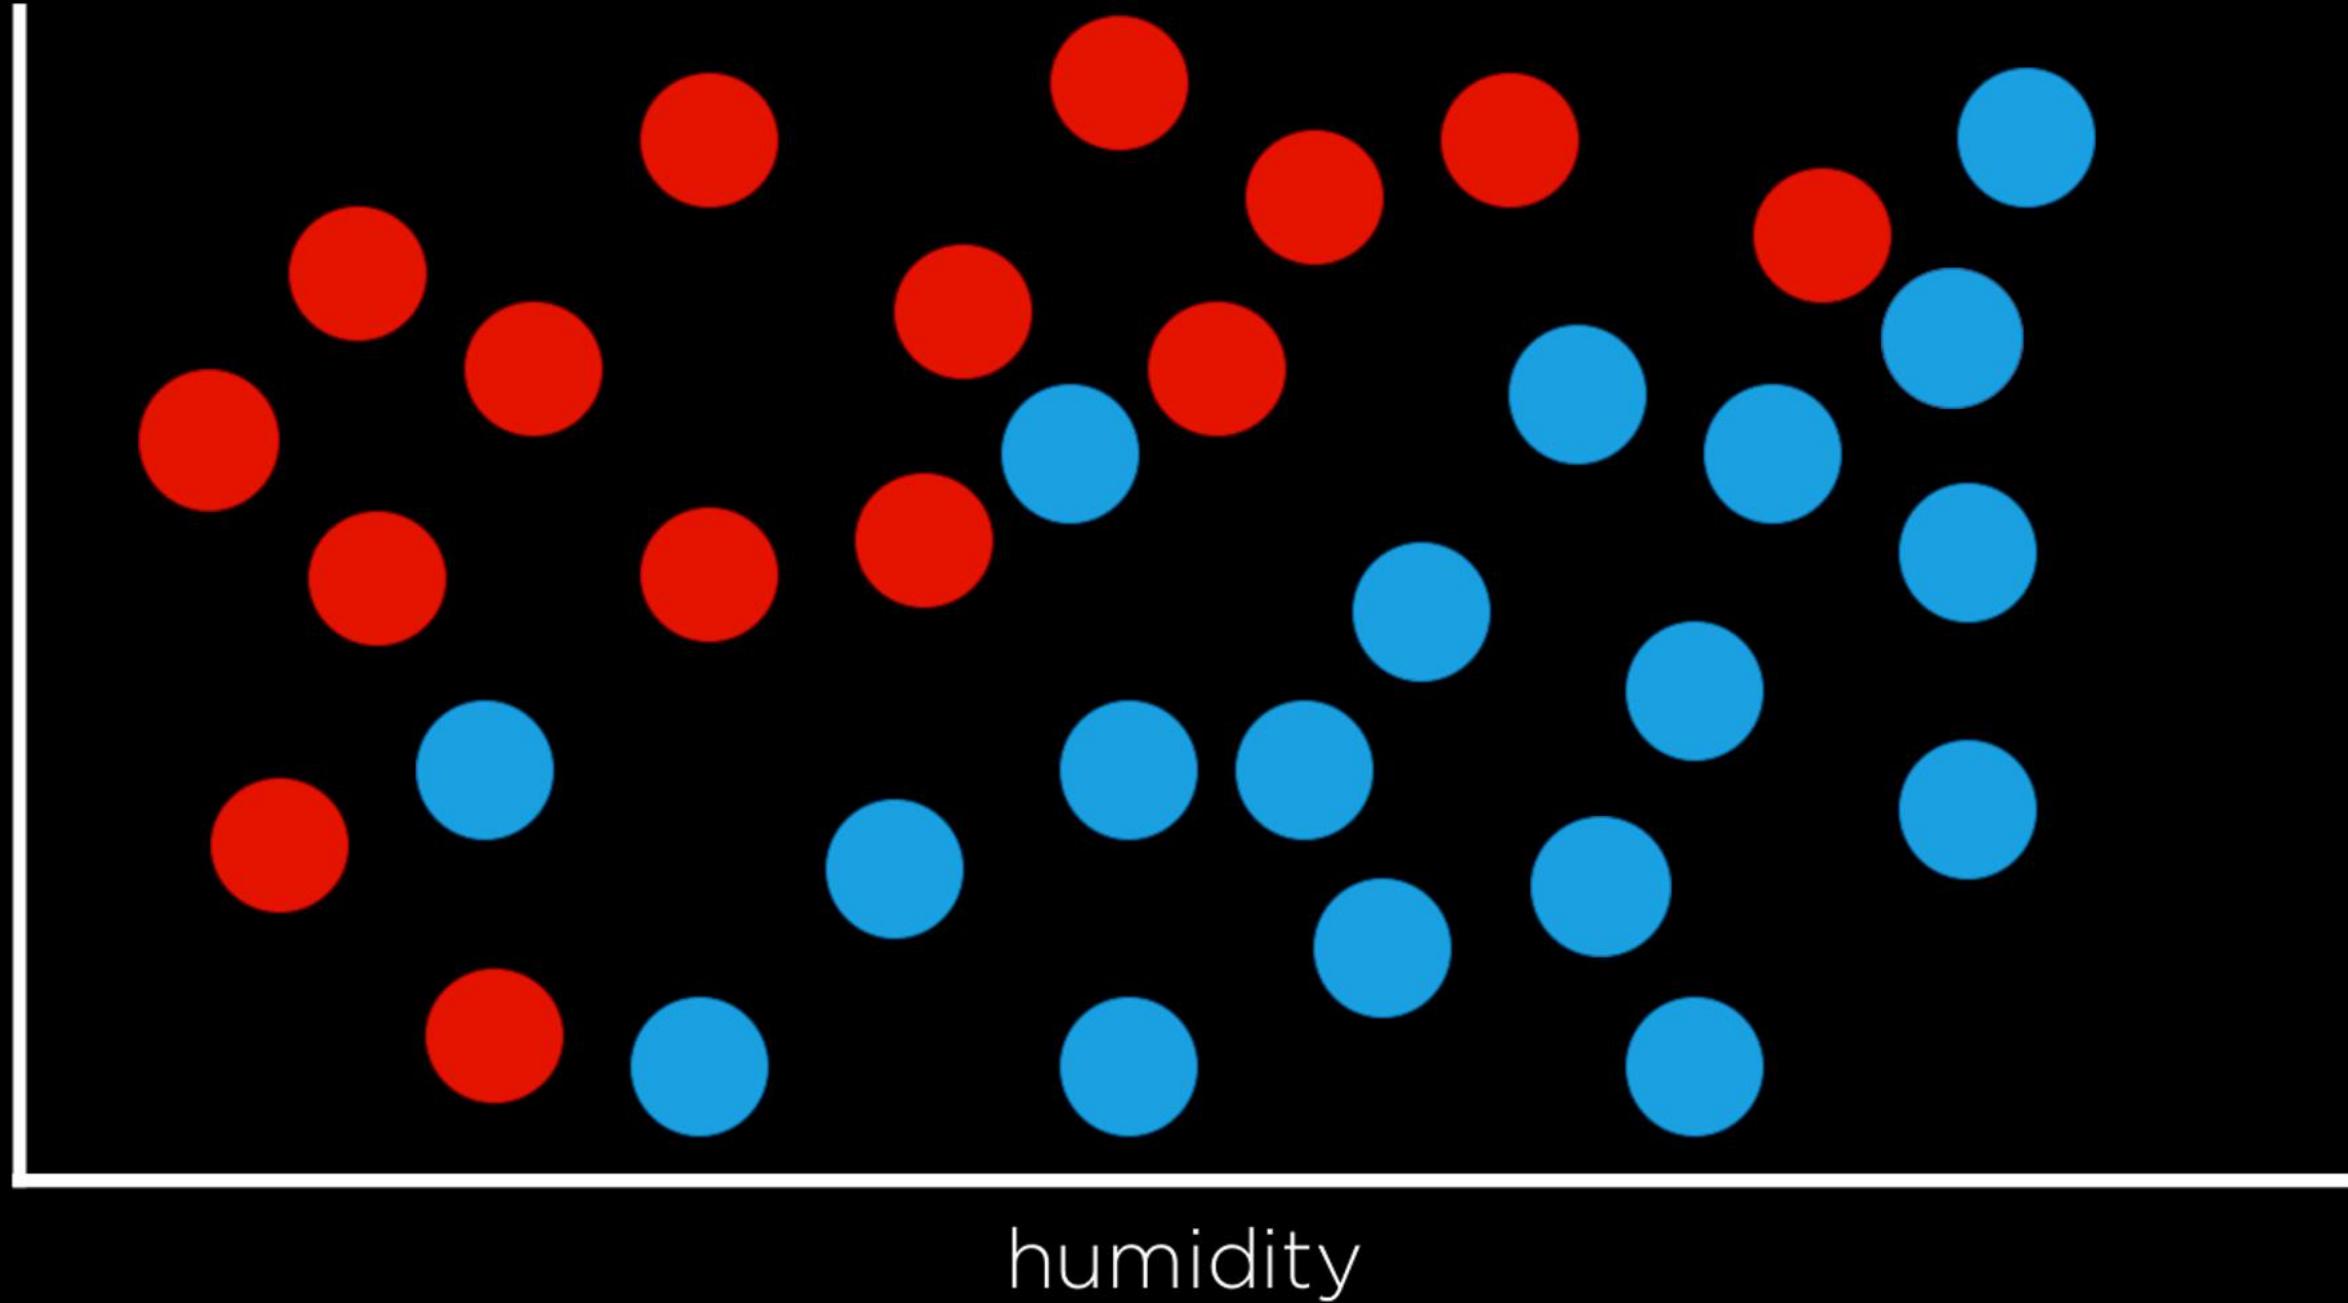
pressure



# **$k$ -nearest-neighbor classification**

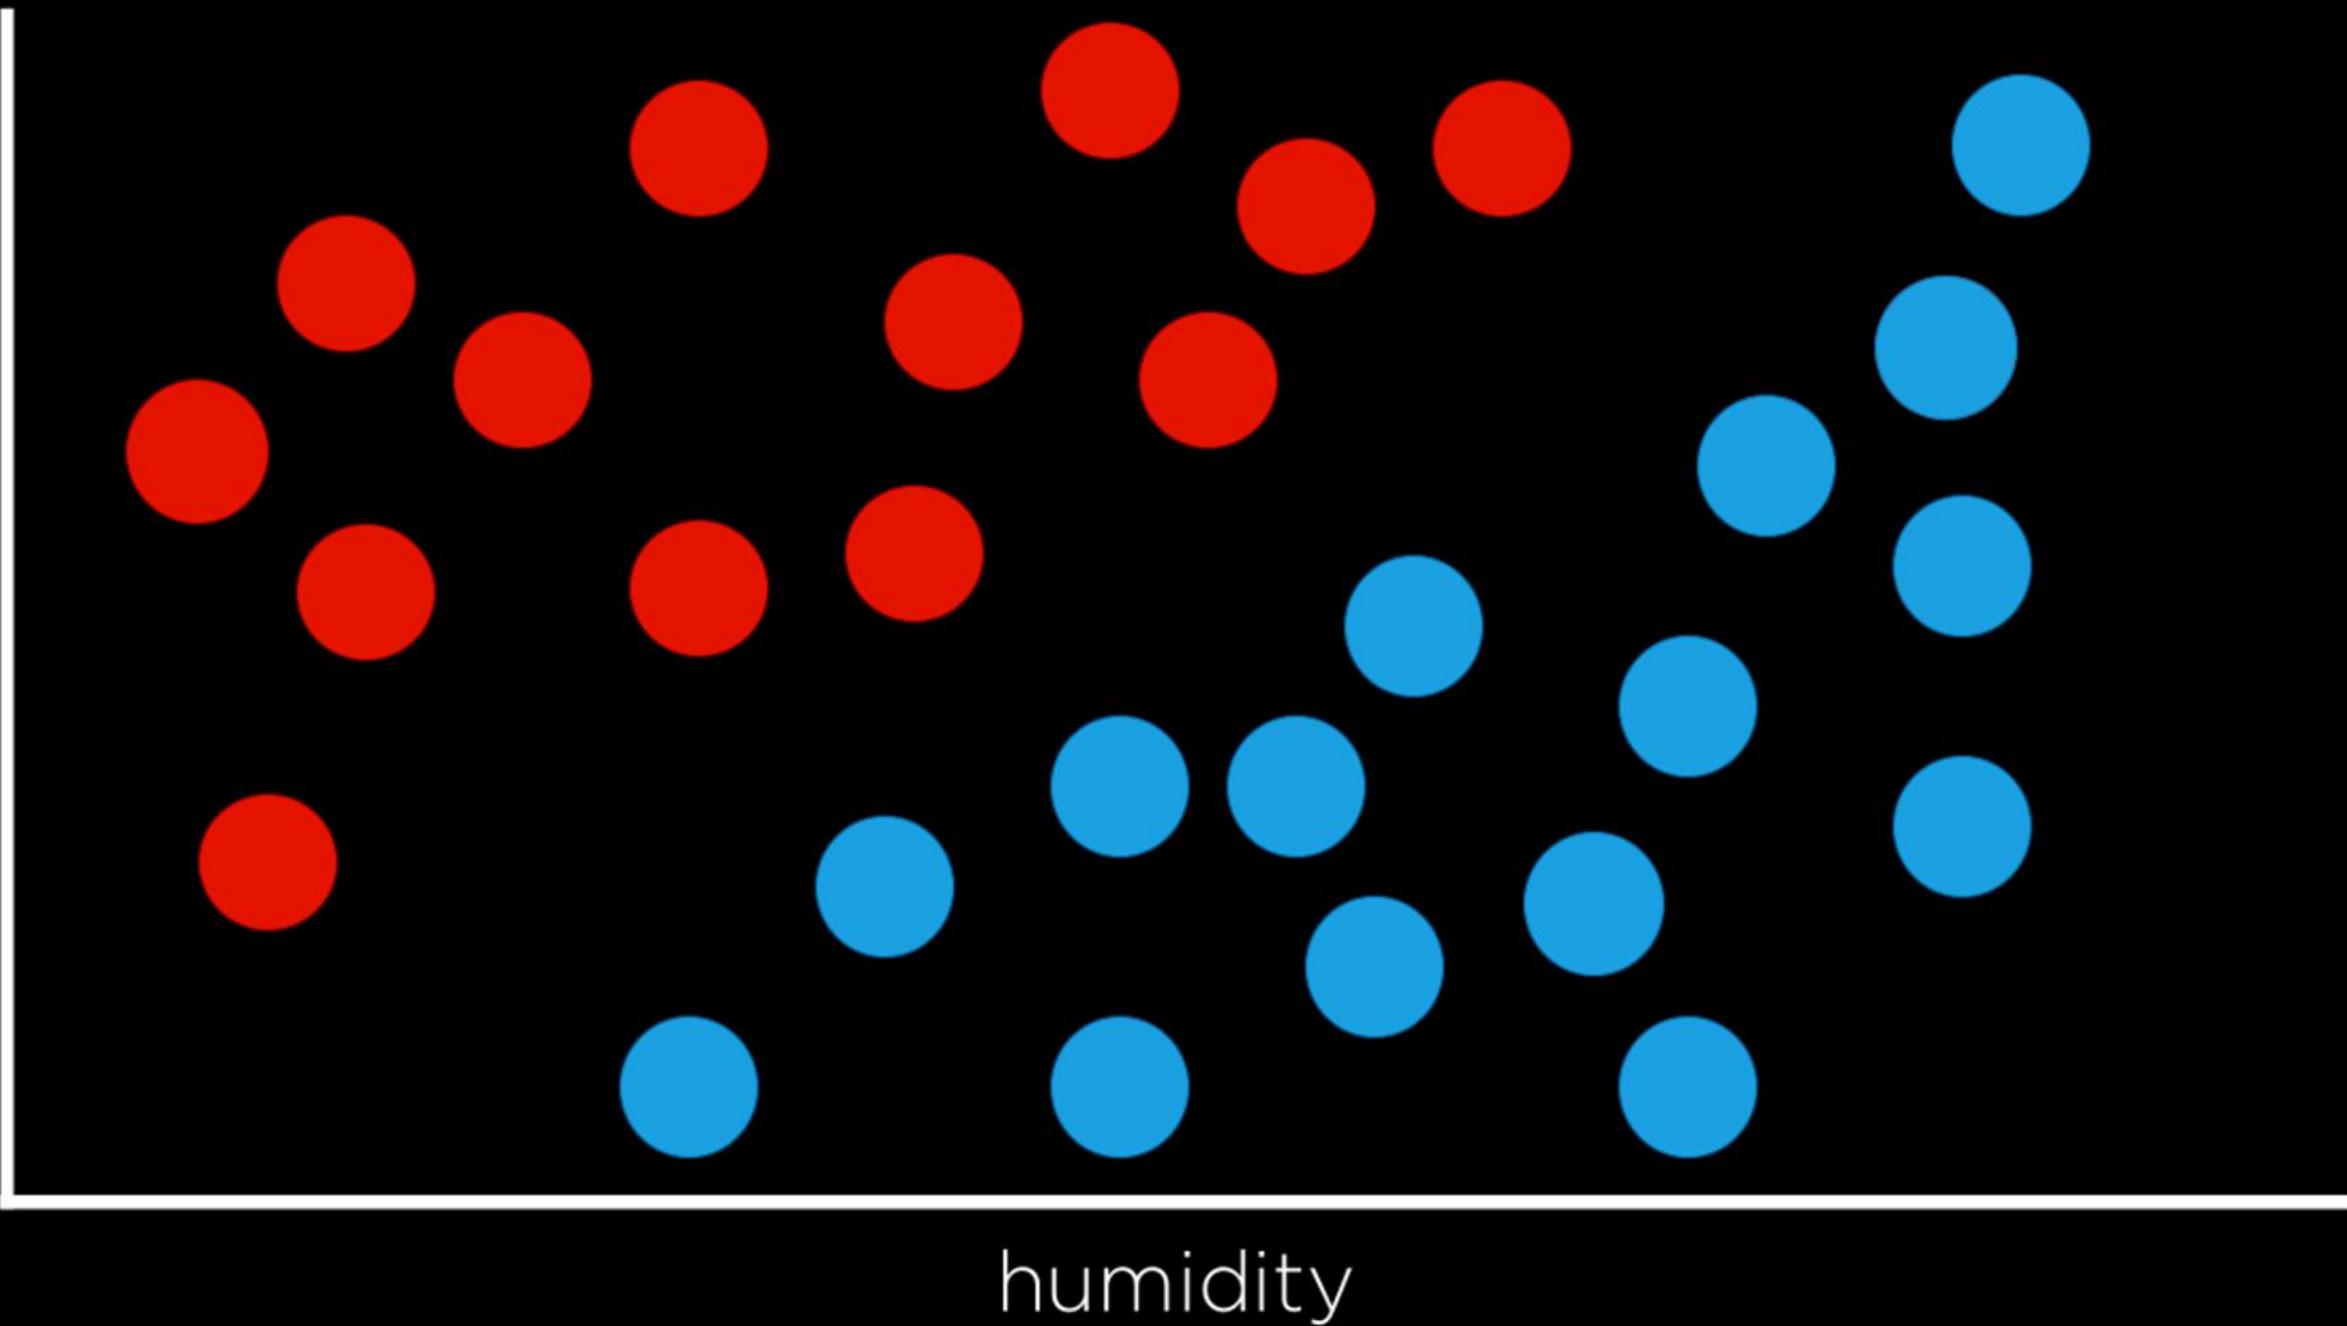
algorithm that, given an input, chooses the most common class out of the  $k$  nearest data points to that input

pressure

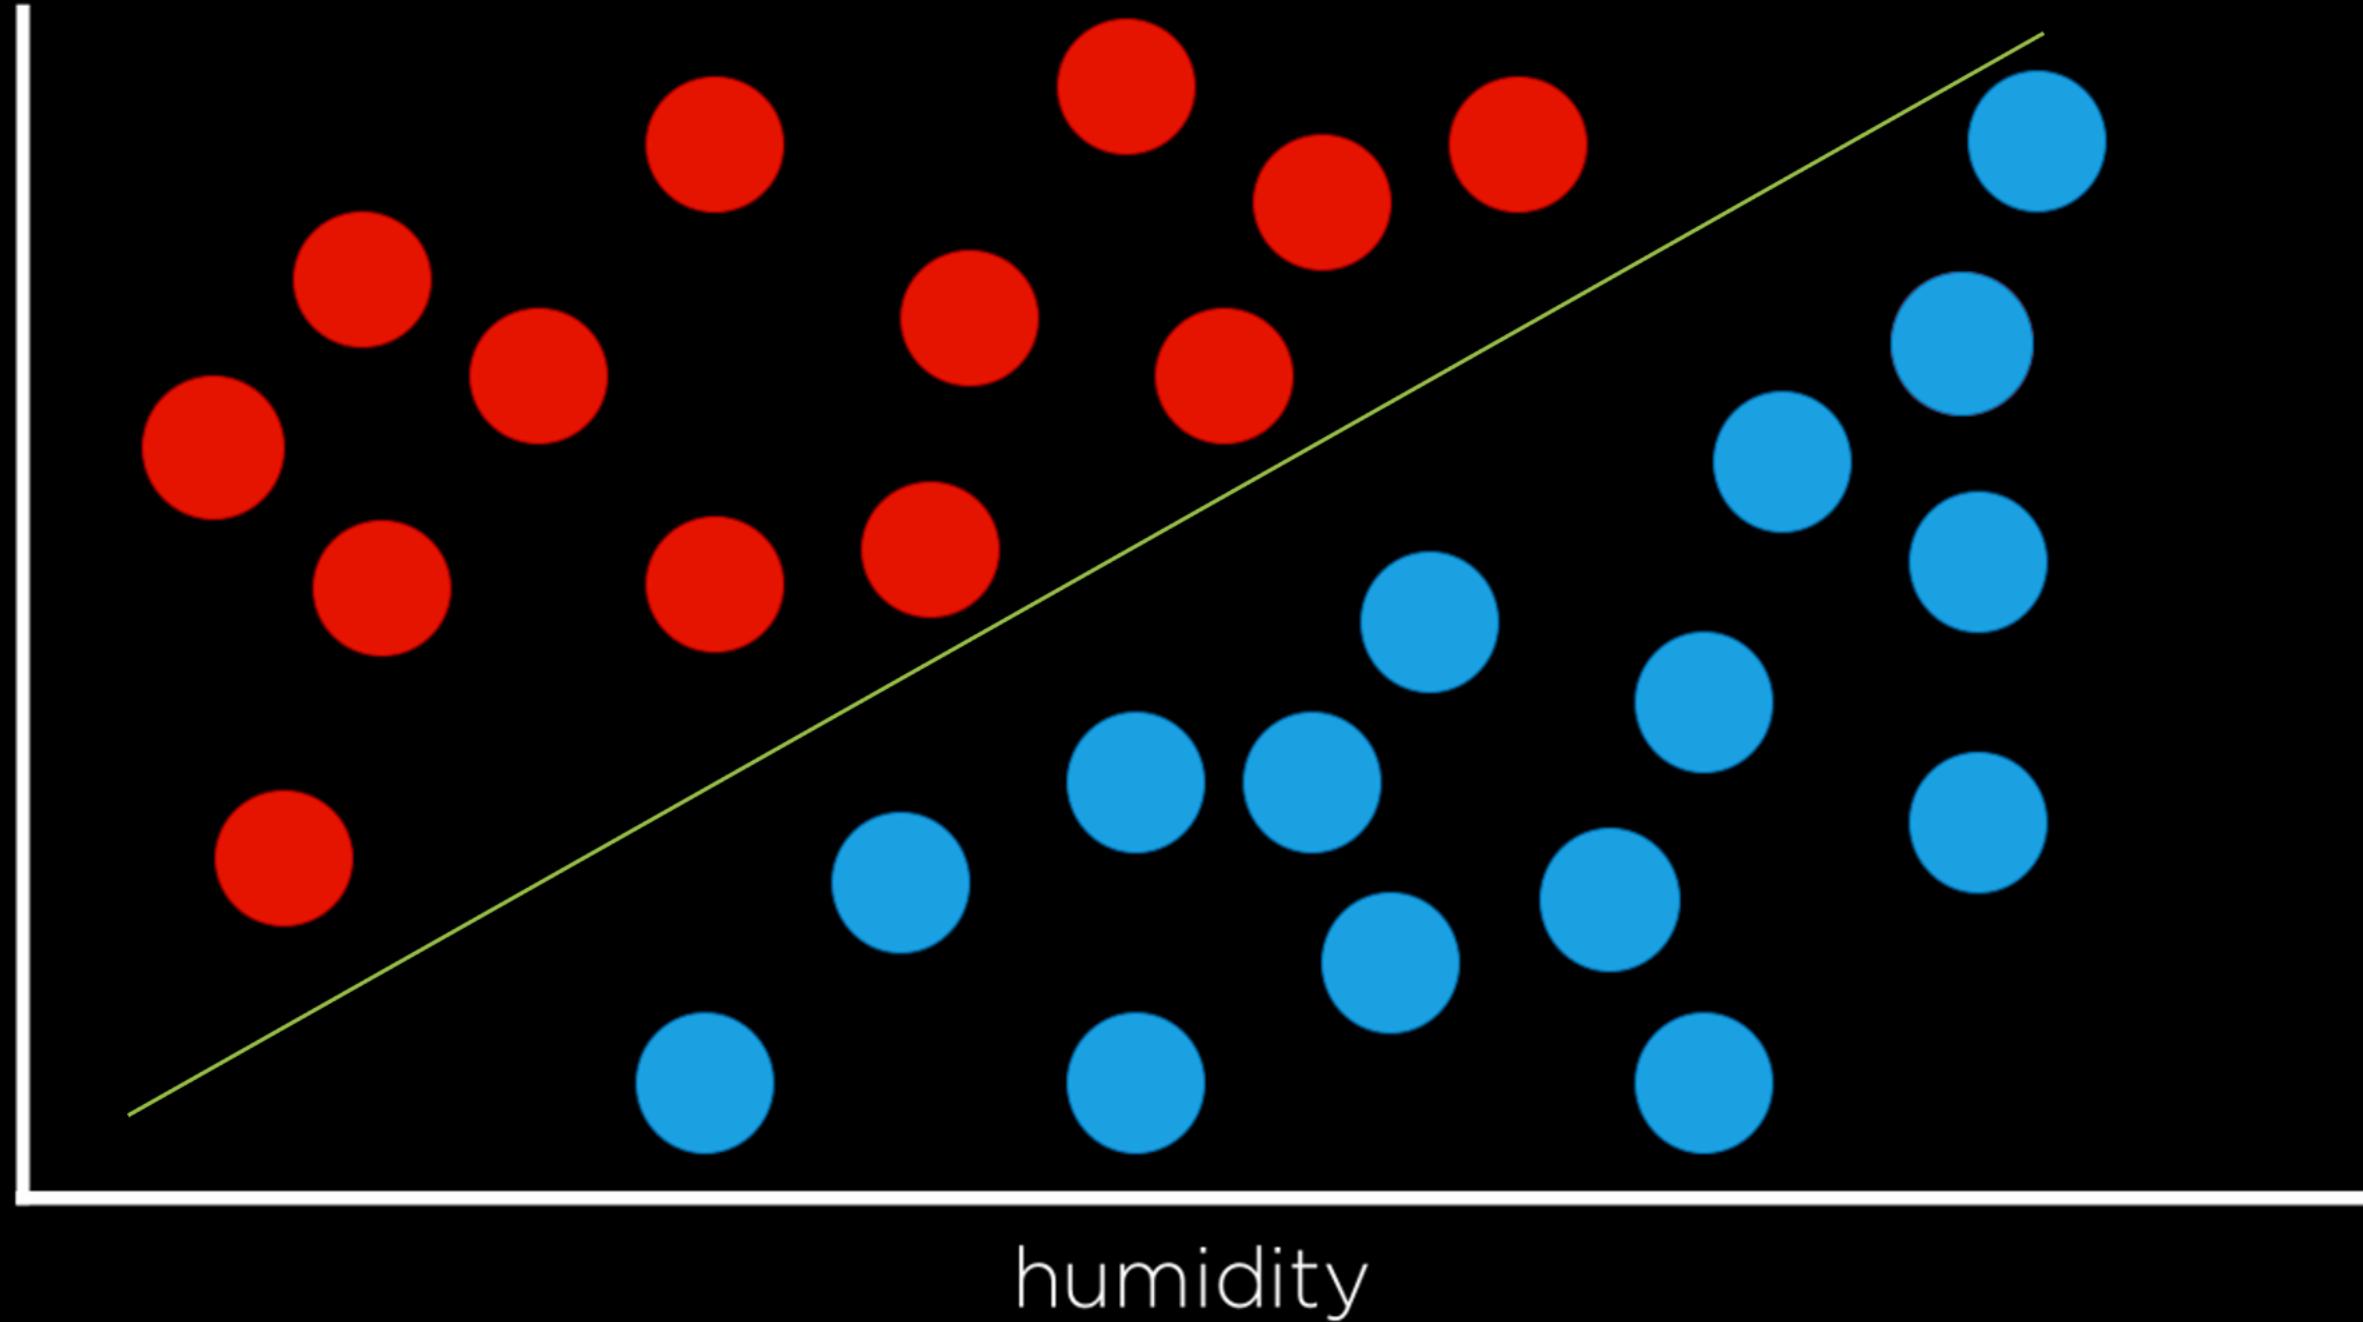


KNN has some drawbacks and challenges, such as computational expense, slow speed, memory and storage issues for large datasets, sensitivity to the choice of  $k$  and the distance metric, and susceptibility to the curse of dimensionality.

pressure

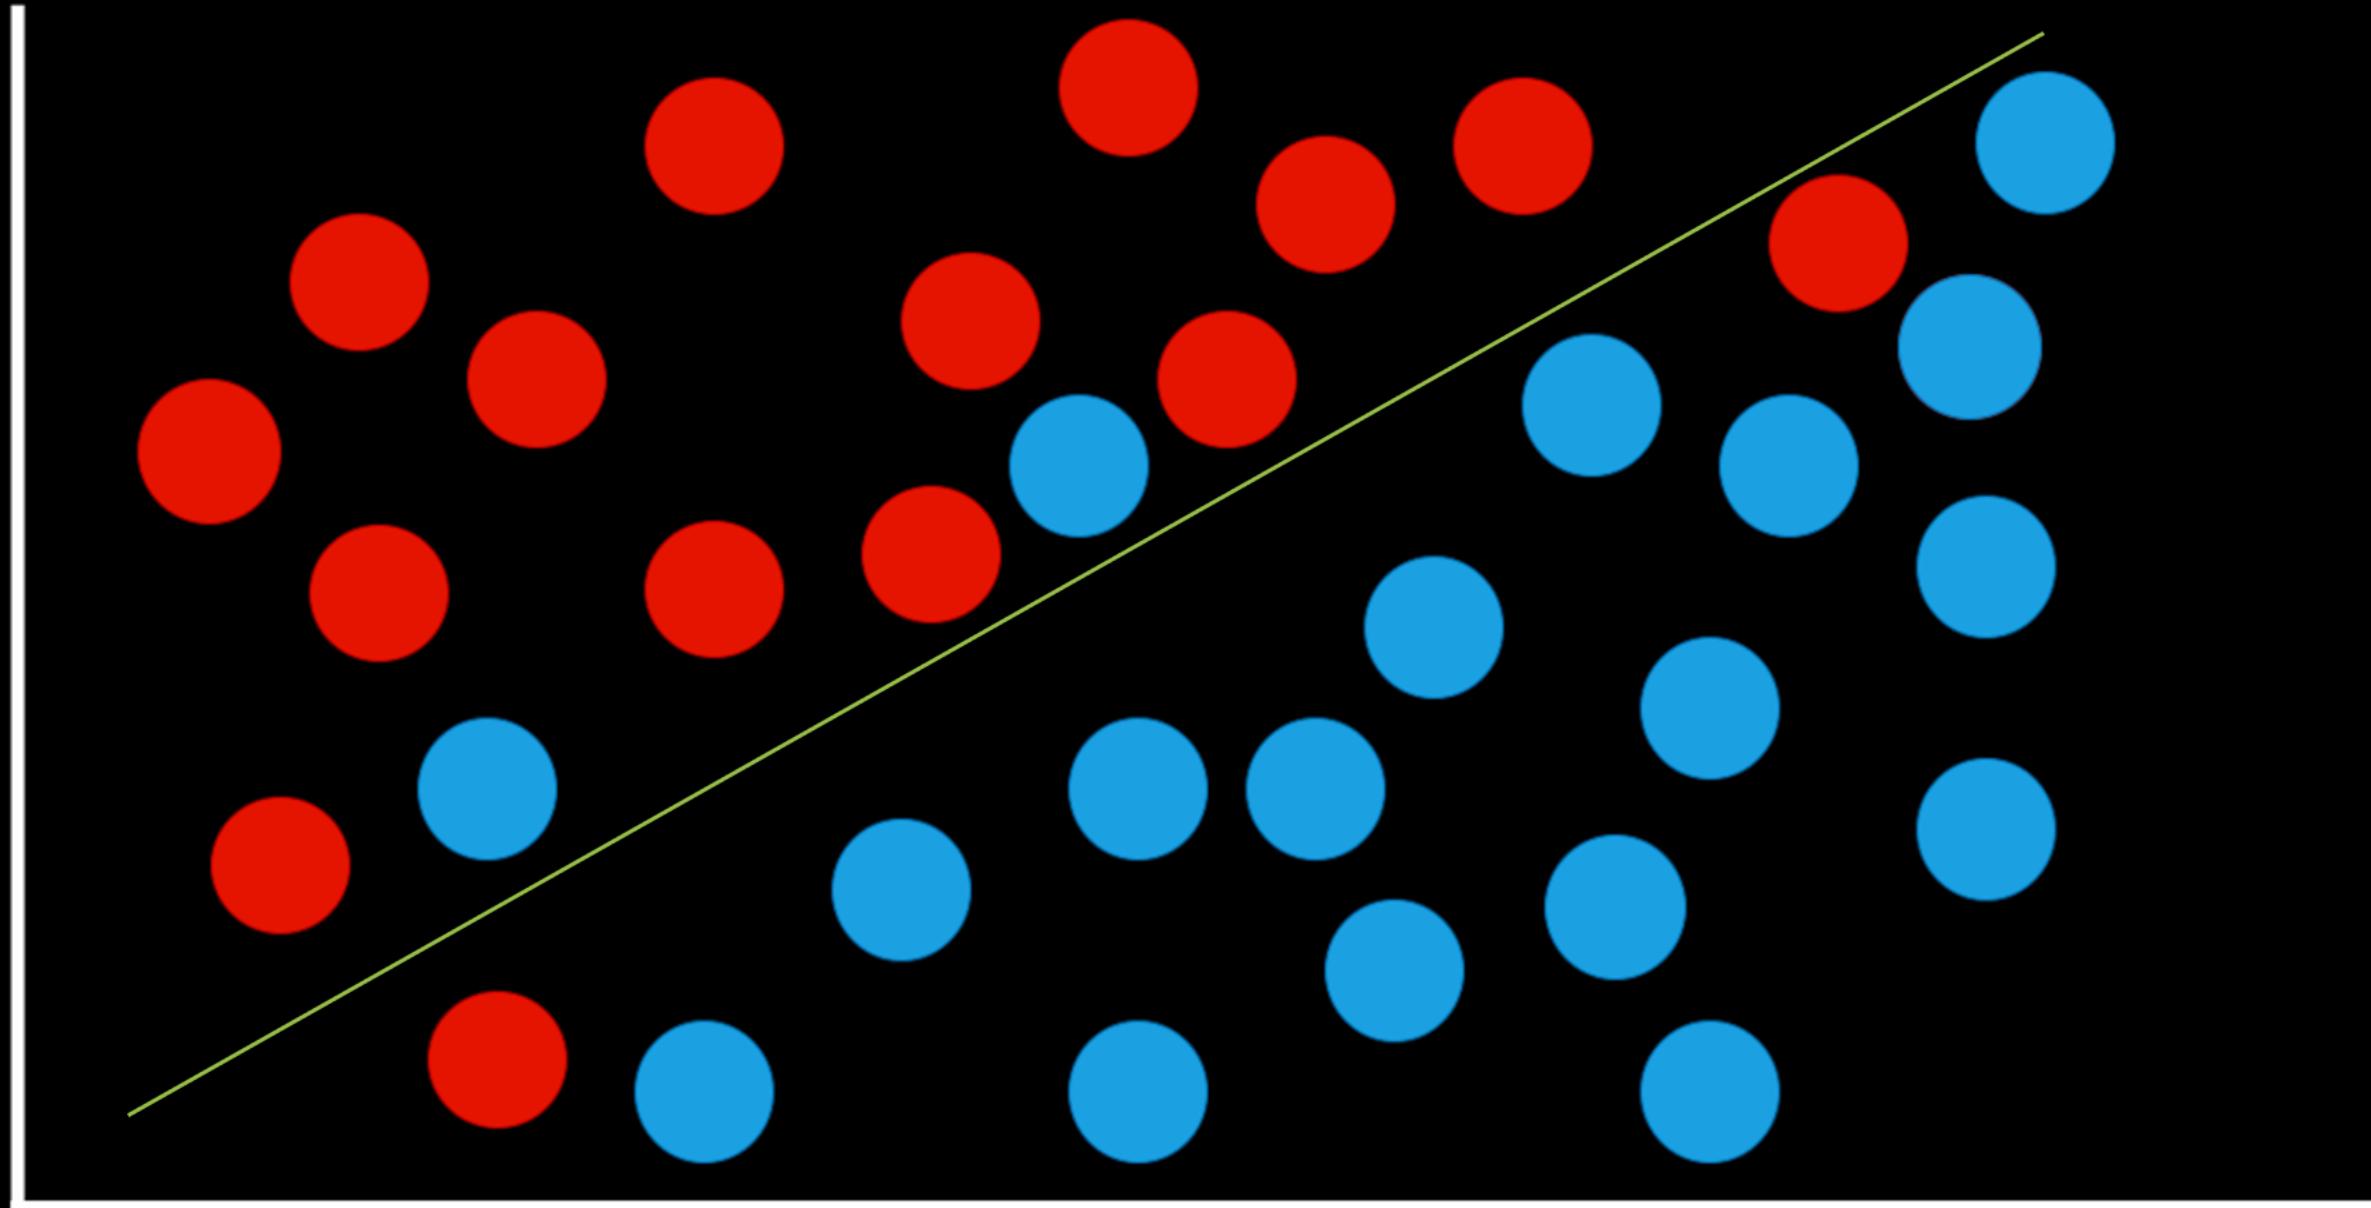


pressure



pressure

humidity



$x_1$  = Humidity

$x_2$  = Pressure

$$h(x_1, x_2) = \begin{cases} \text{Rain if } w_0 + w_1x_1 + w_2x_2 \geq 0 \\ \text{No Rain otherwise} \end{cases}$$

Weight Vector  $\mathbf{w}$ :  $(w_0, w_1, w_2)$

Input Vector  $\mathbf{x}$ :  $(1, x_1, x_2)$

$\mathbf{w} \cdot \mathbf{x}$ :  $w_0 + w_1x_1 + w_2x_2$

$$h(x_1, x_2) = \begin{cases} 1 & \text{if } w_0 + w_1x_1 + w_2x_2 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Weight Vector  $\mathbf{w}$ :  $(w_0, w_1, w_2)$

Input Vector  $\mathbf{x}$ :  $(1, x_1, x_2)$

$\mathbf{w} \cdot \mathbf{x}$ :  $w_0 + w_1x_1 + w_2x_2$

$$h_{\mathbf{w}}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

# perceptron learning rule

Given data point  $(\mathbf{x}, y)$ , update each weight according to:

$$w_i = w_i + \alpha(y - h_{\mathbf{w}}(\mathbf{x})) \times x_i$$

# perceptron learning rule

Given data point  $(\mathbf{x}, y)$ , update each weight according to:

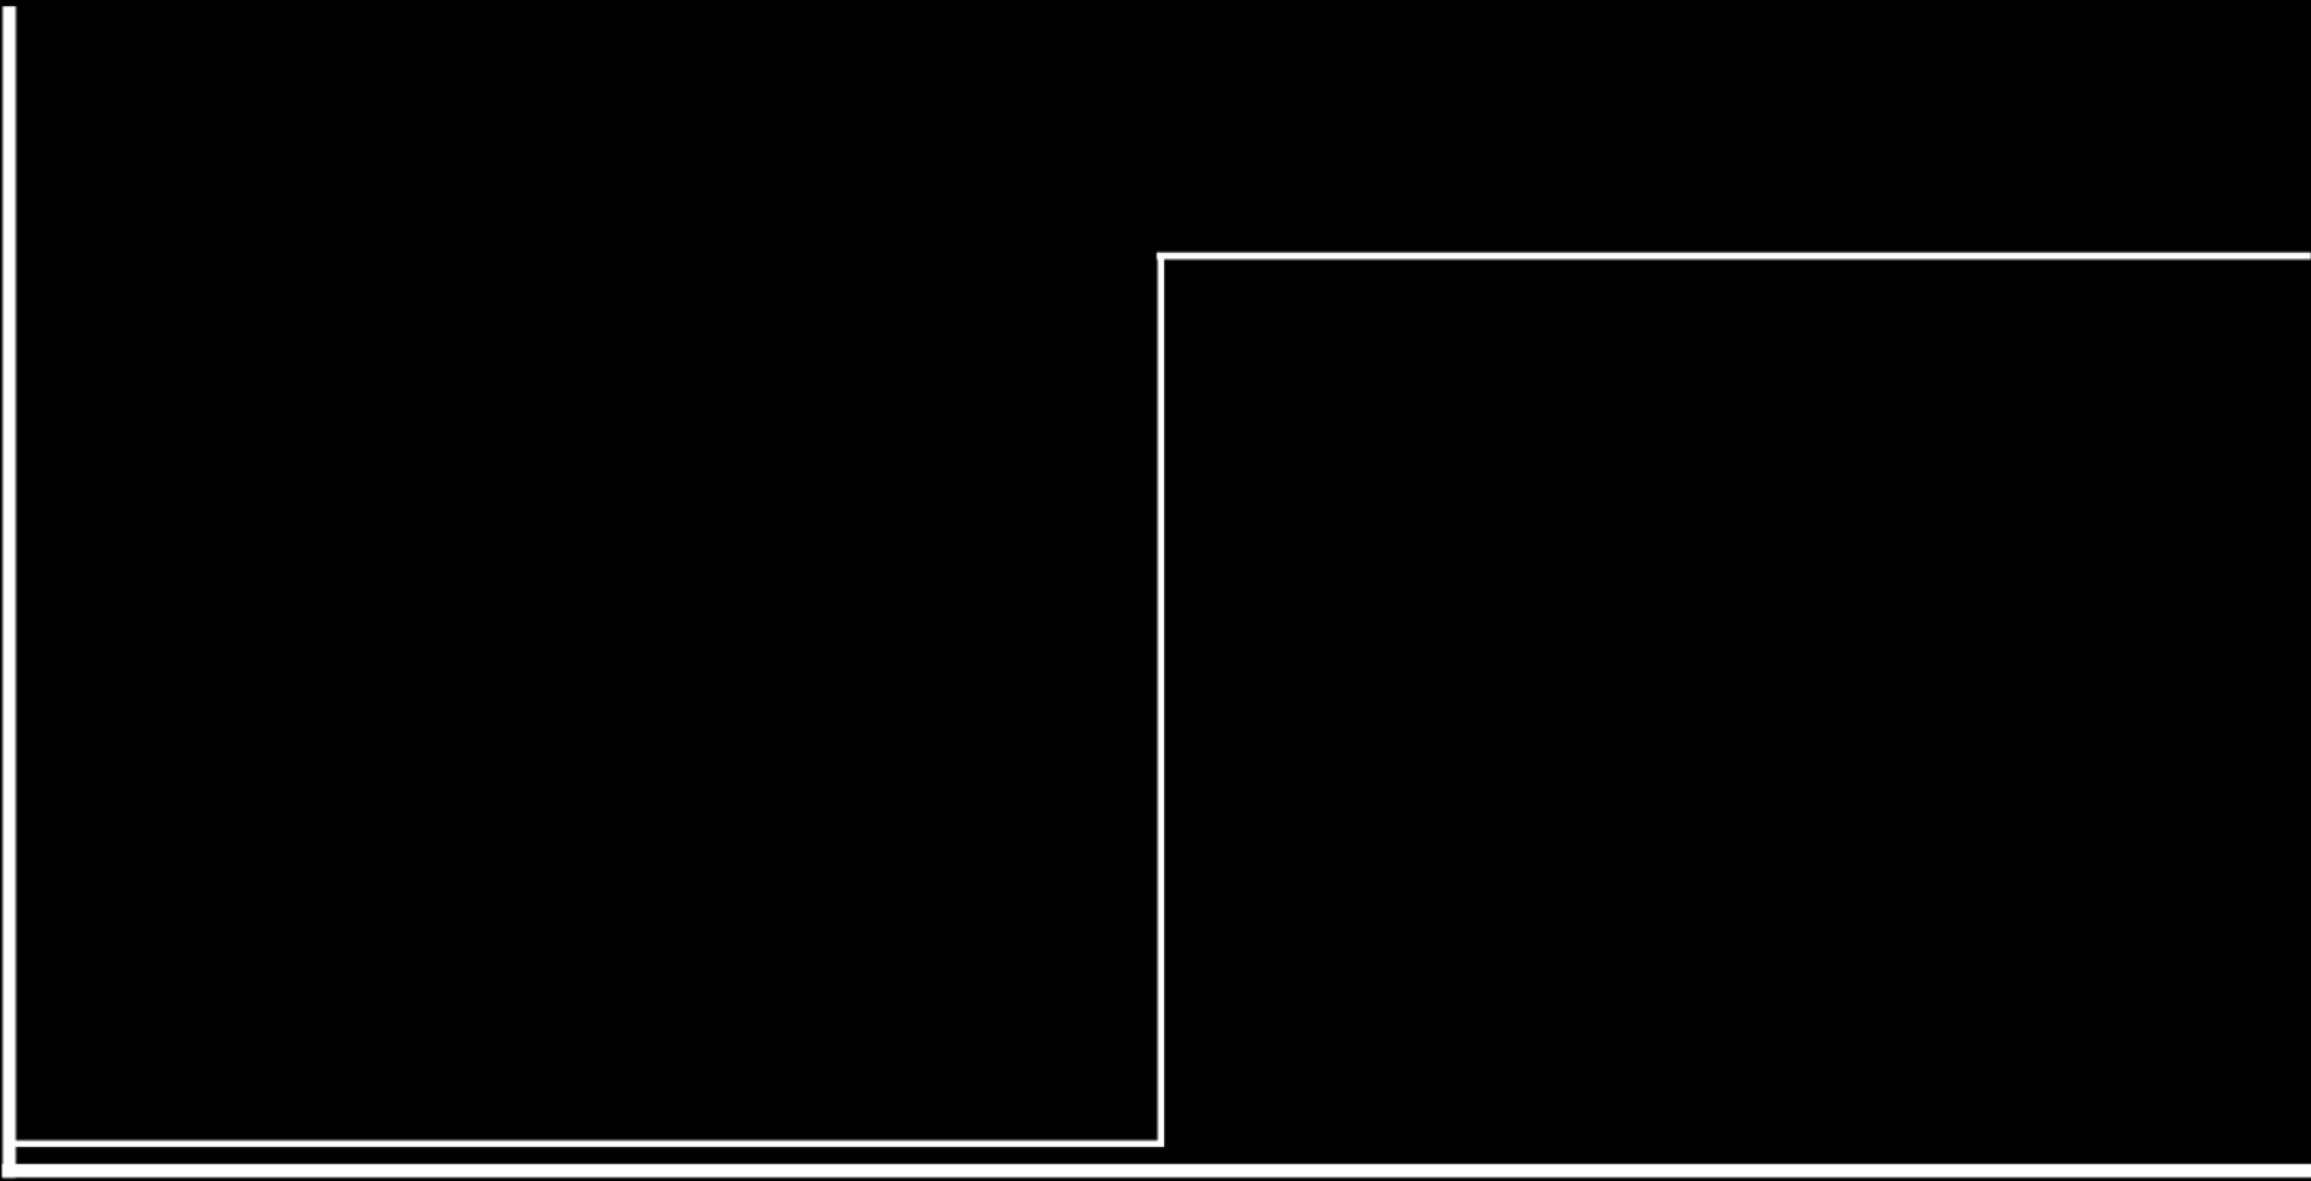
$$w_i = w_i + \alpha(\text{actual value} - \text{estimate}) \times x_i$$

Output

1

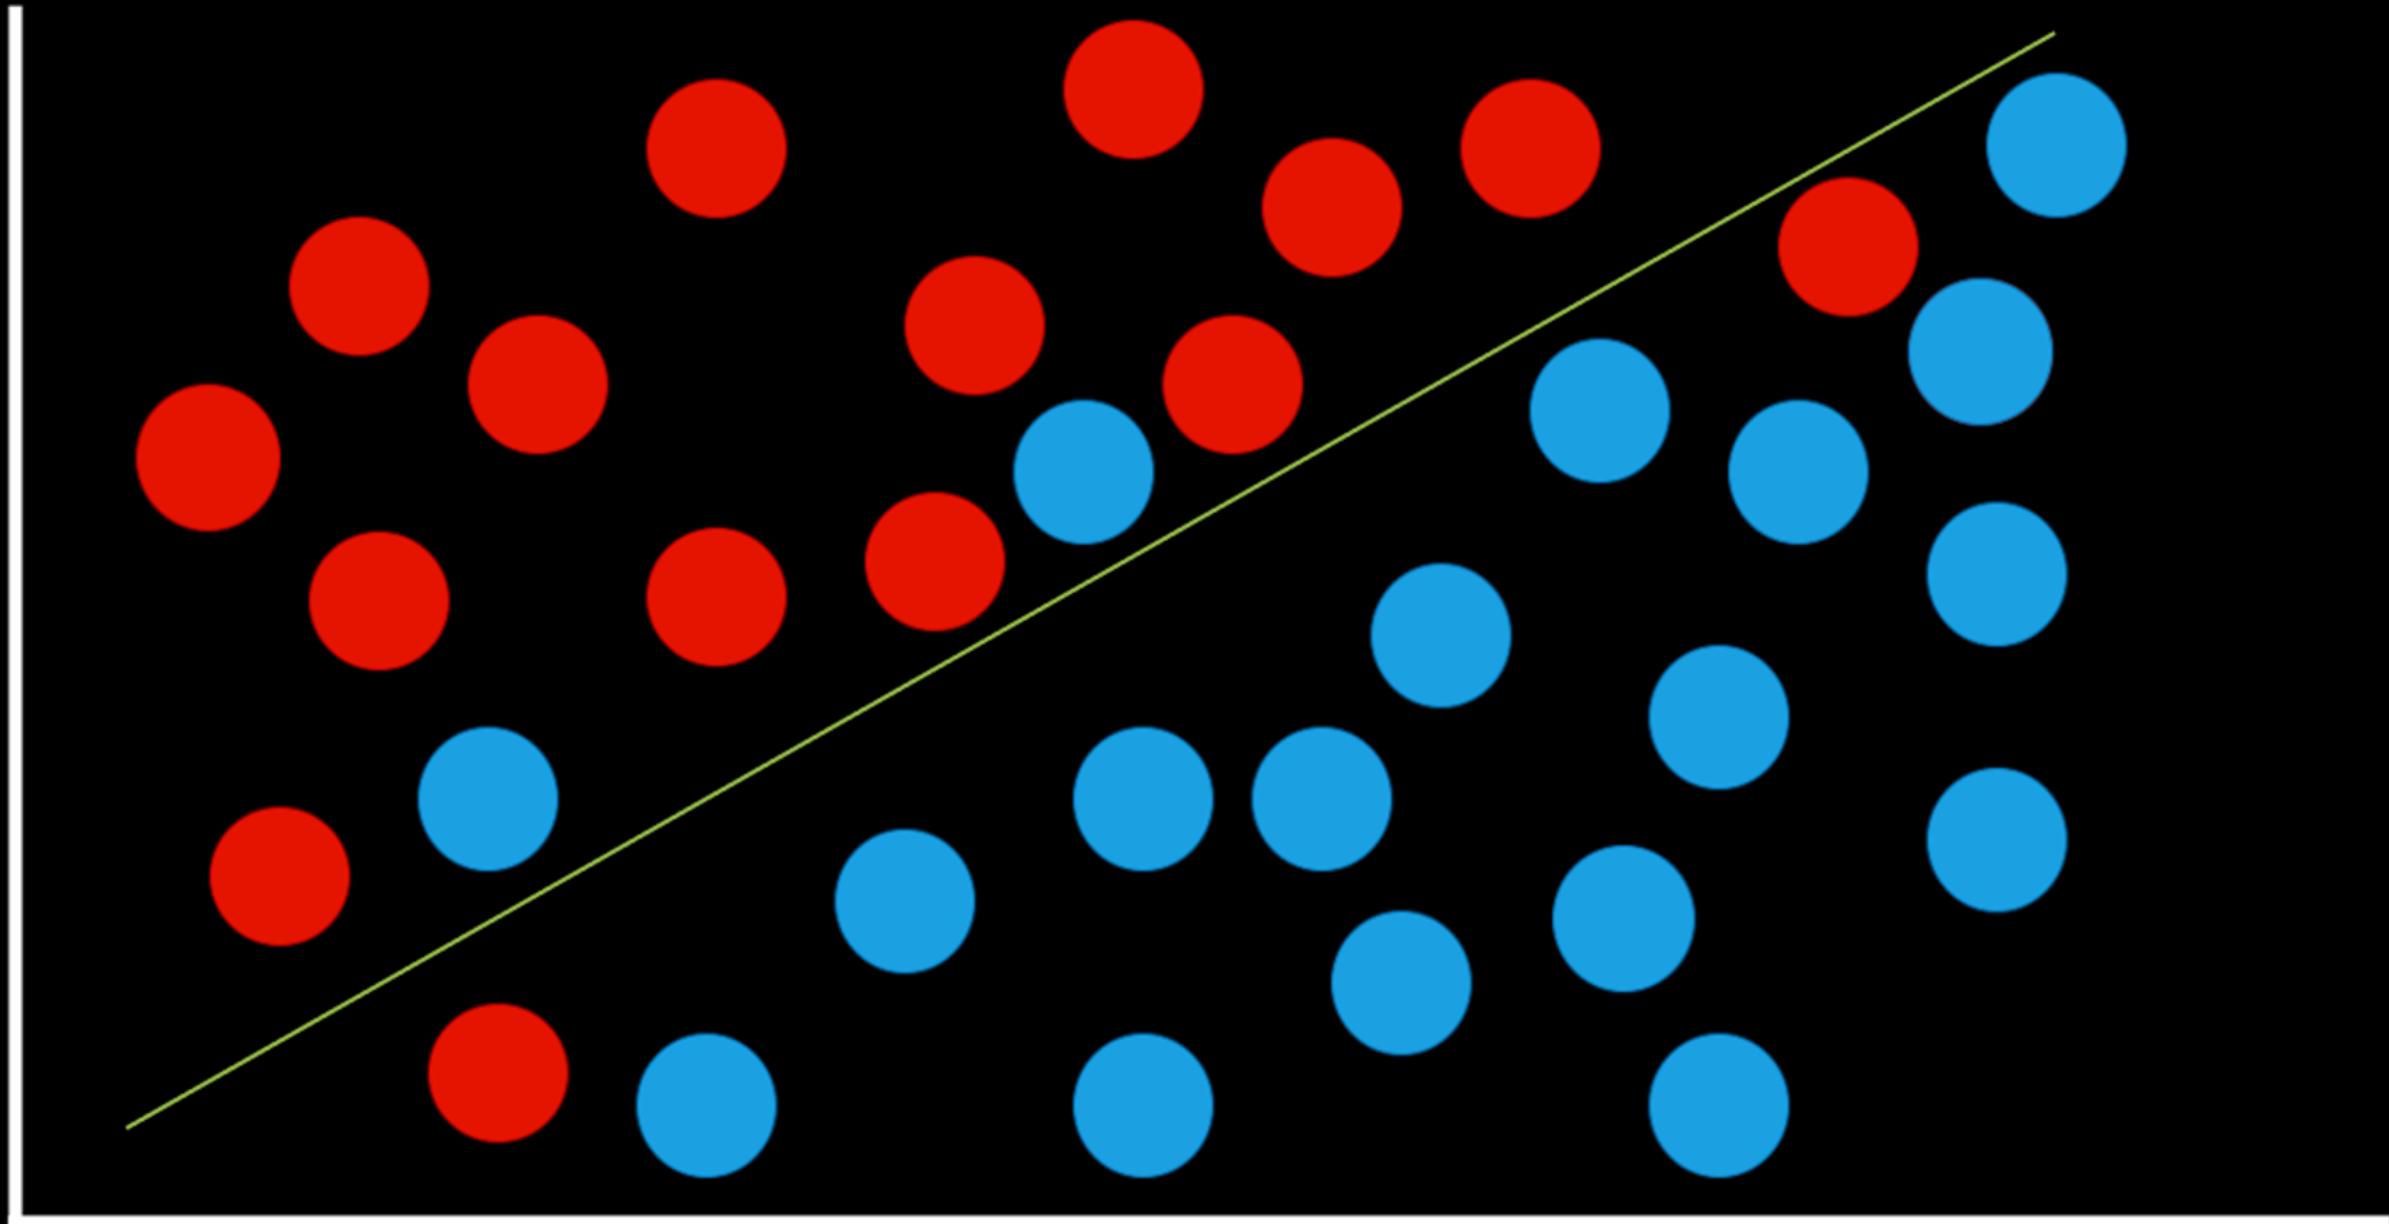
0

$\mathbf{W} \cdot \mathbf{X}$



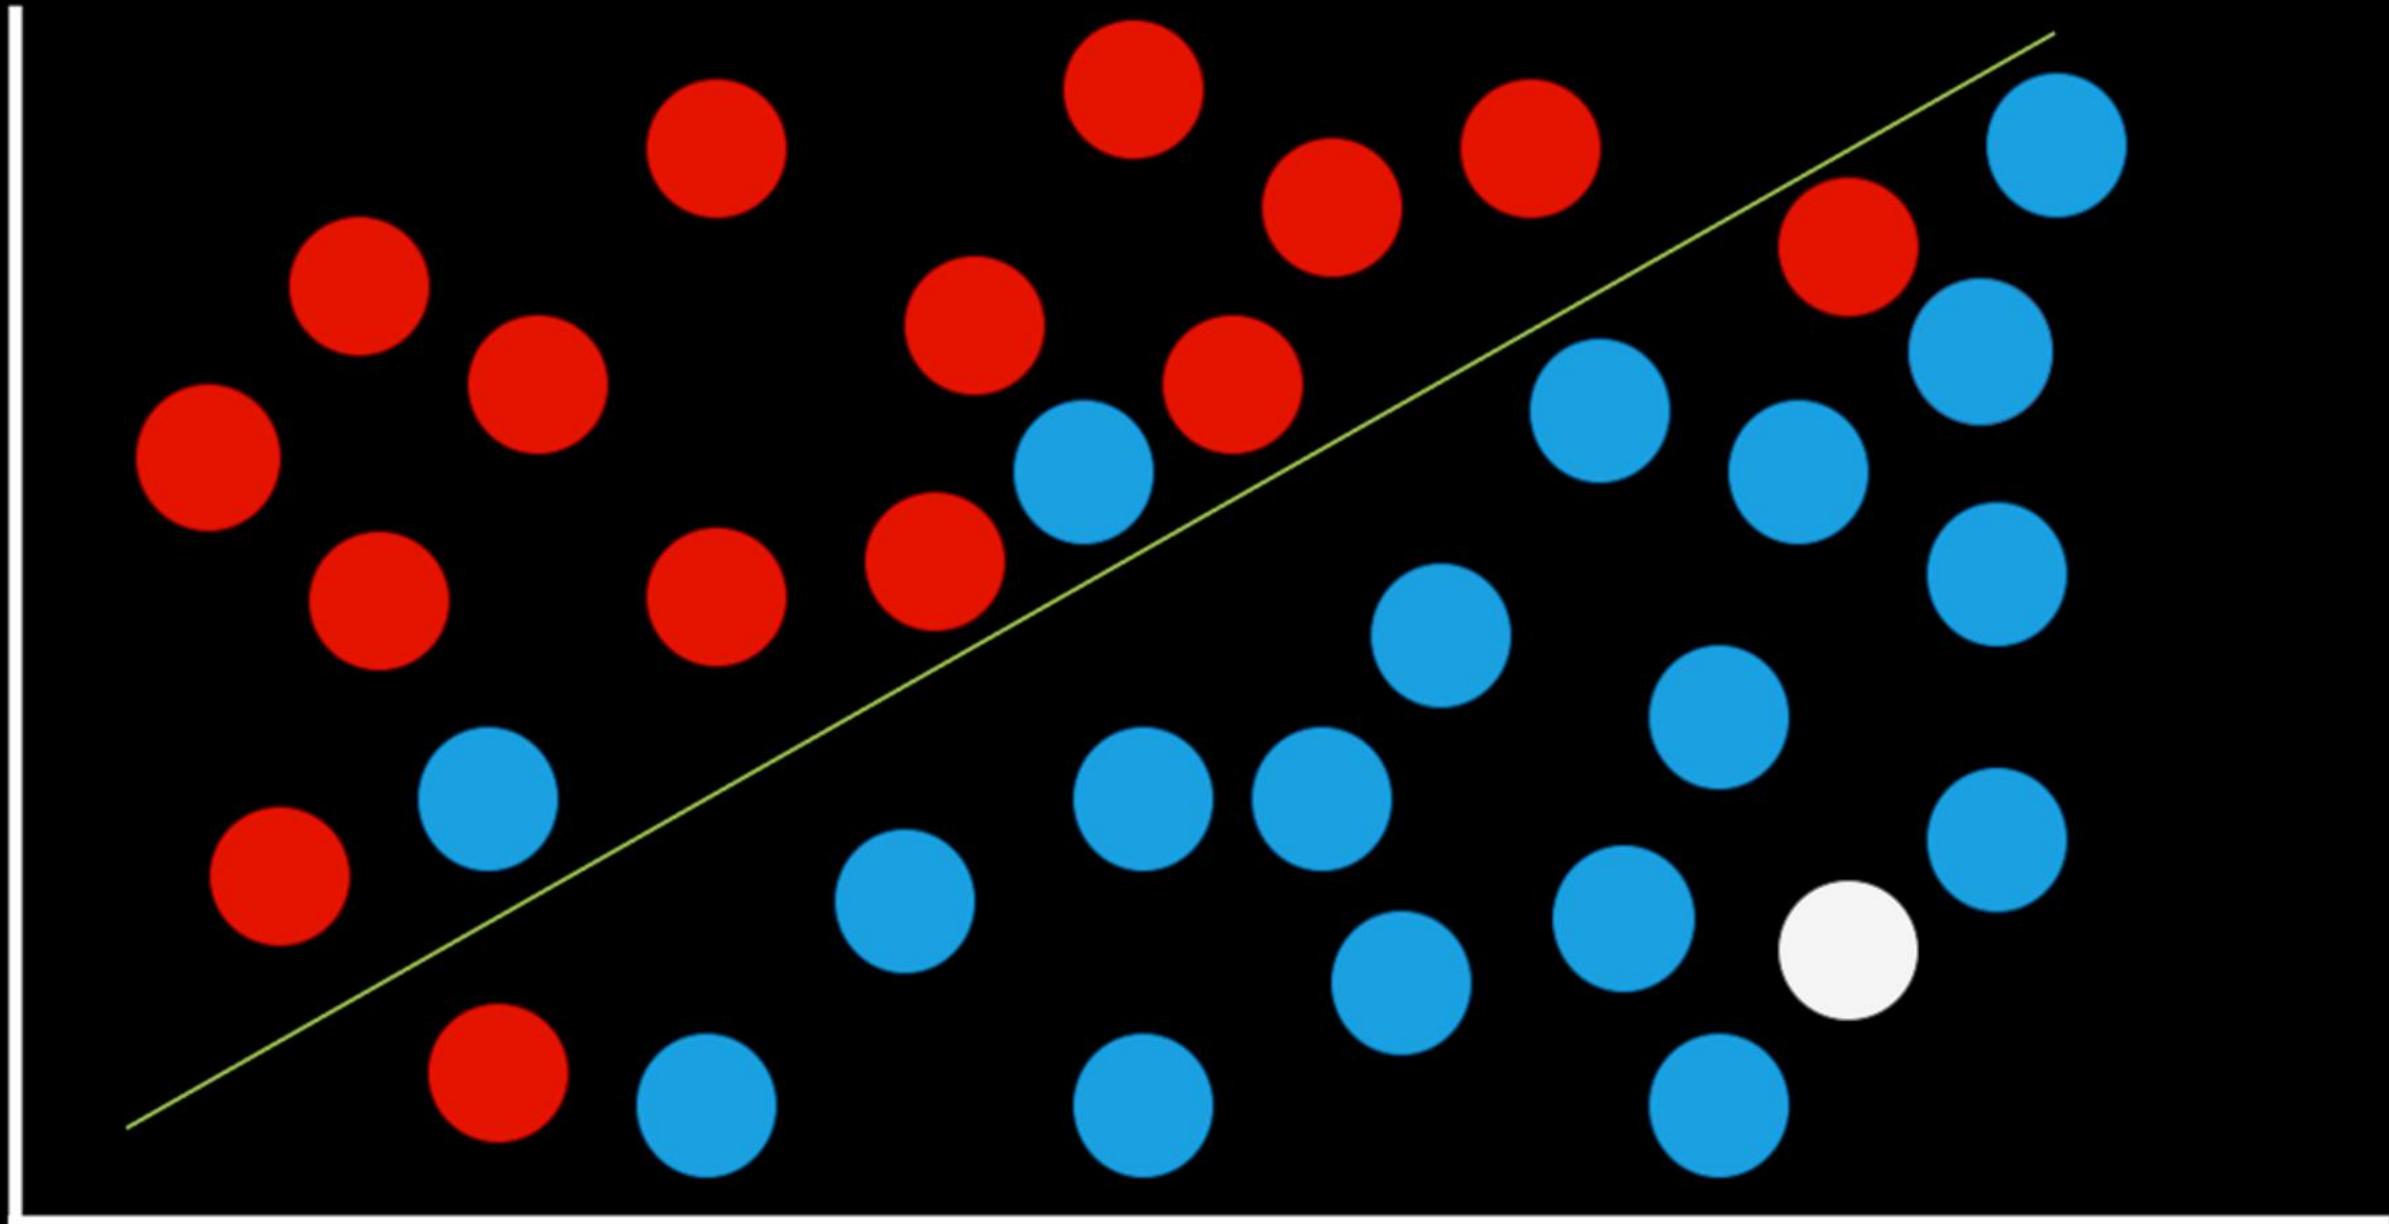
pressure

humidity



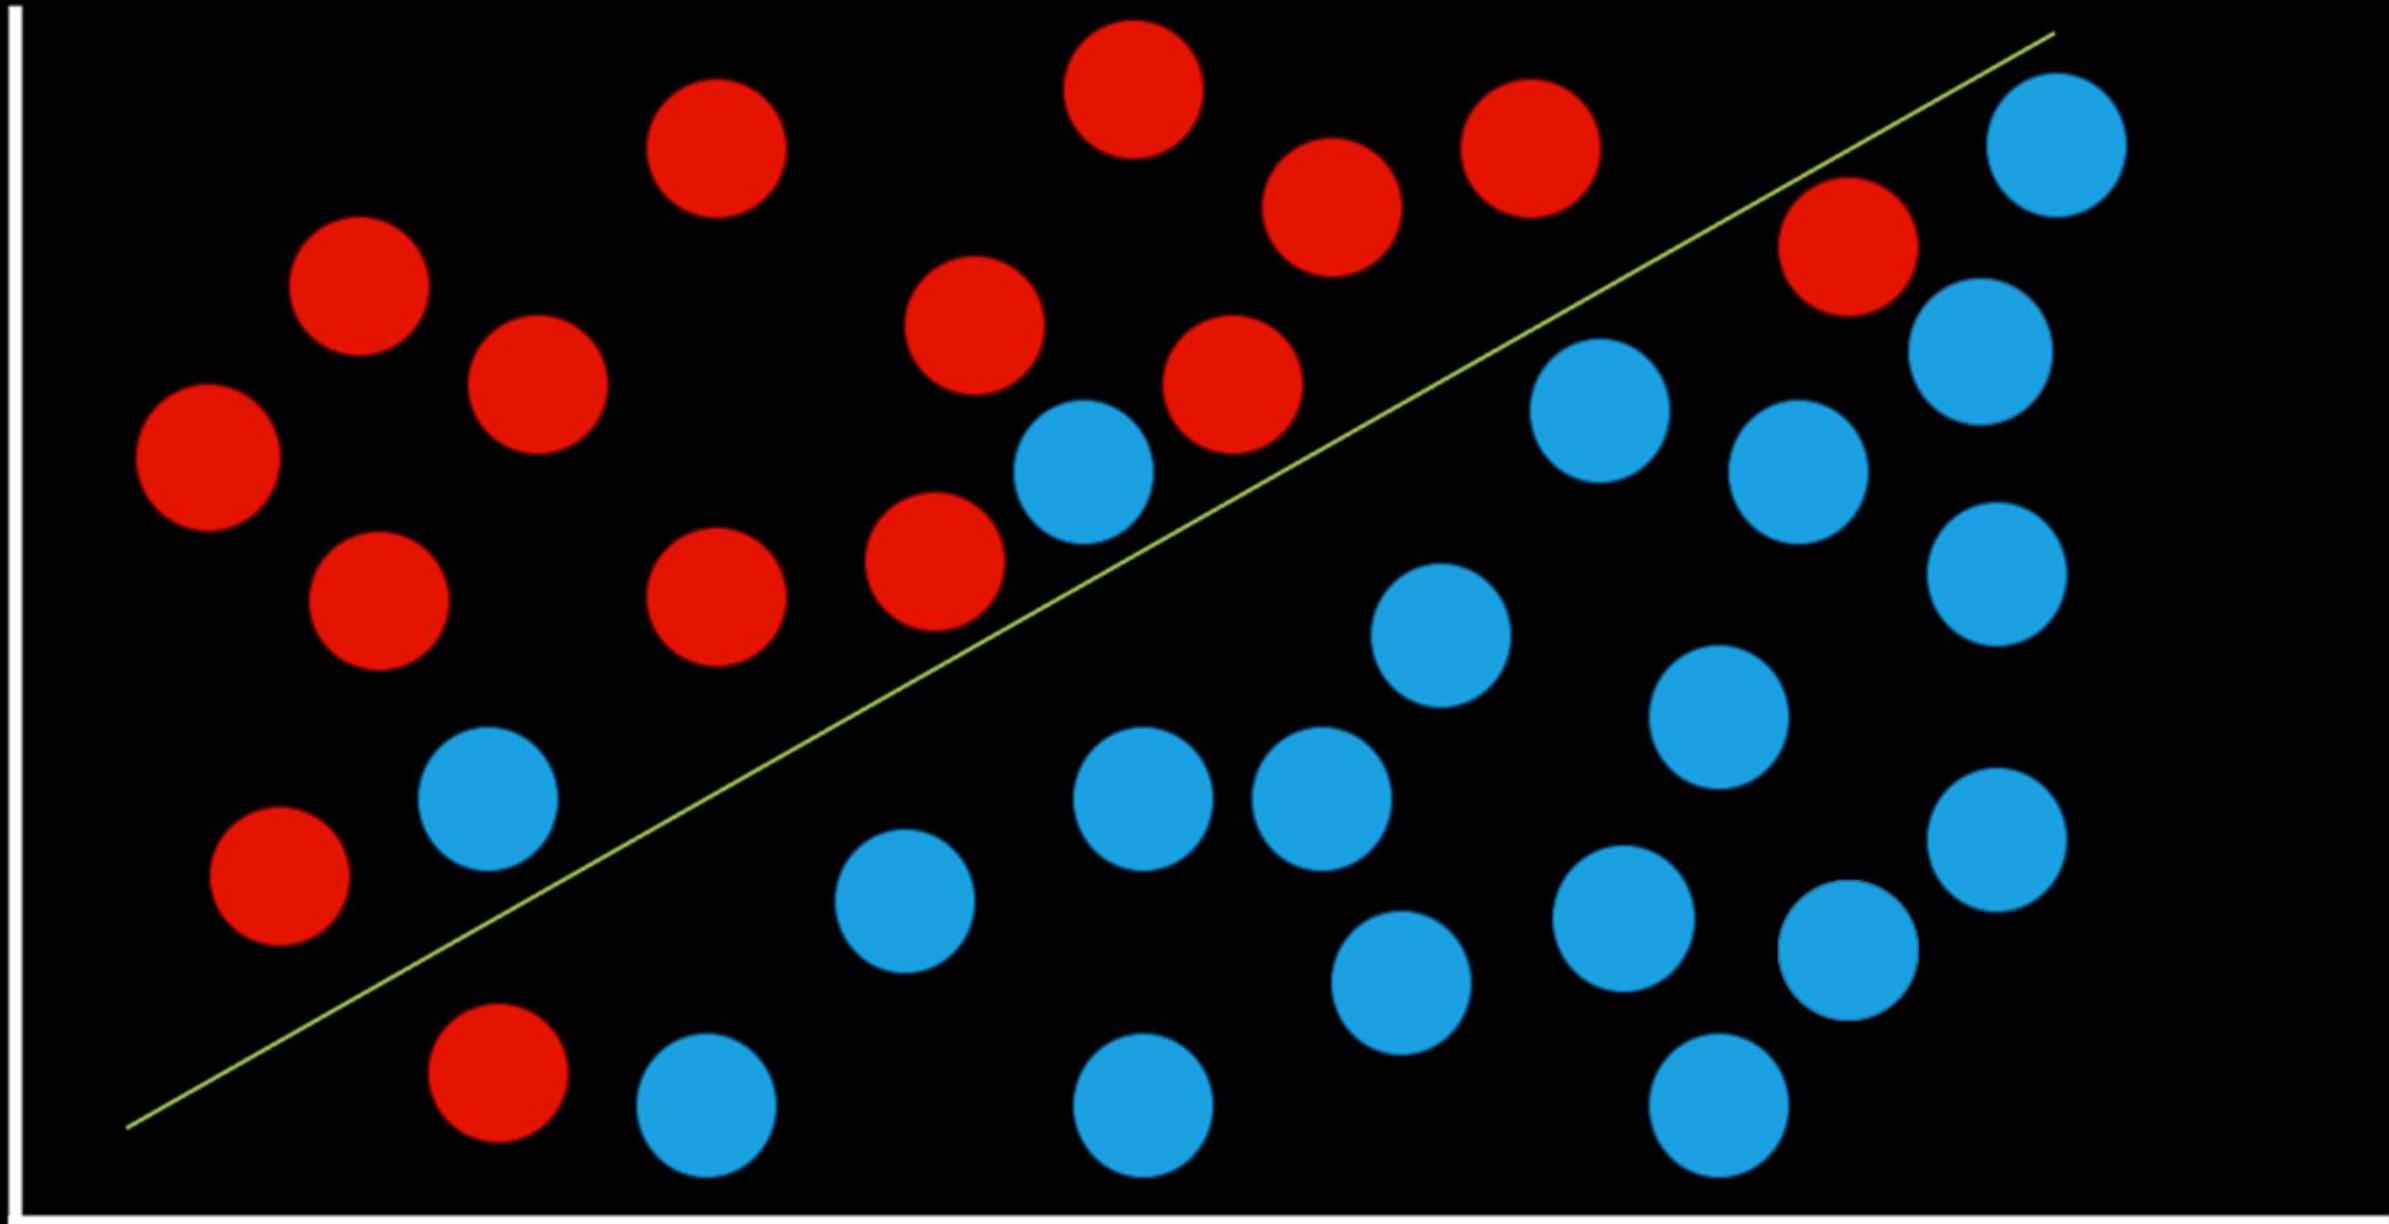
pressure

humidity



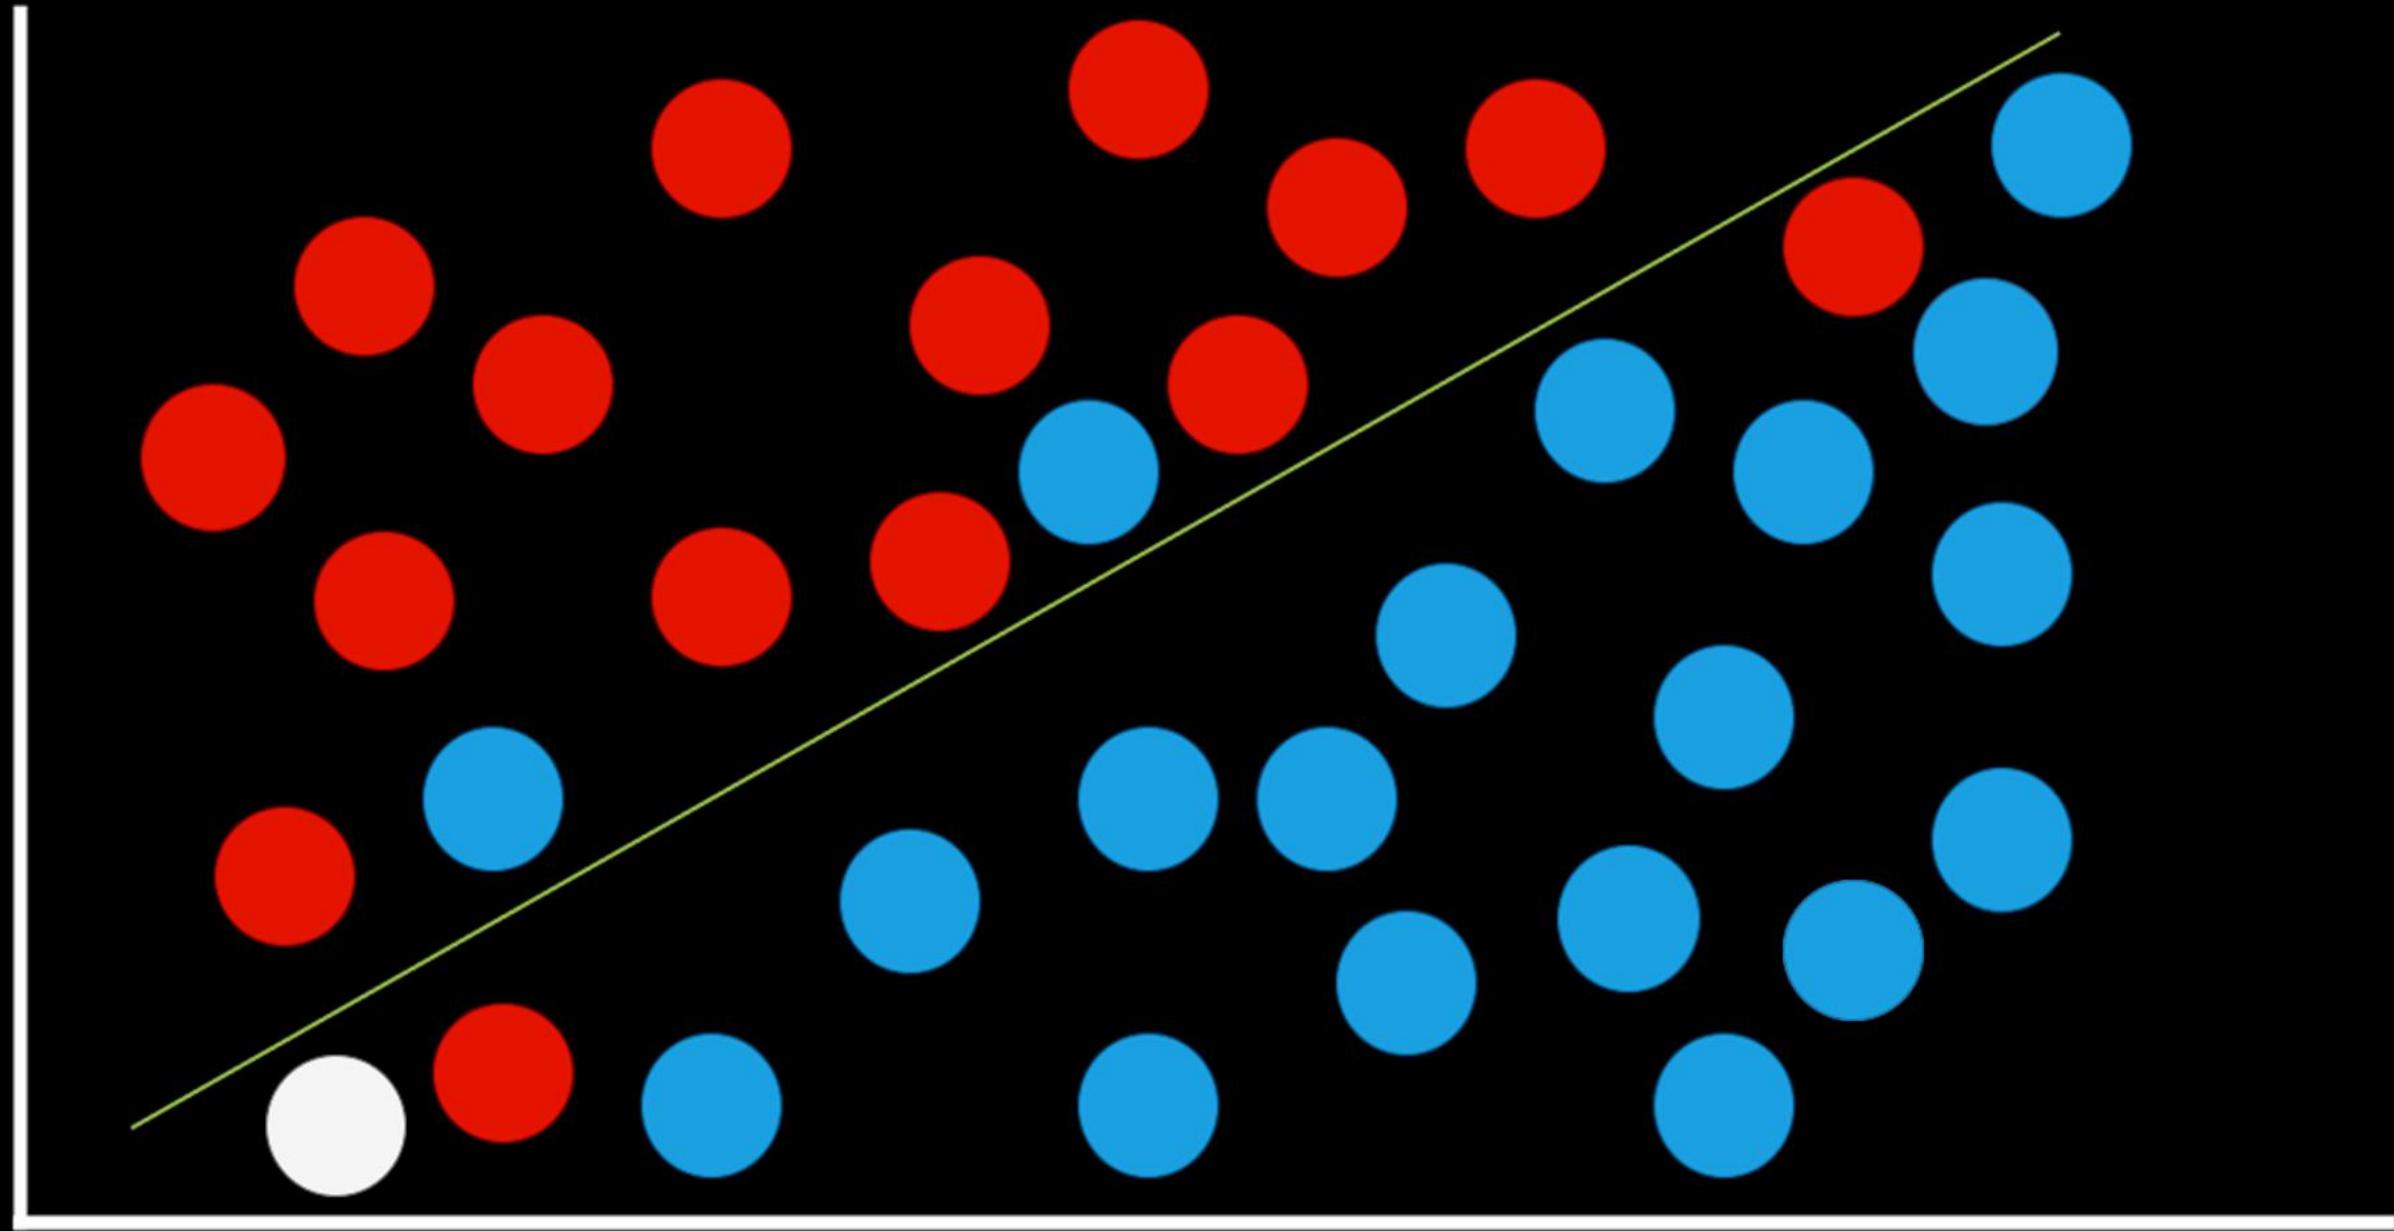
pressure

humidity



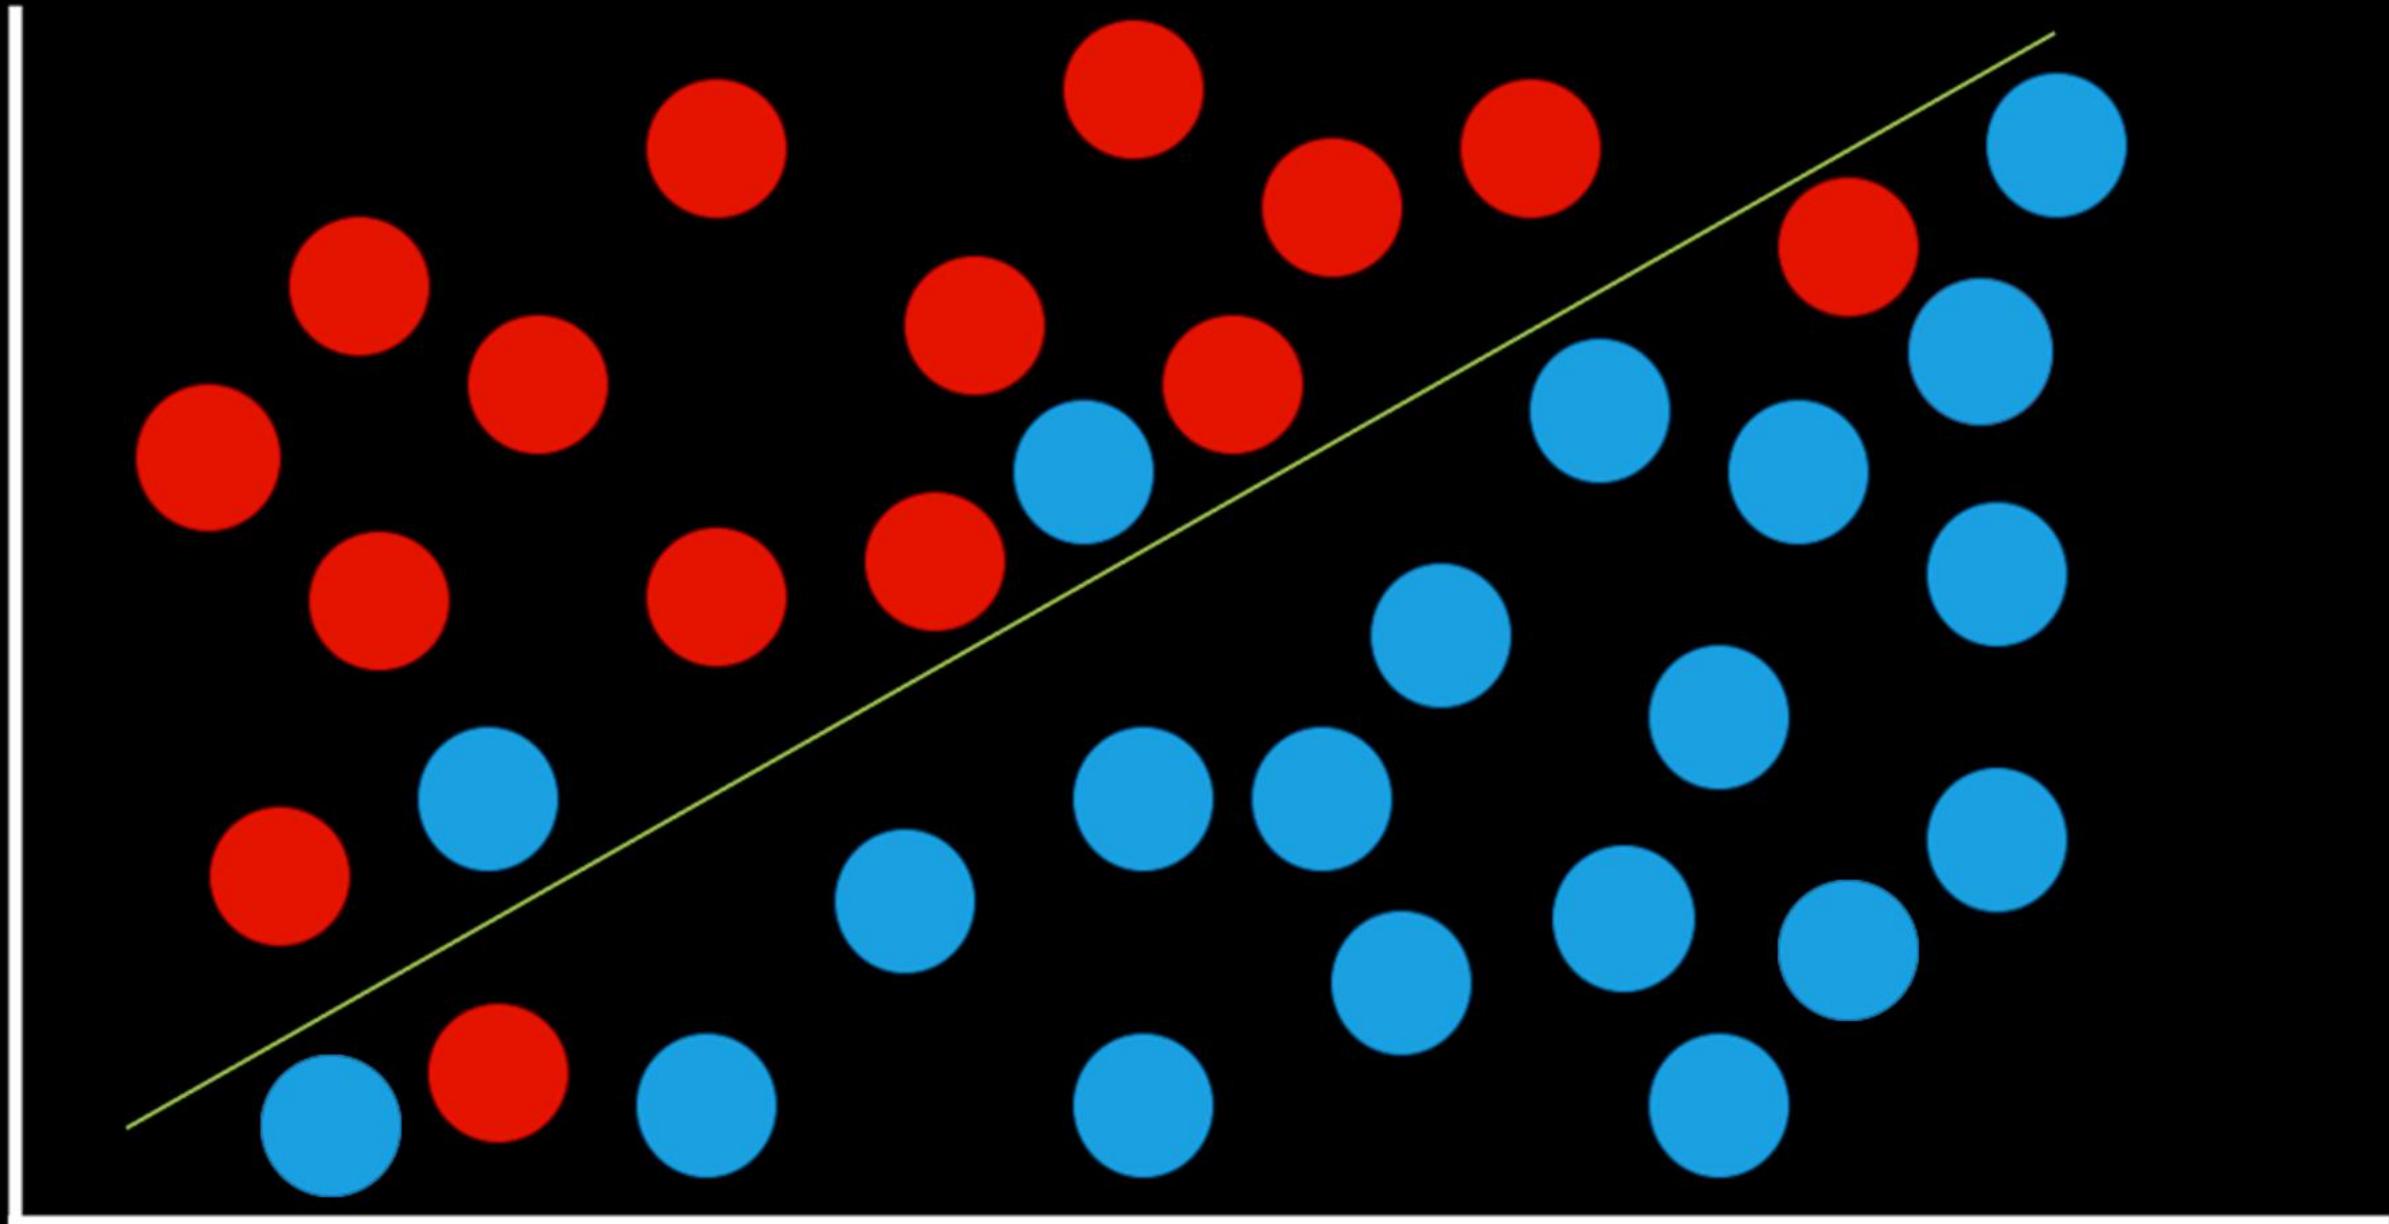
pressure

humidity



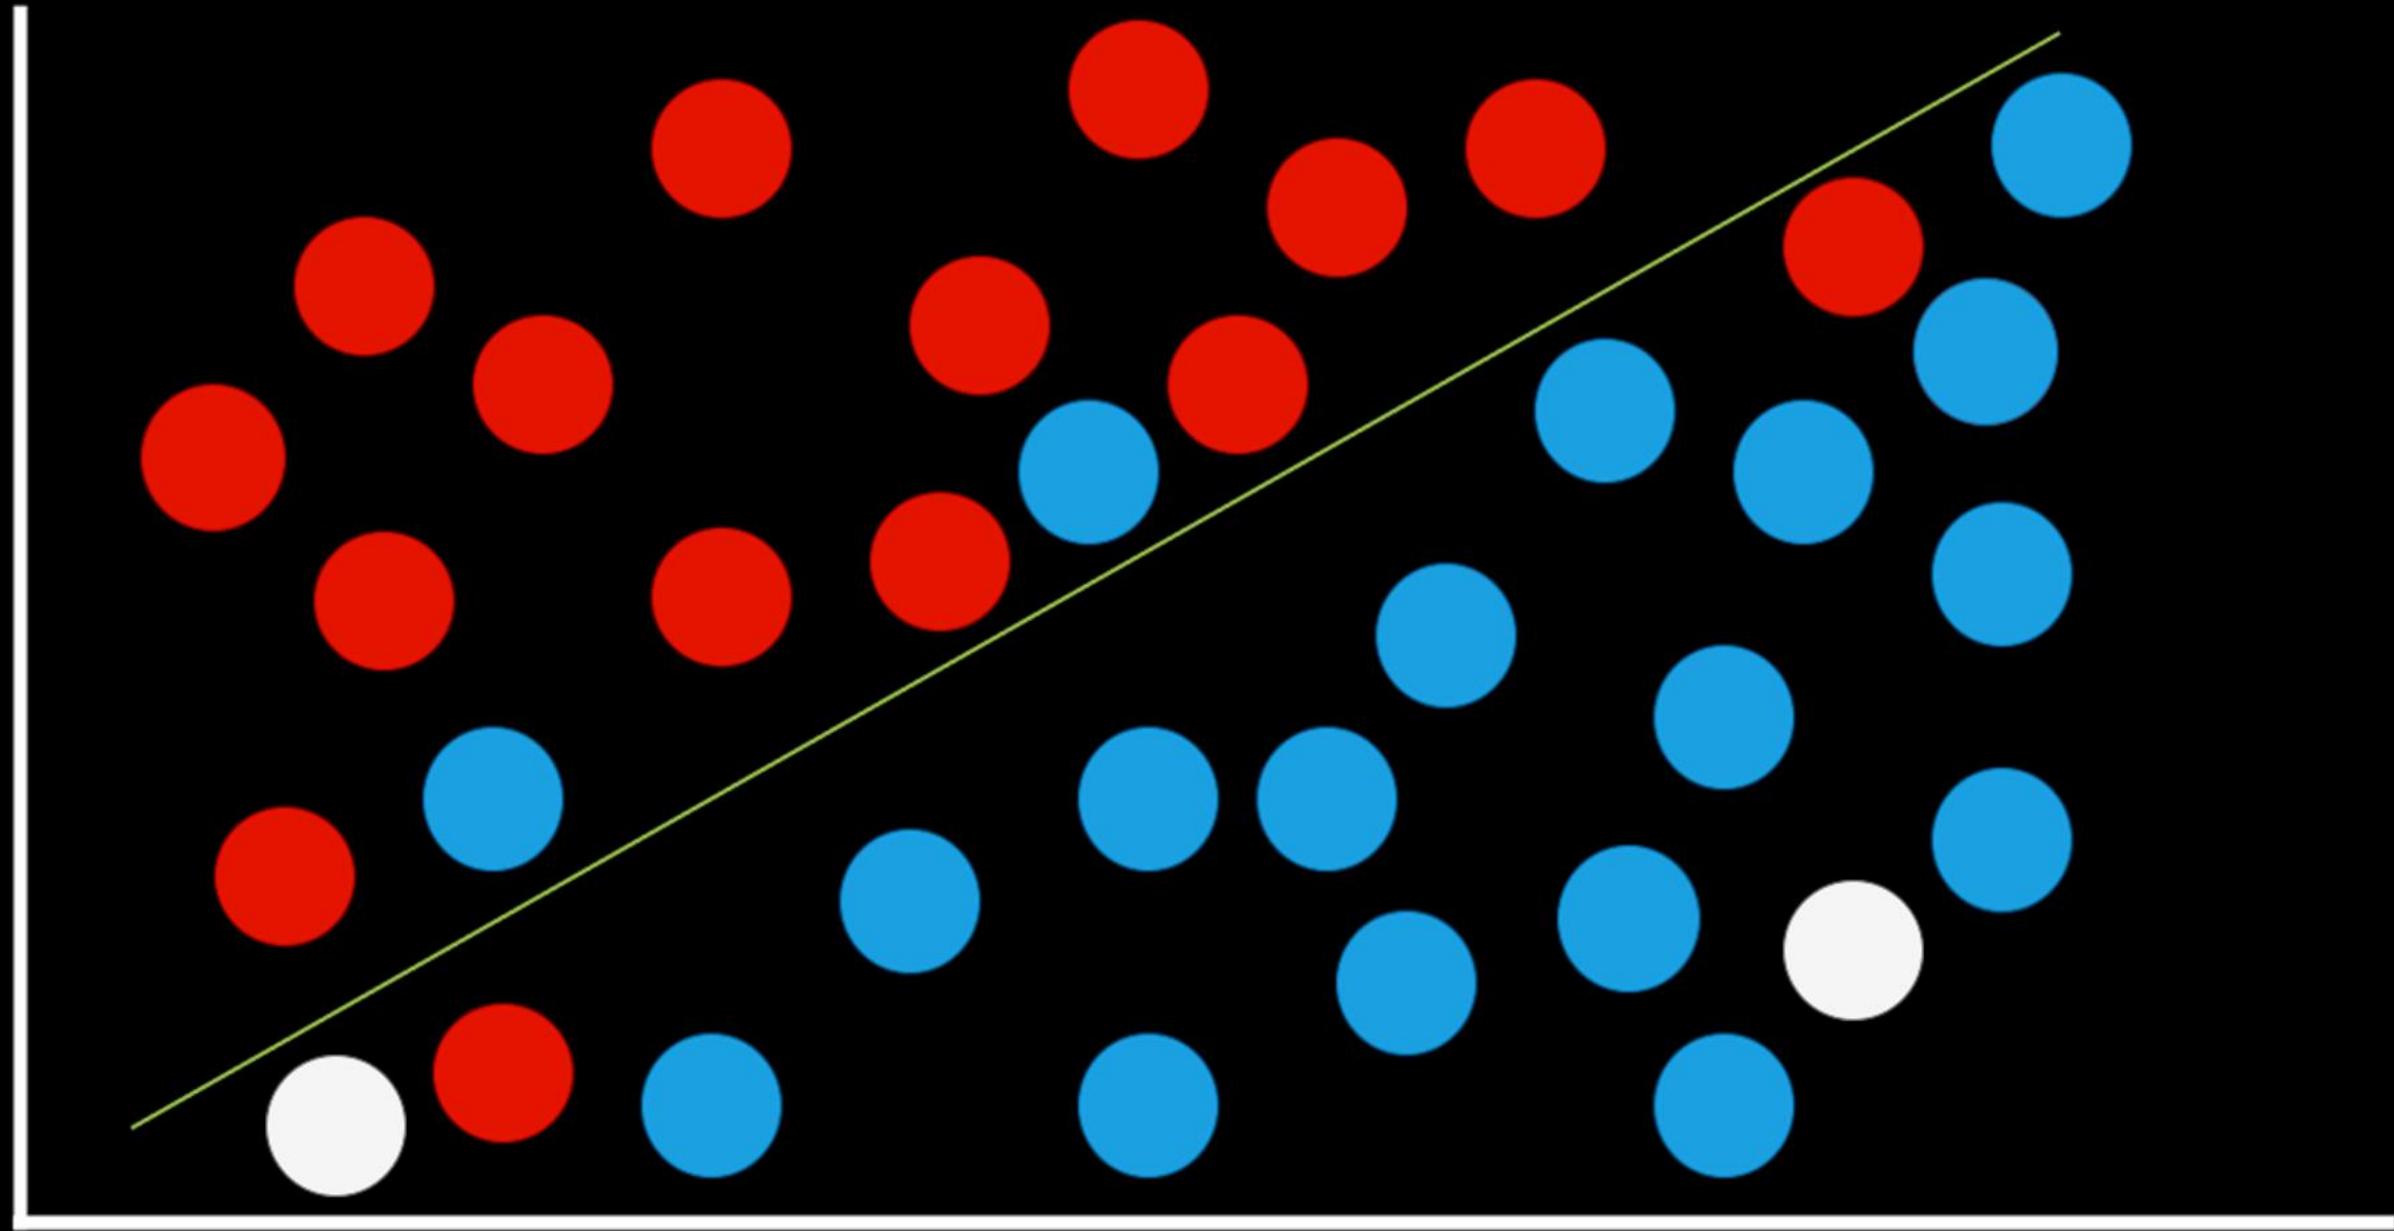
pressure

humidity



pressure

humidity



hard threshold

Output

1

0

$\mathbf{w} \cdot \mathbf{x}$

# soft threshold

output

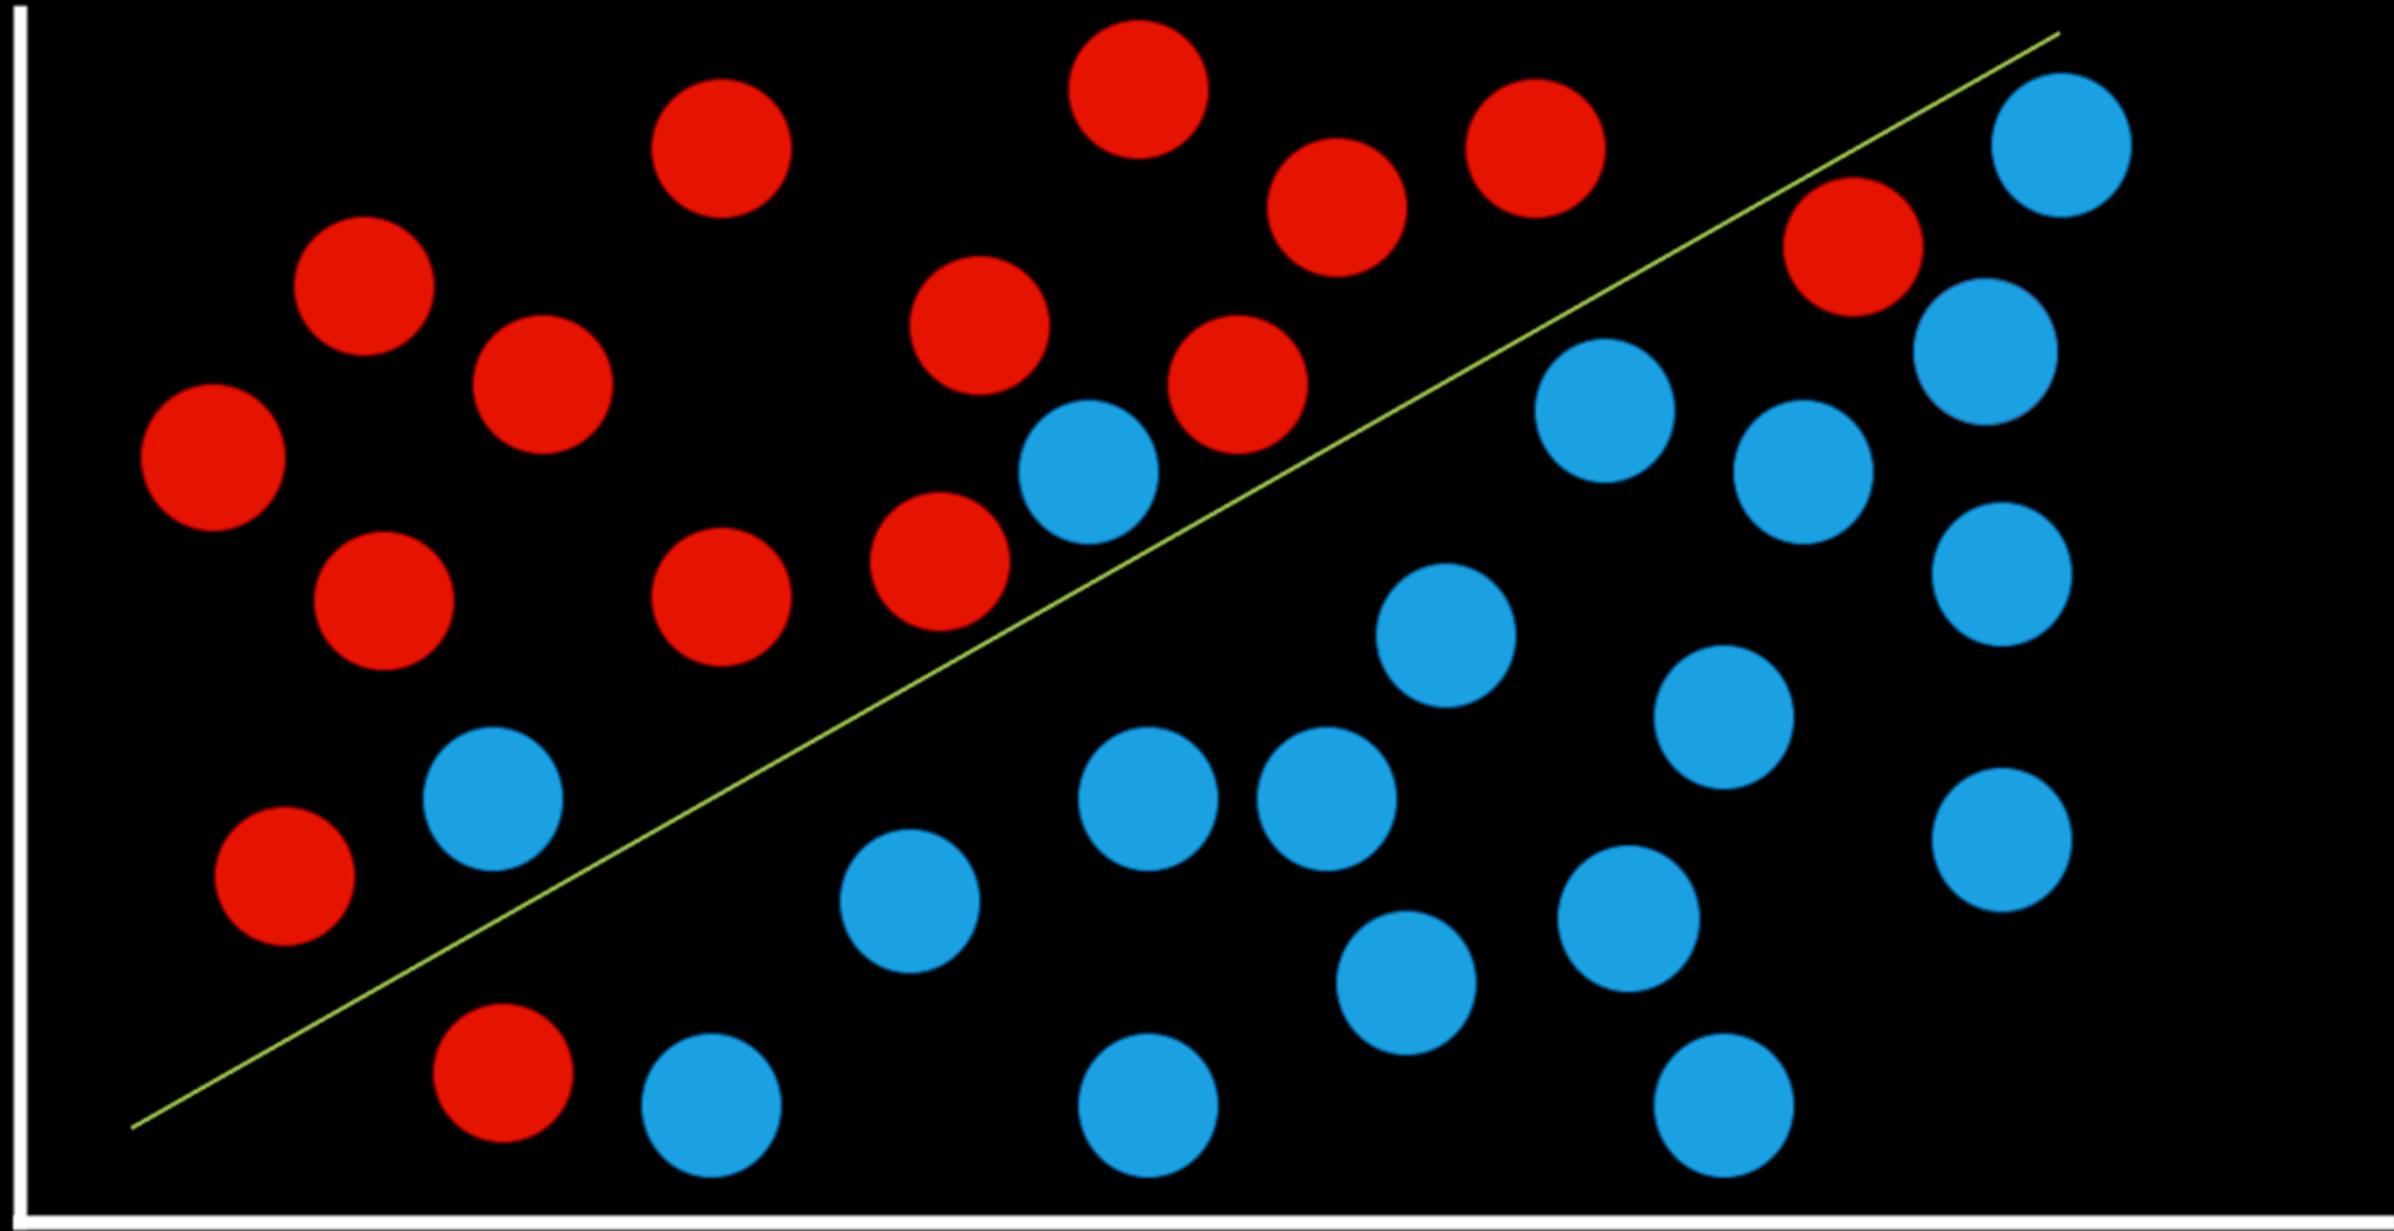
1

0

$w \cdot x$

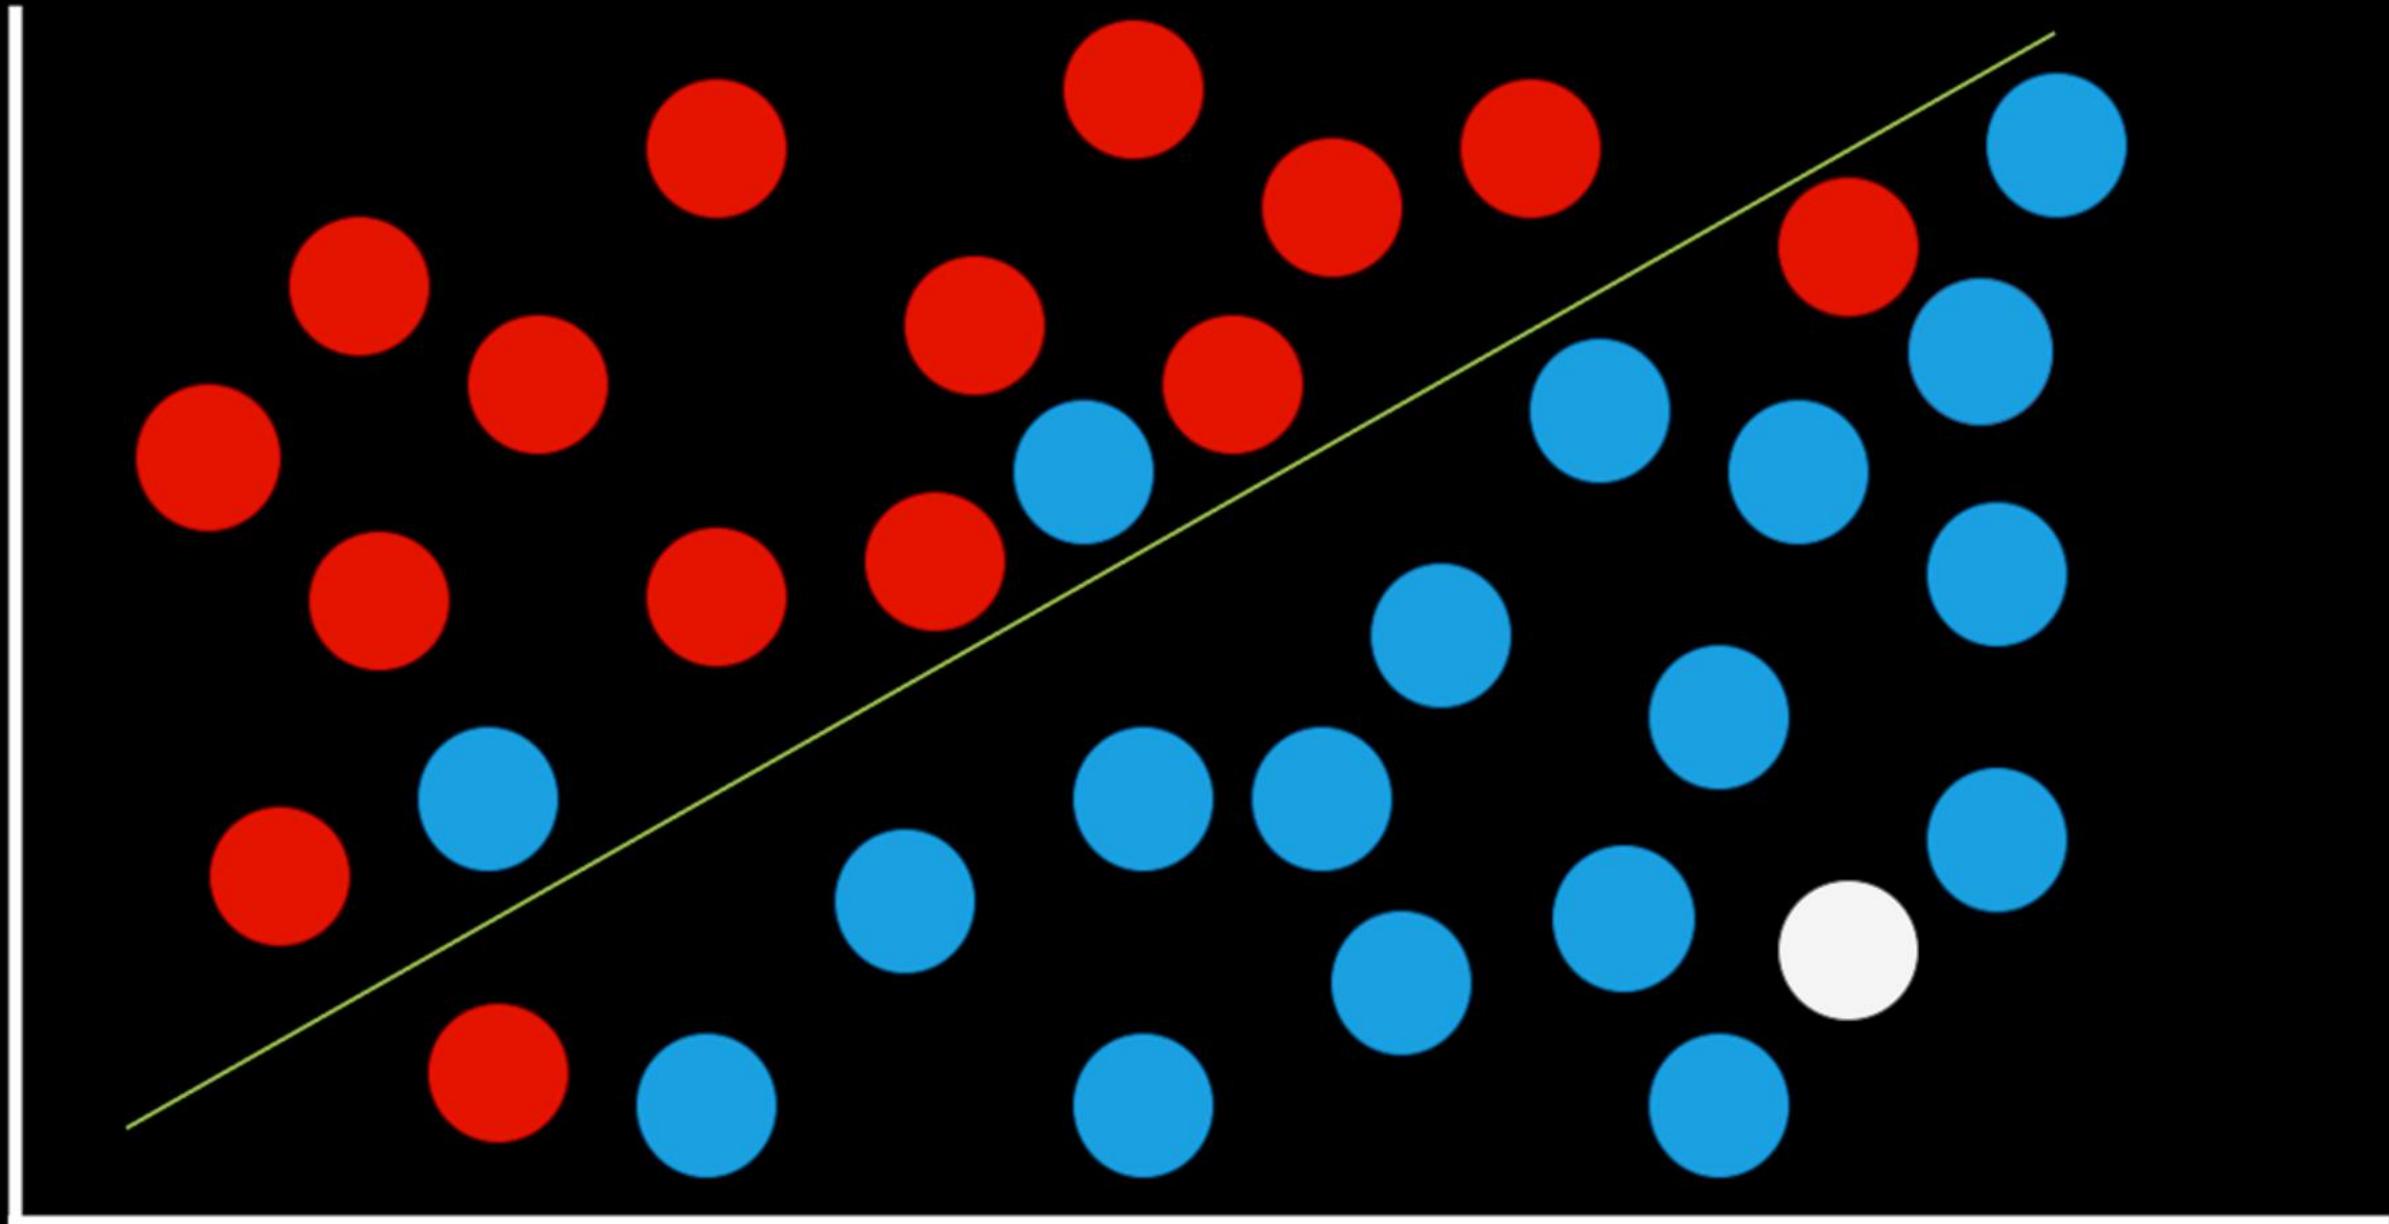
pressure

humidity



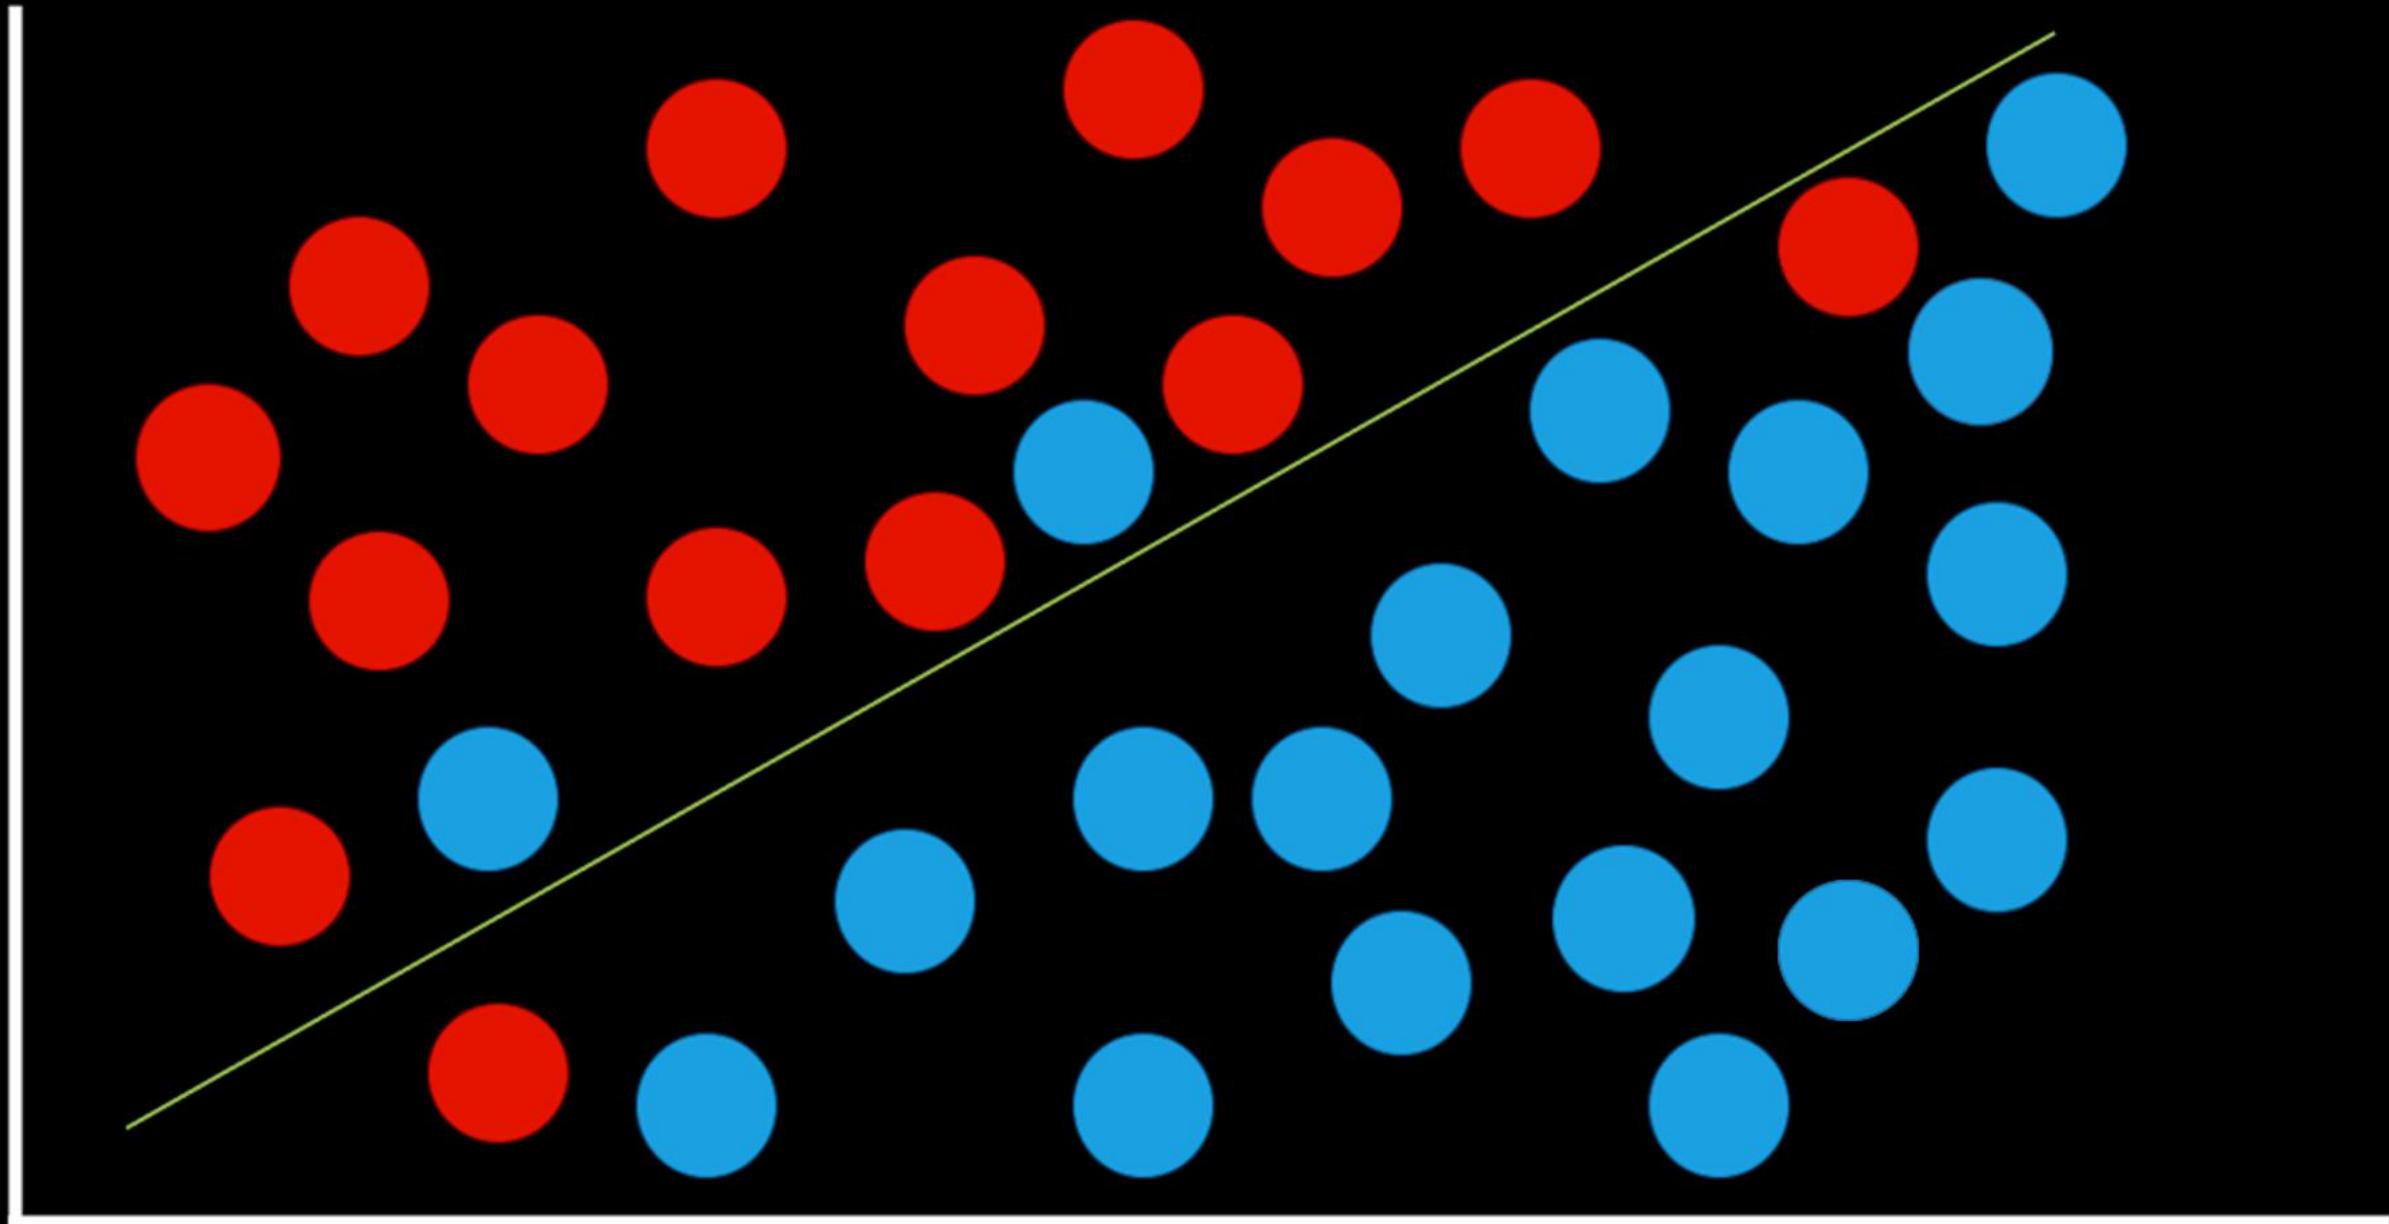
pressure

humidity



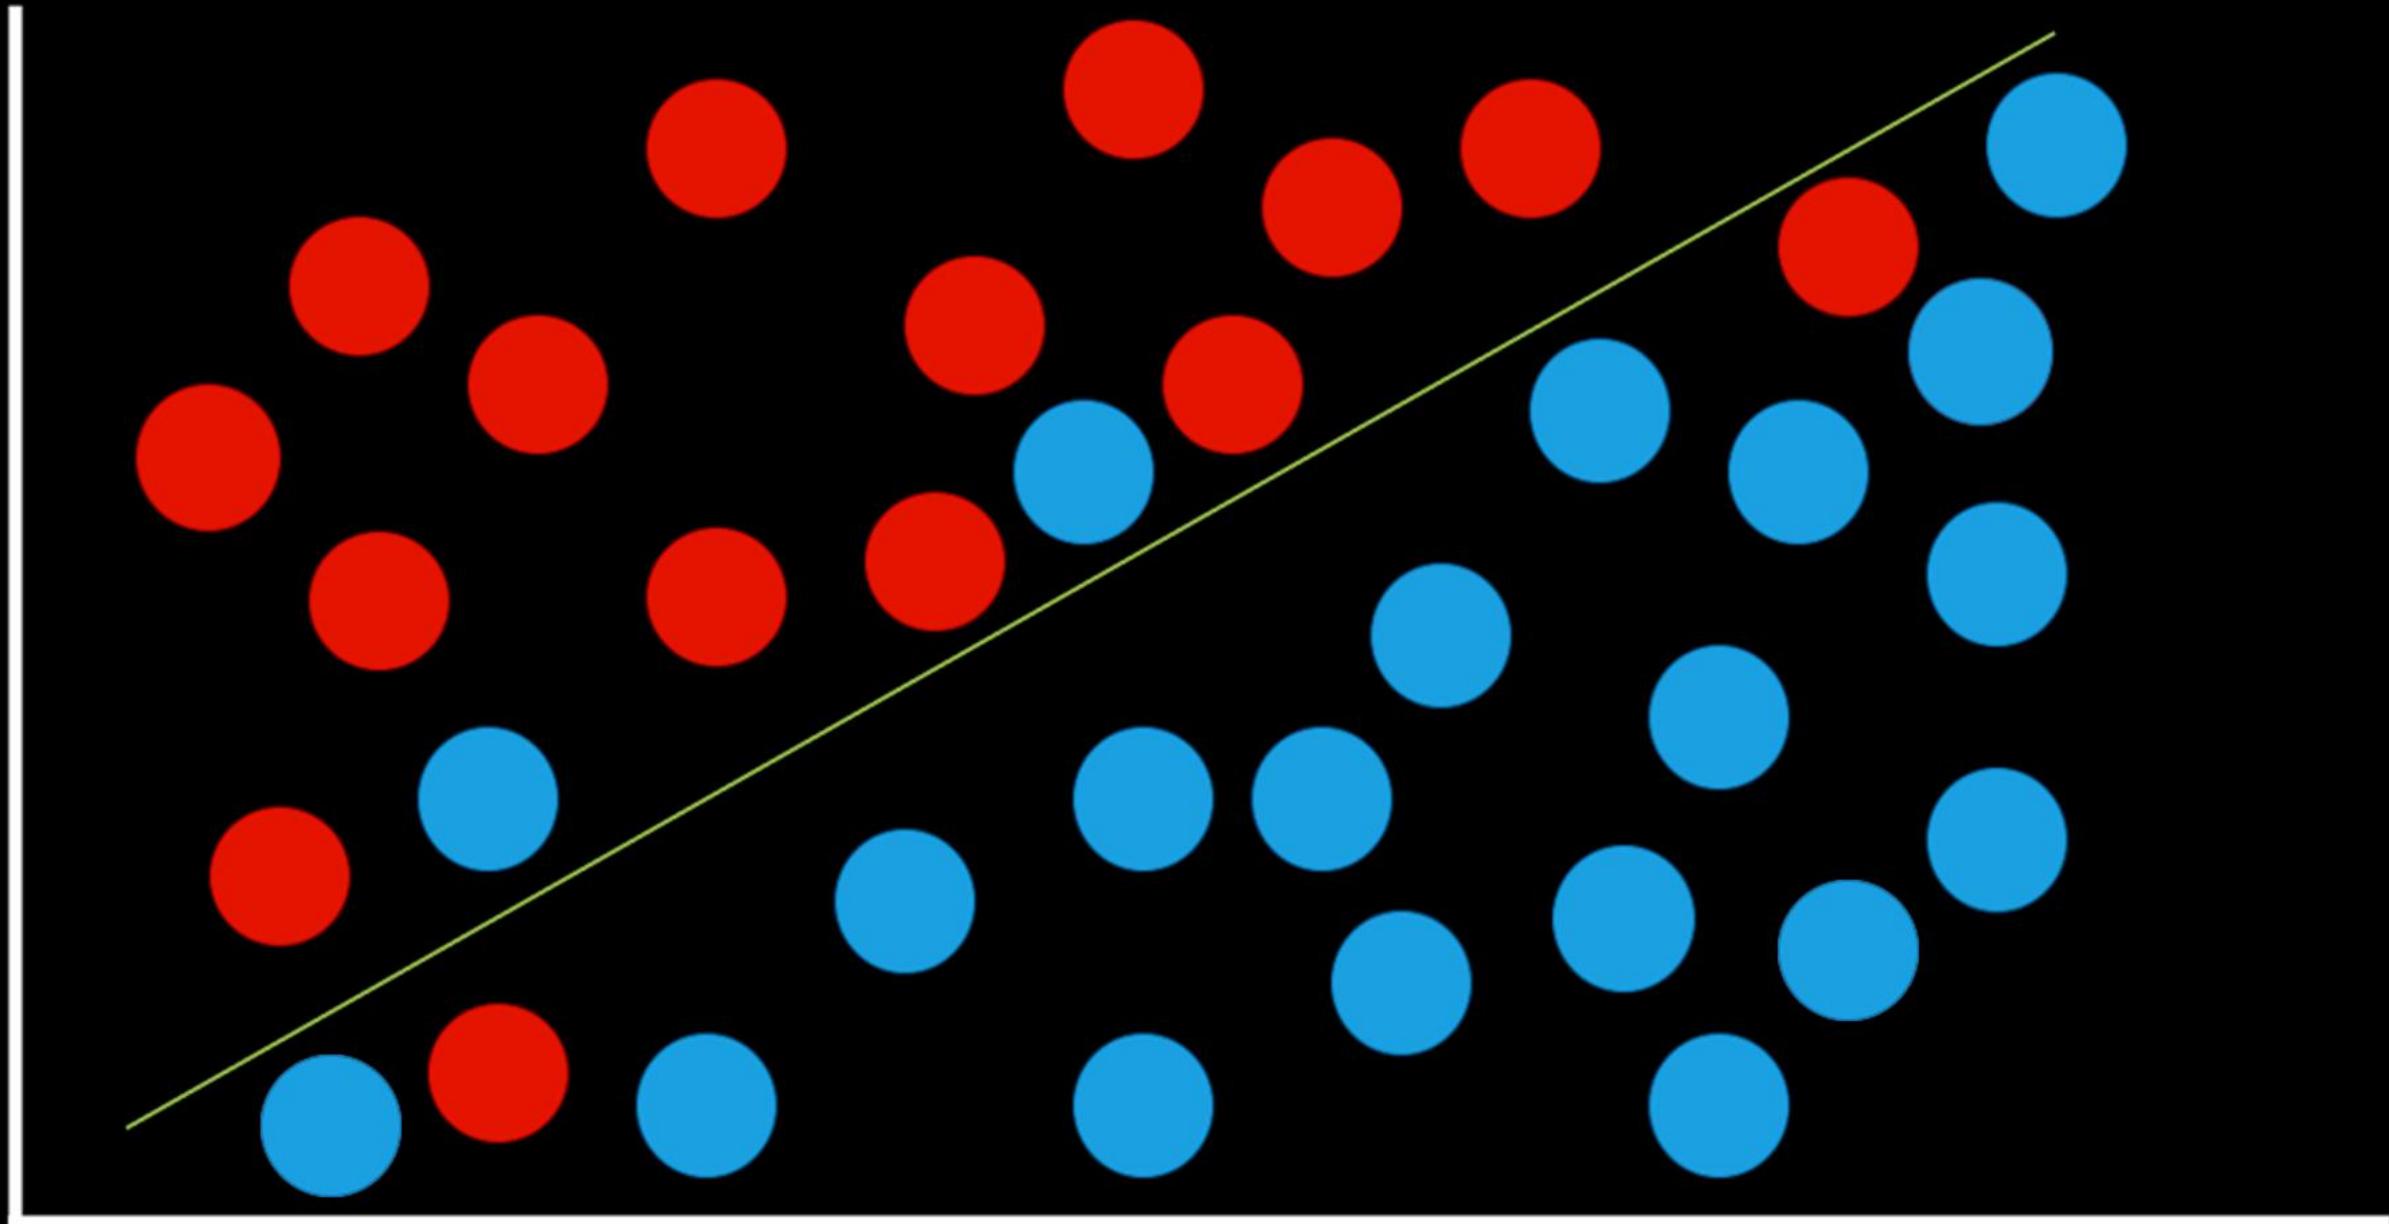
pressure

humidity



pressure

humidity



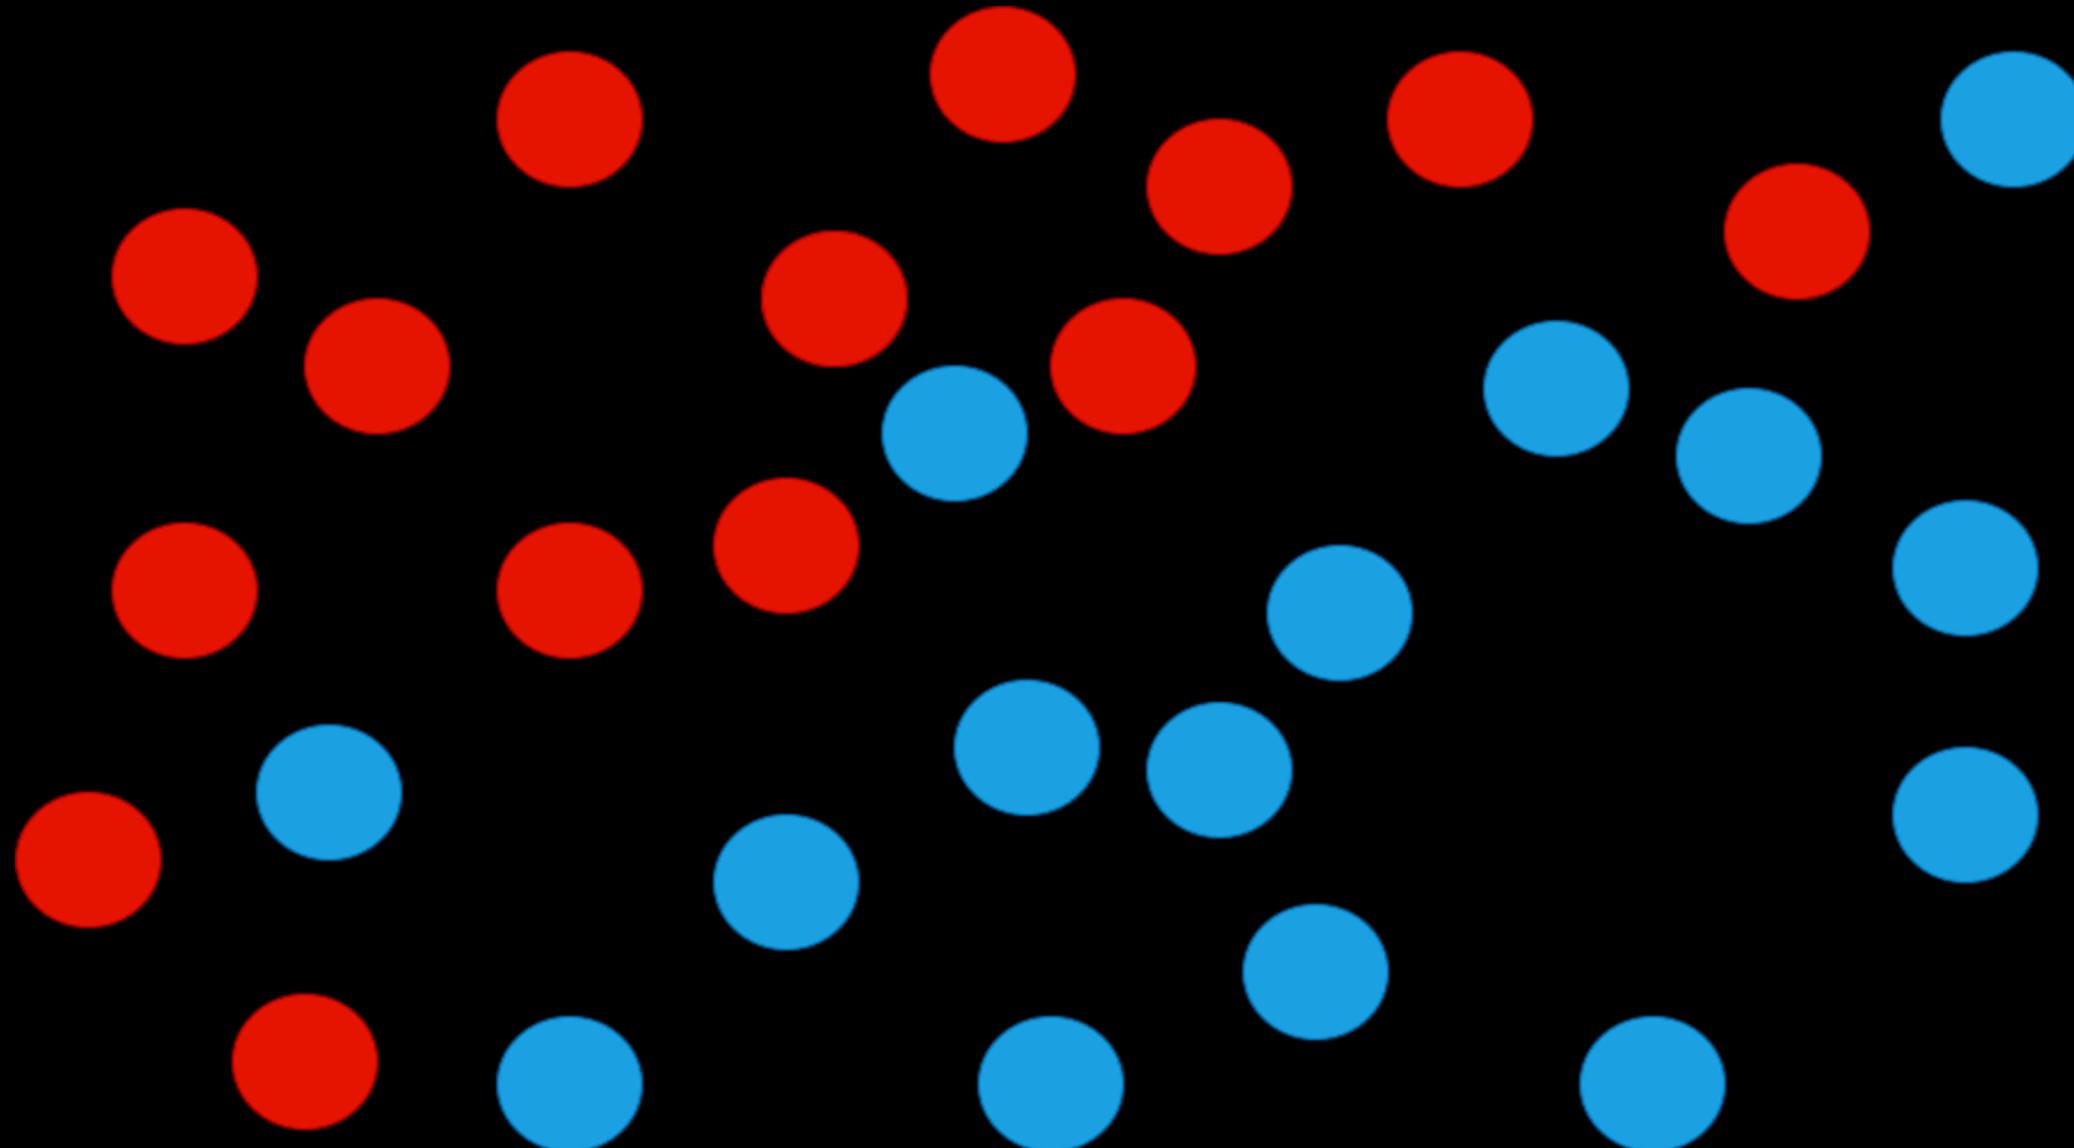
# Evaluating Hypotheses

# loss function

function that expresses how poorly our hypothesis performs

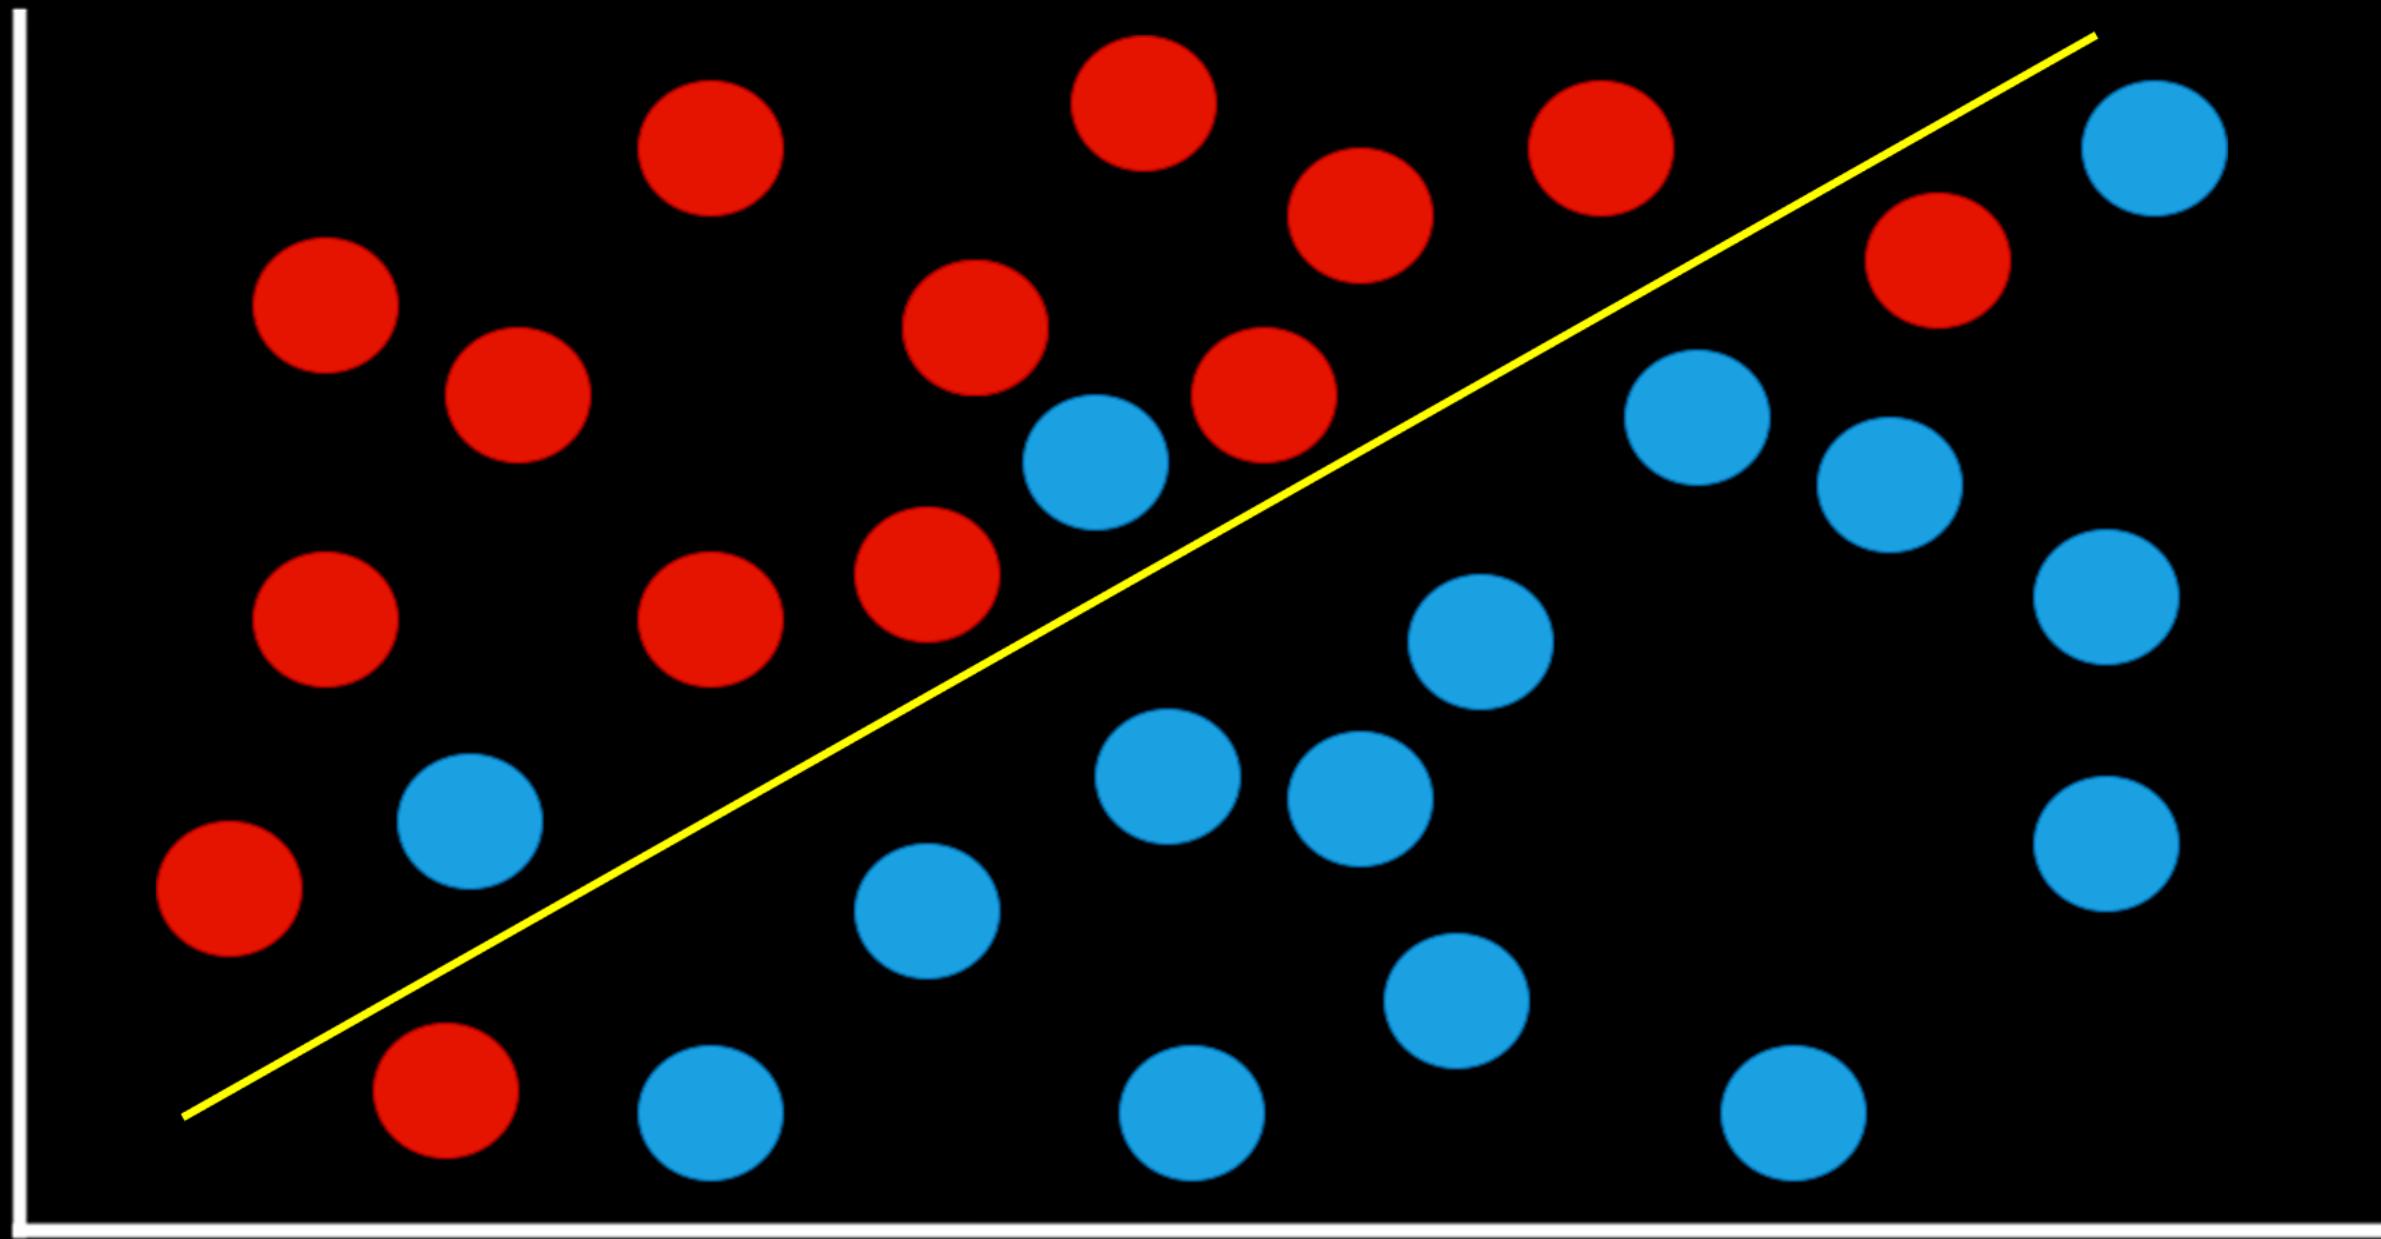
pressure

humidity



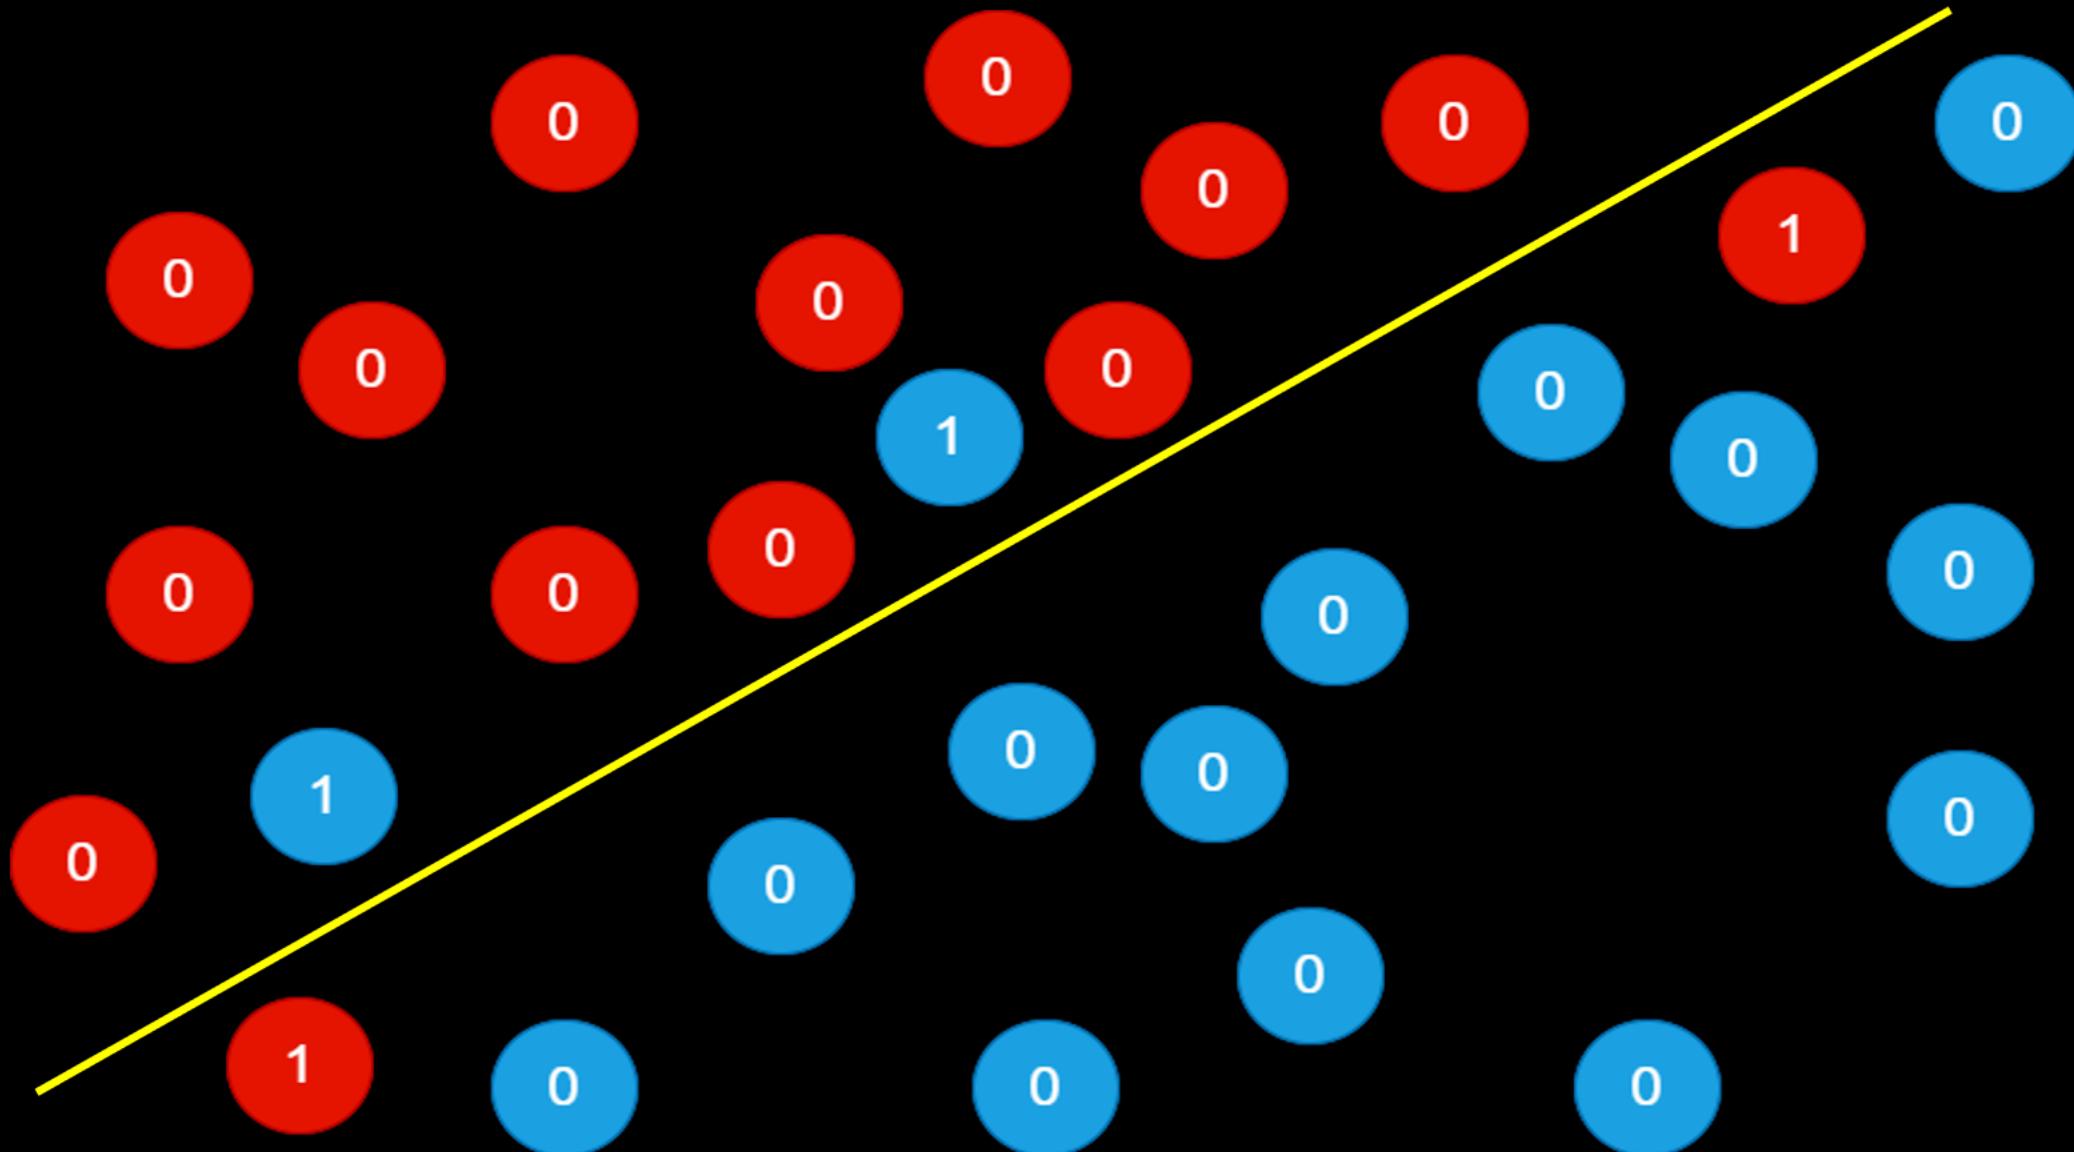
pressure

humidity

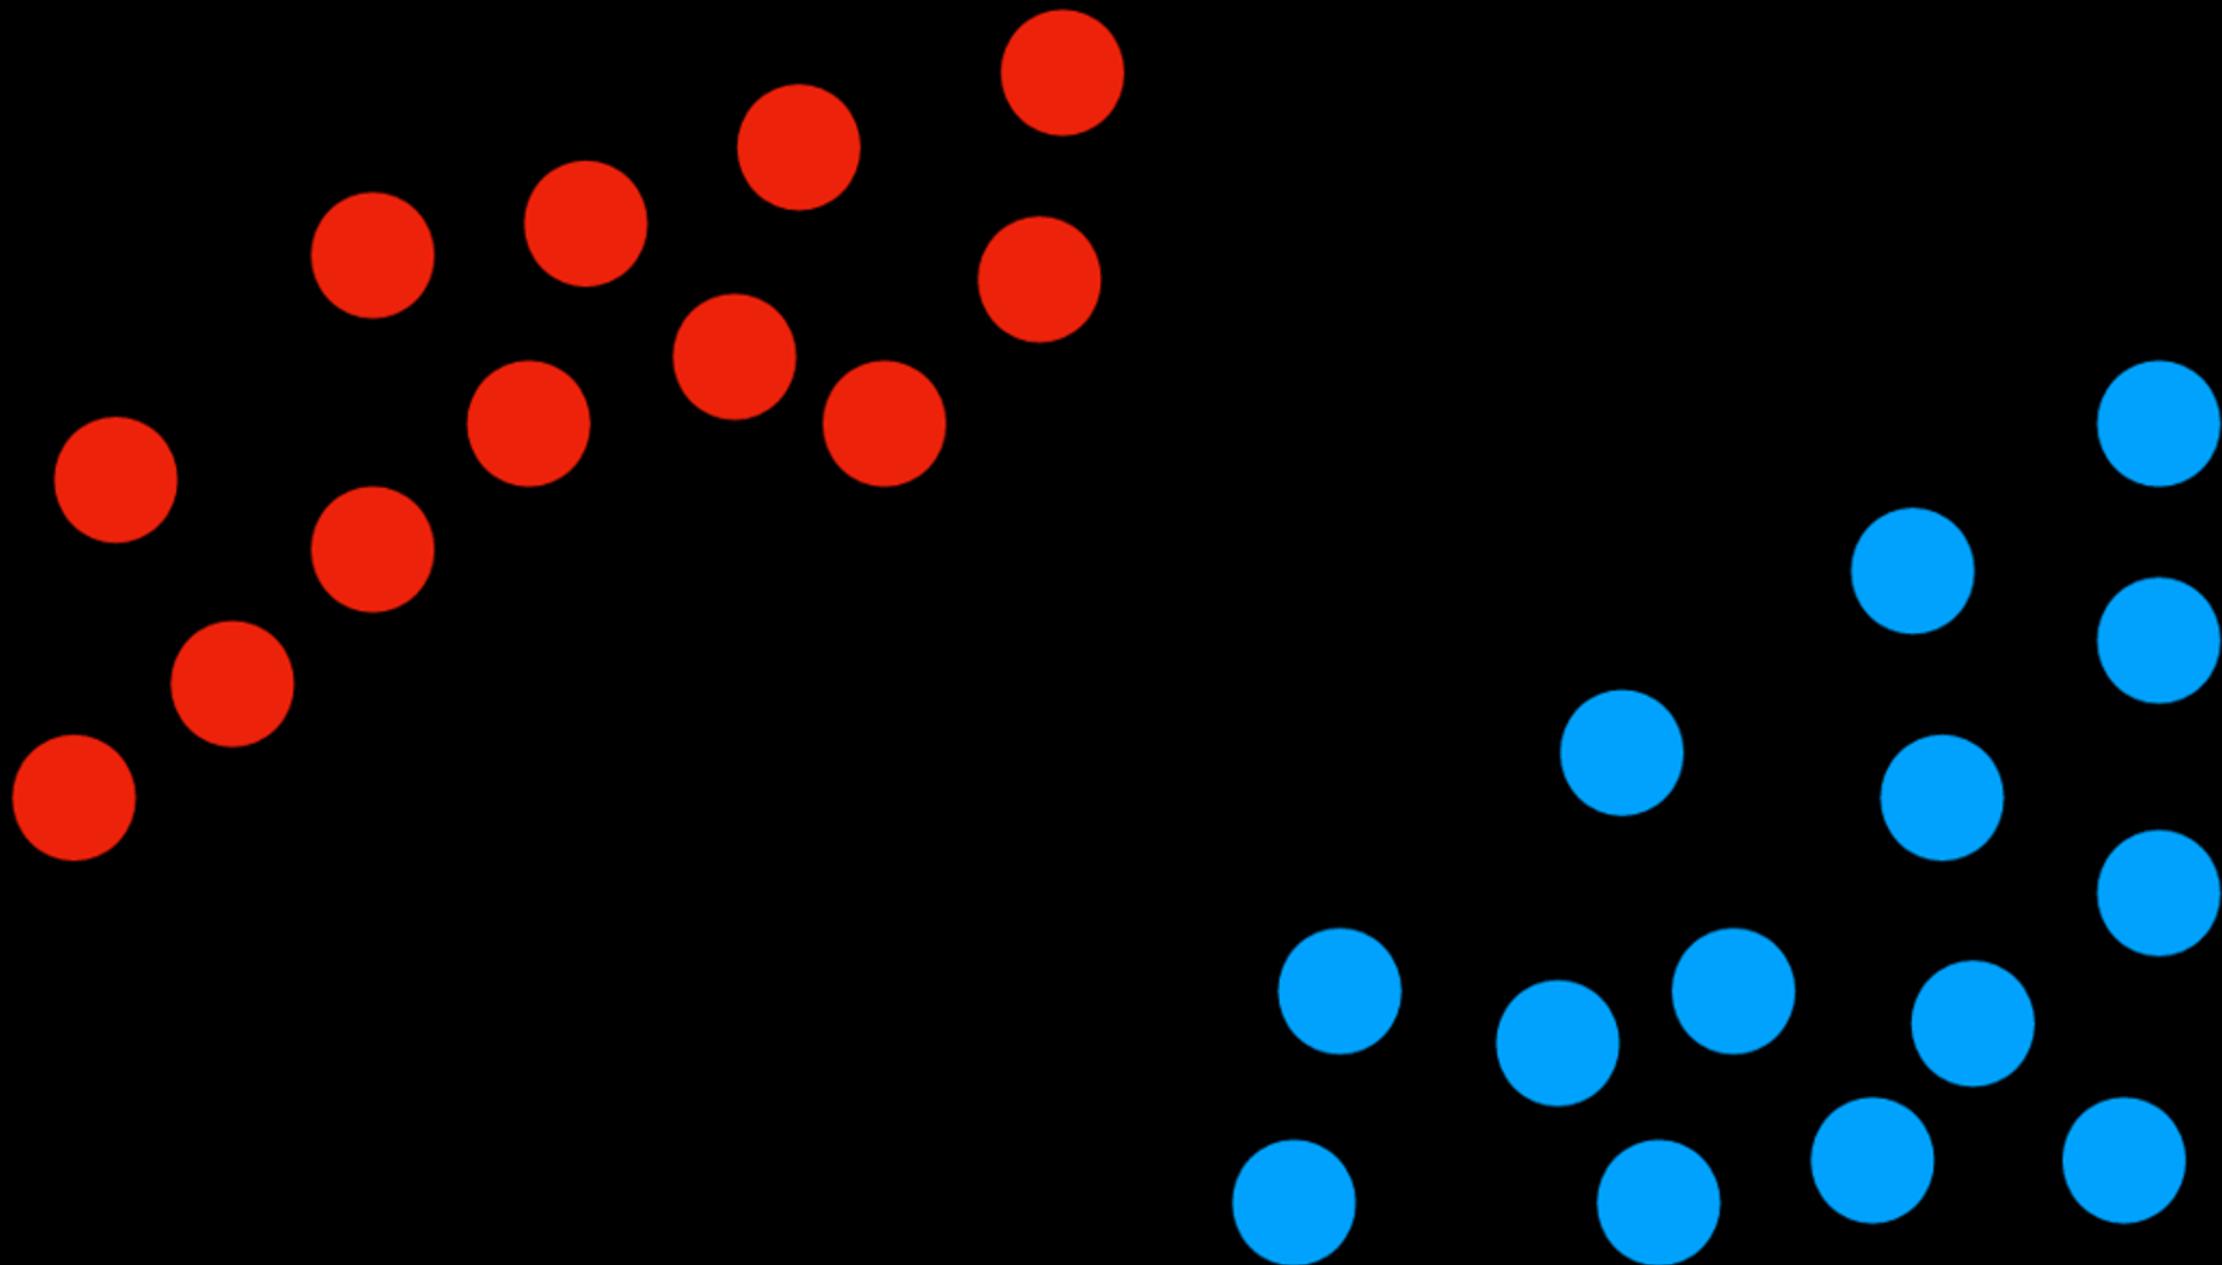


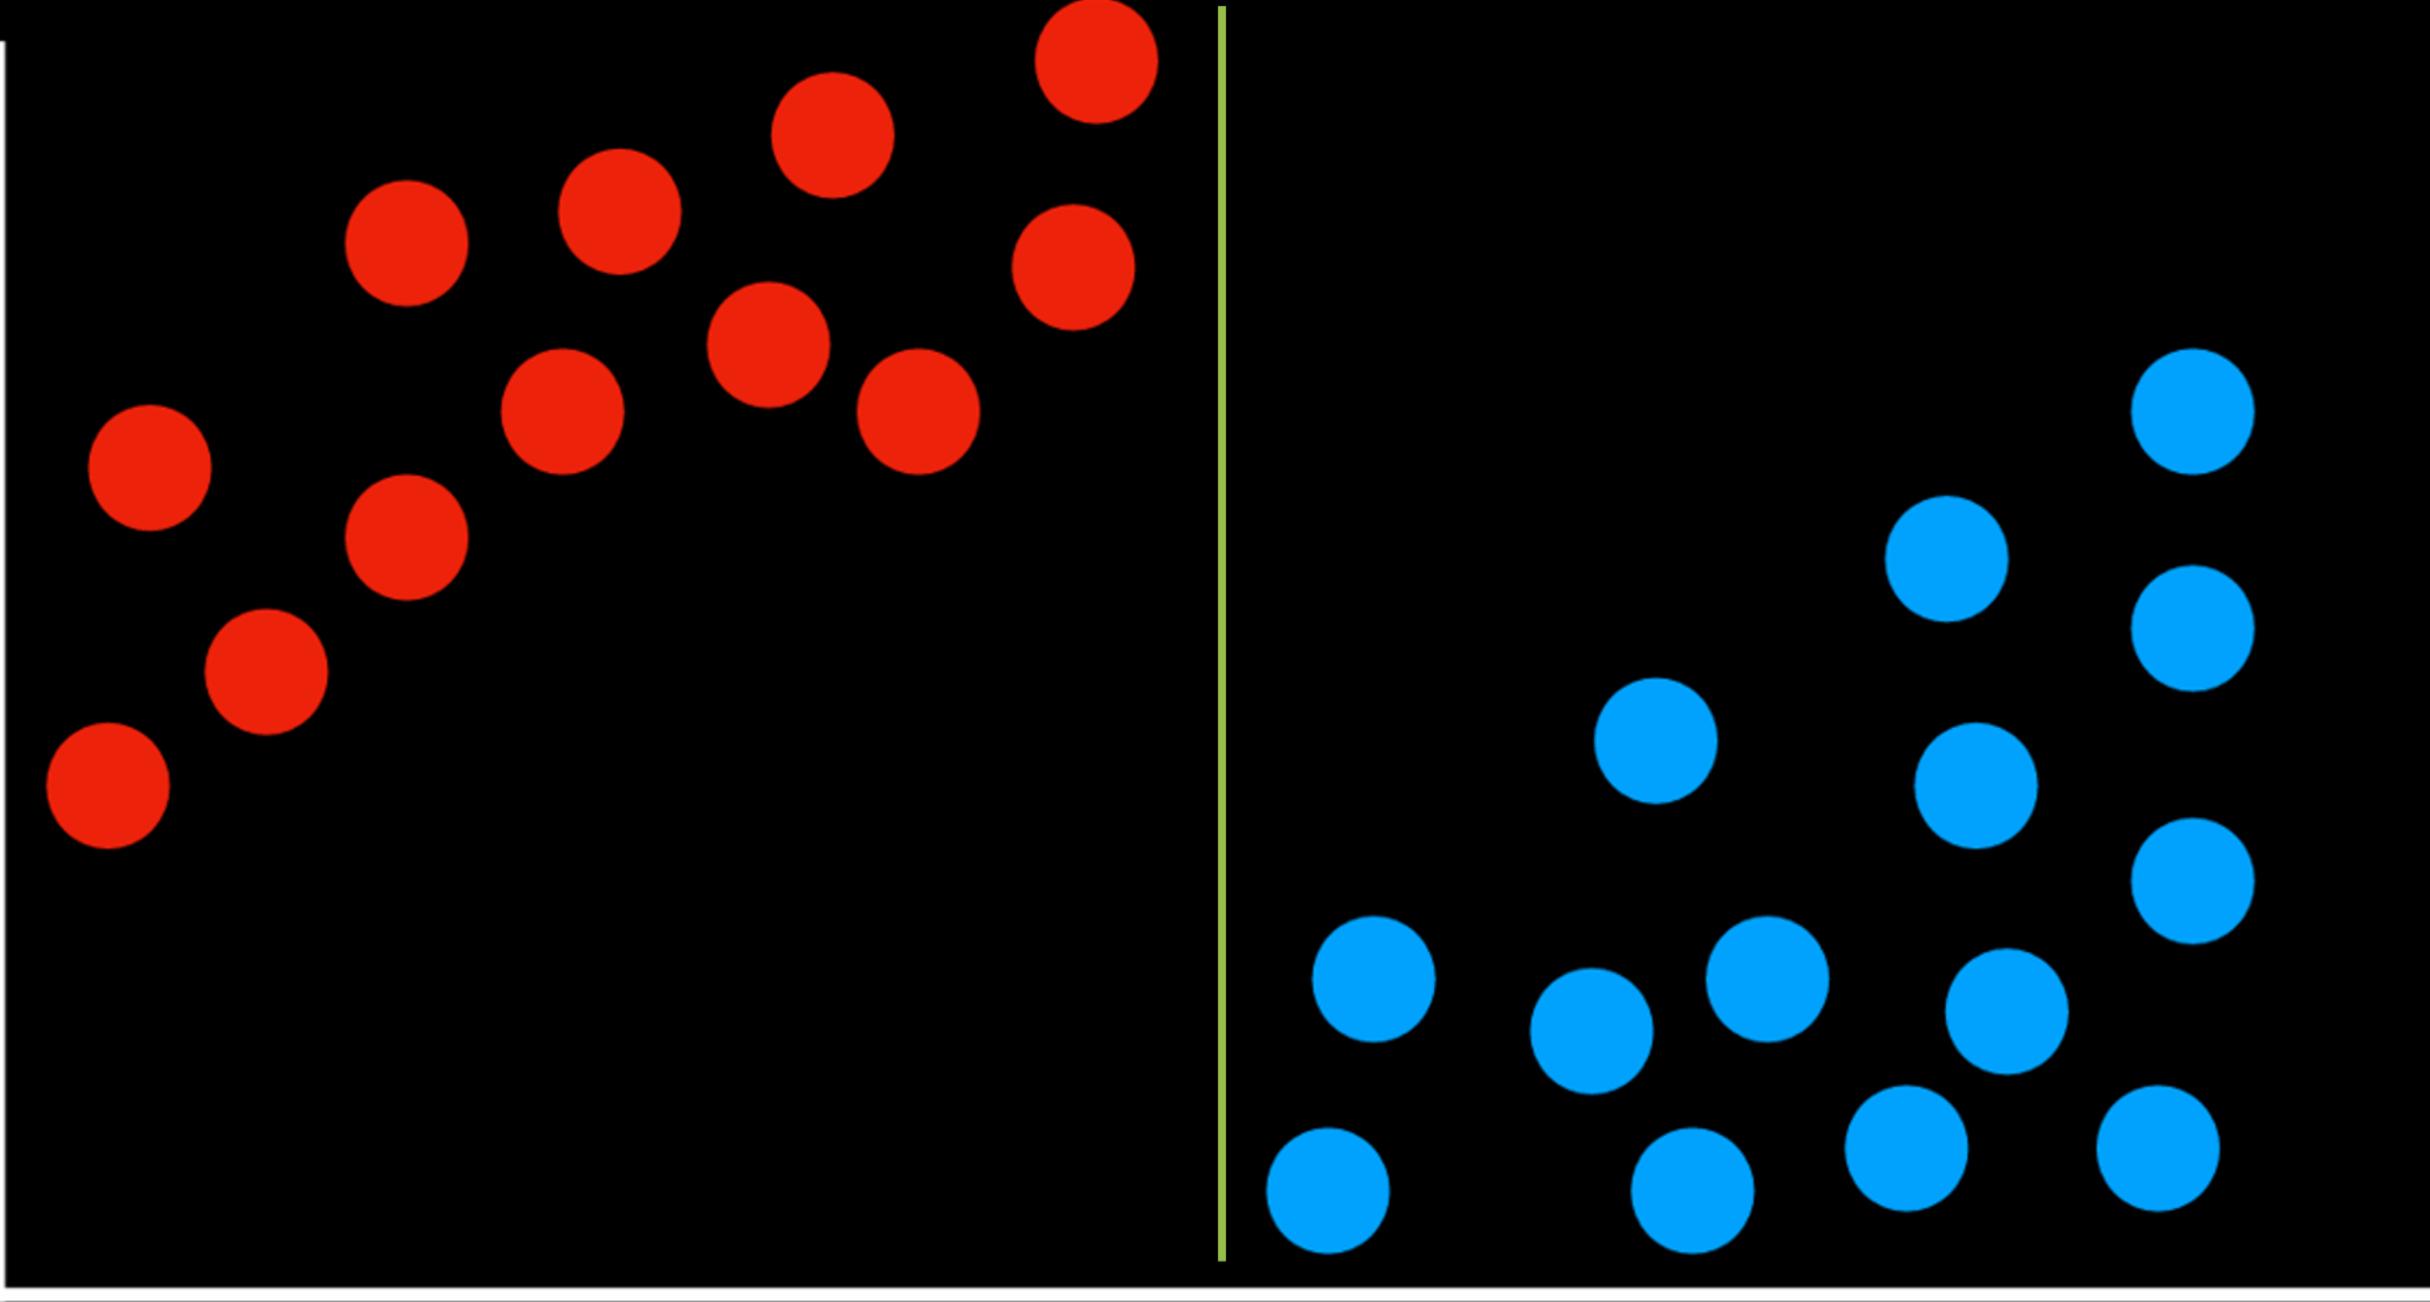
pressure

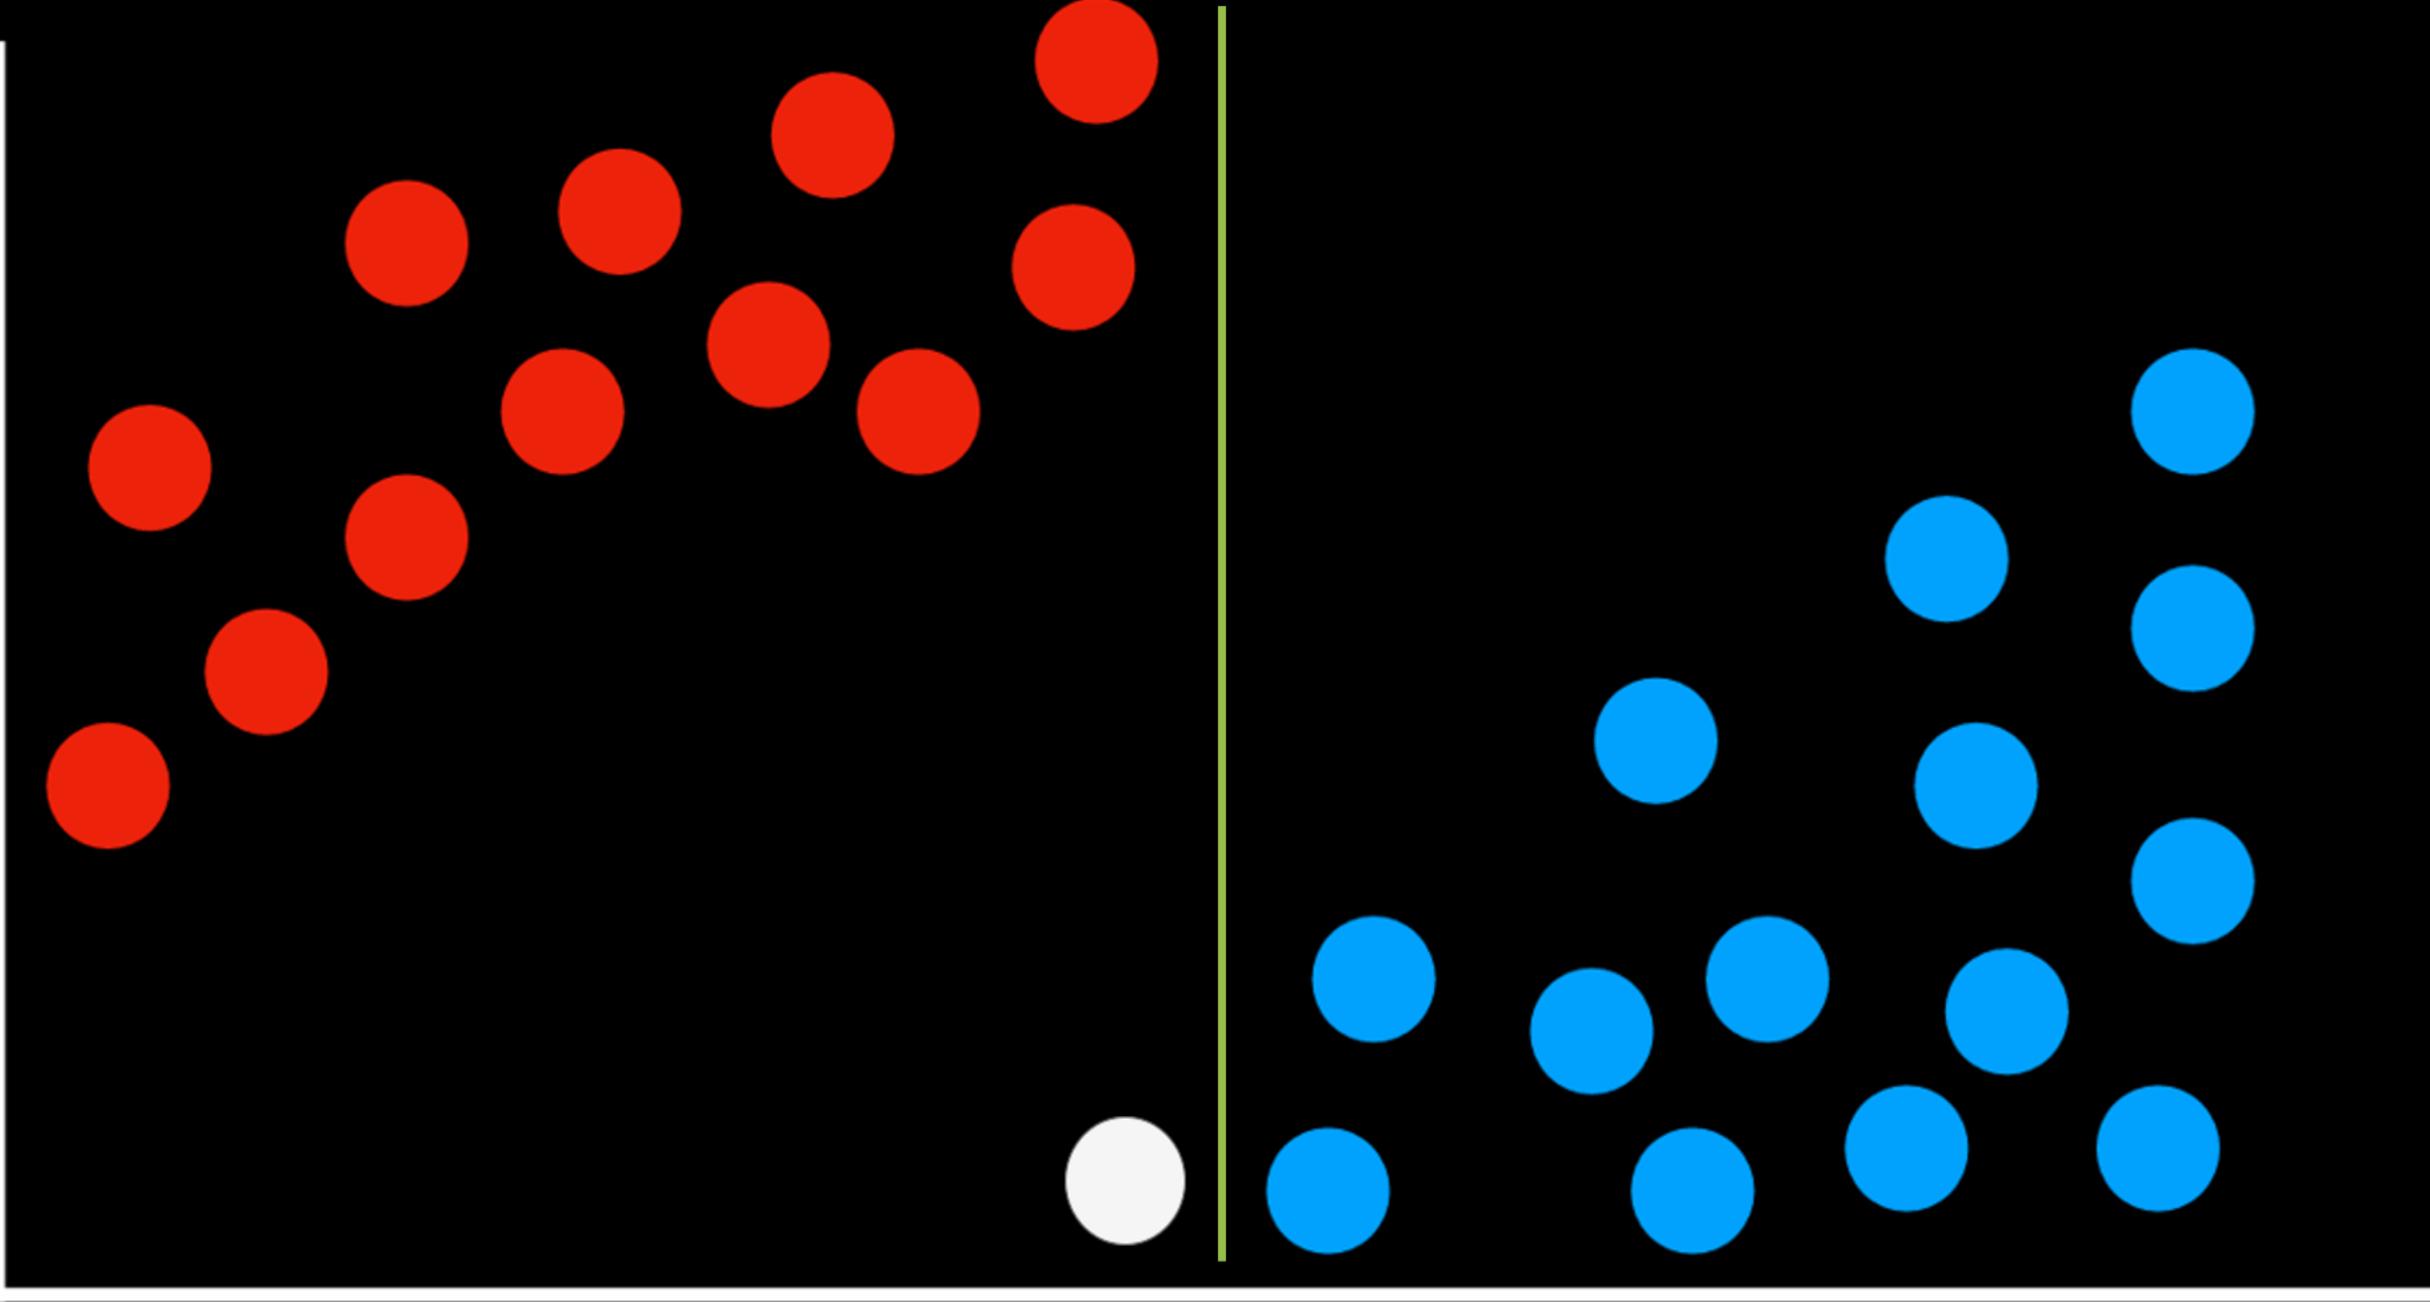
humidity

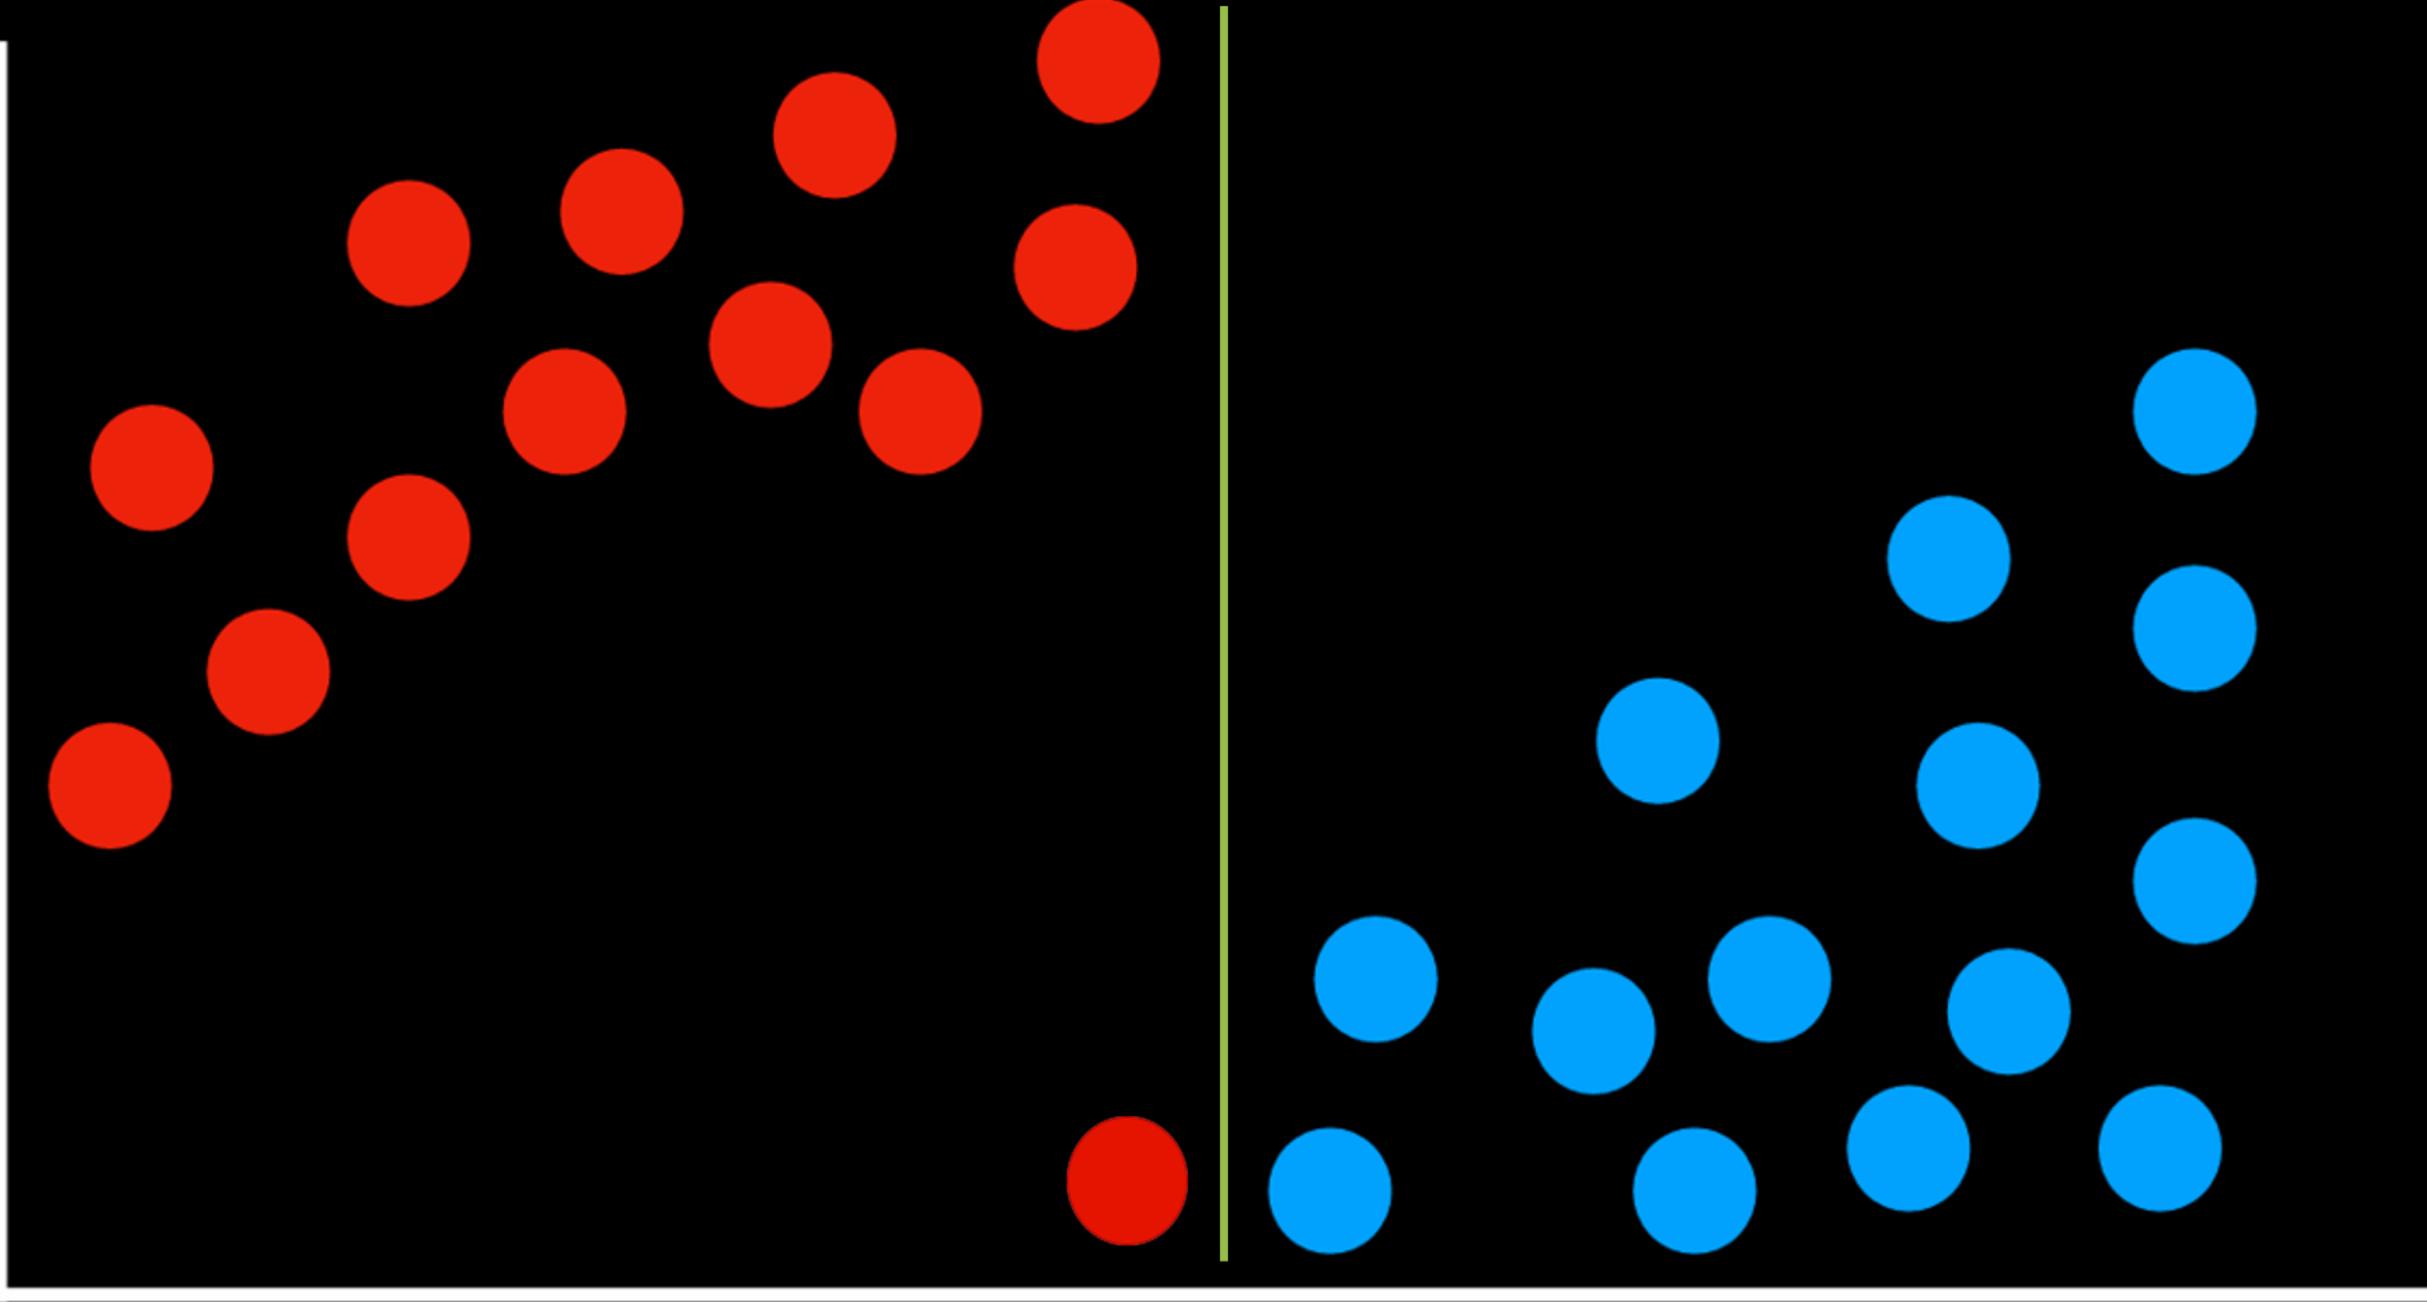


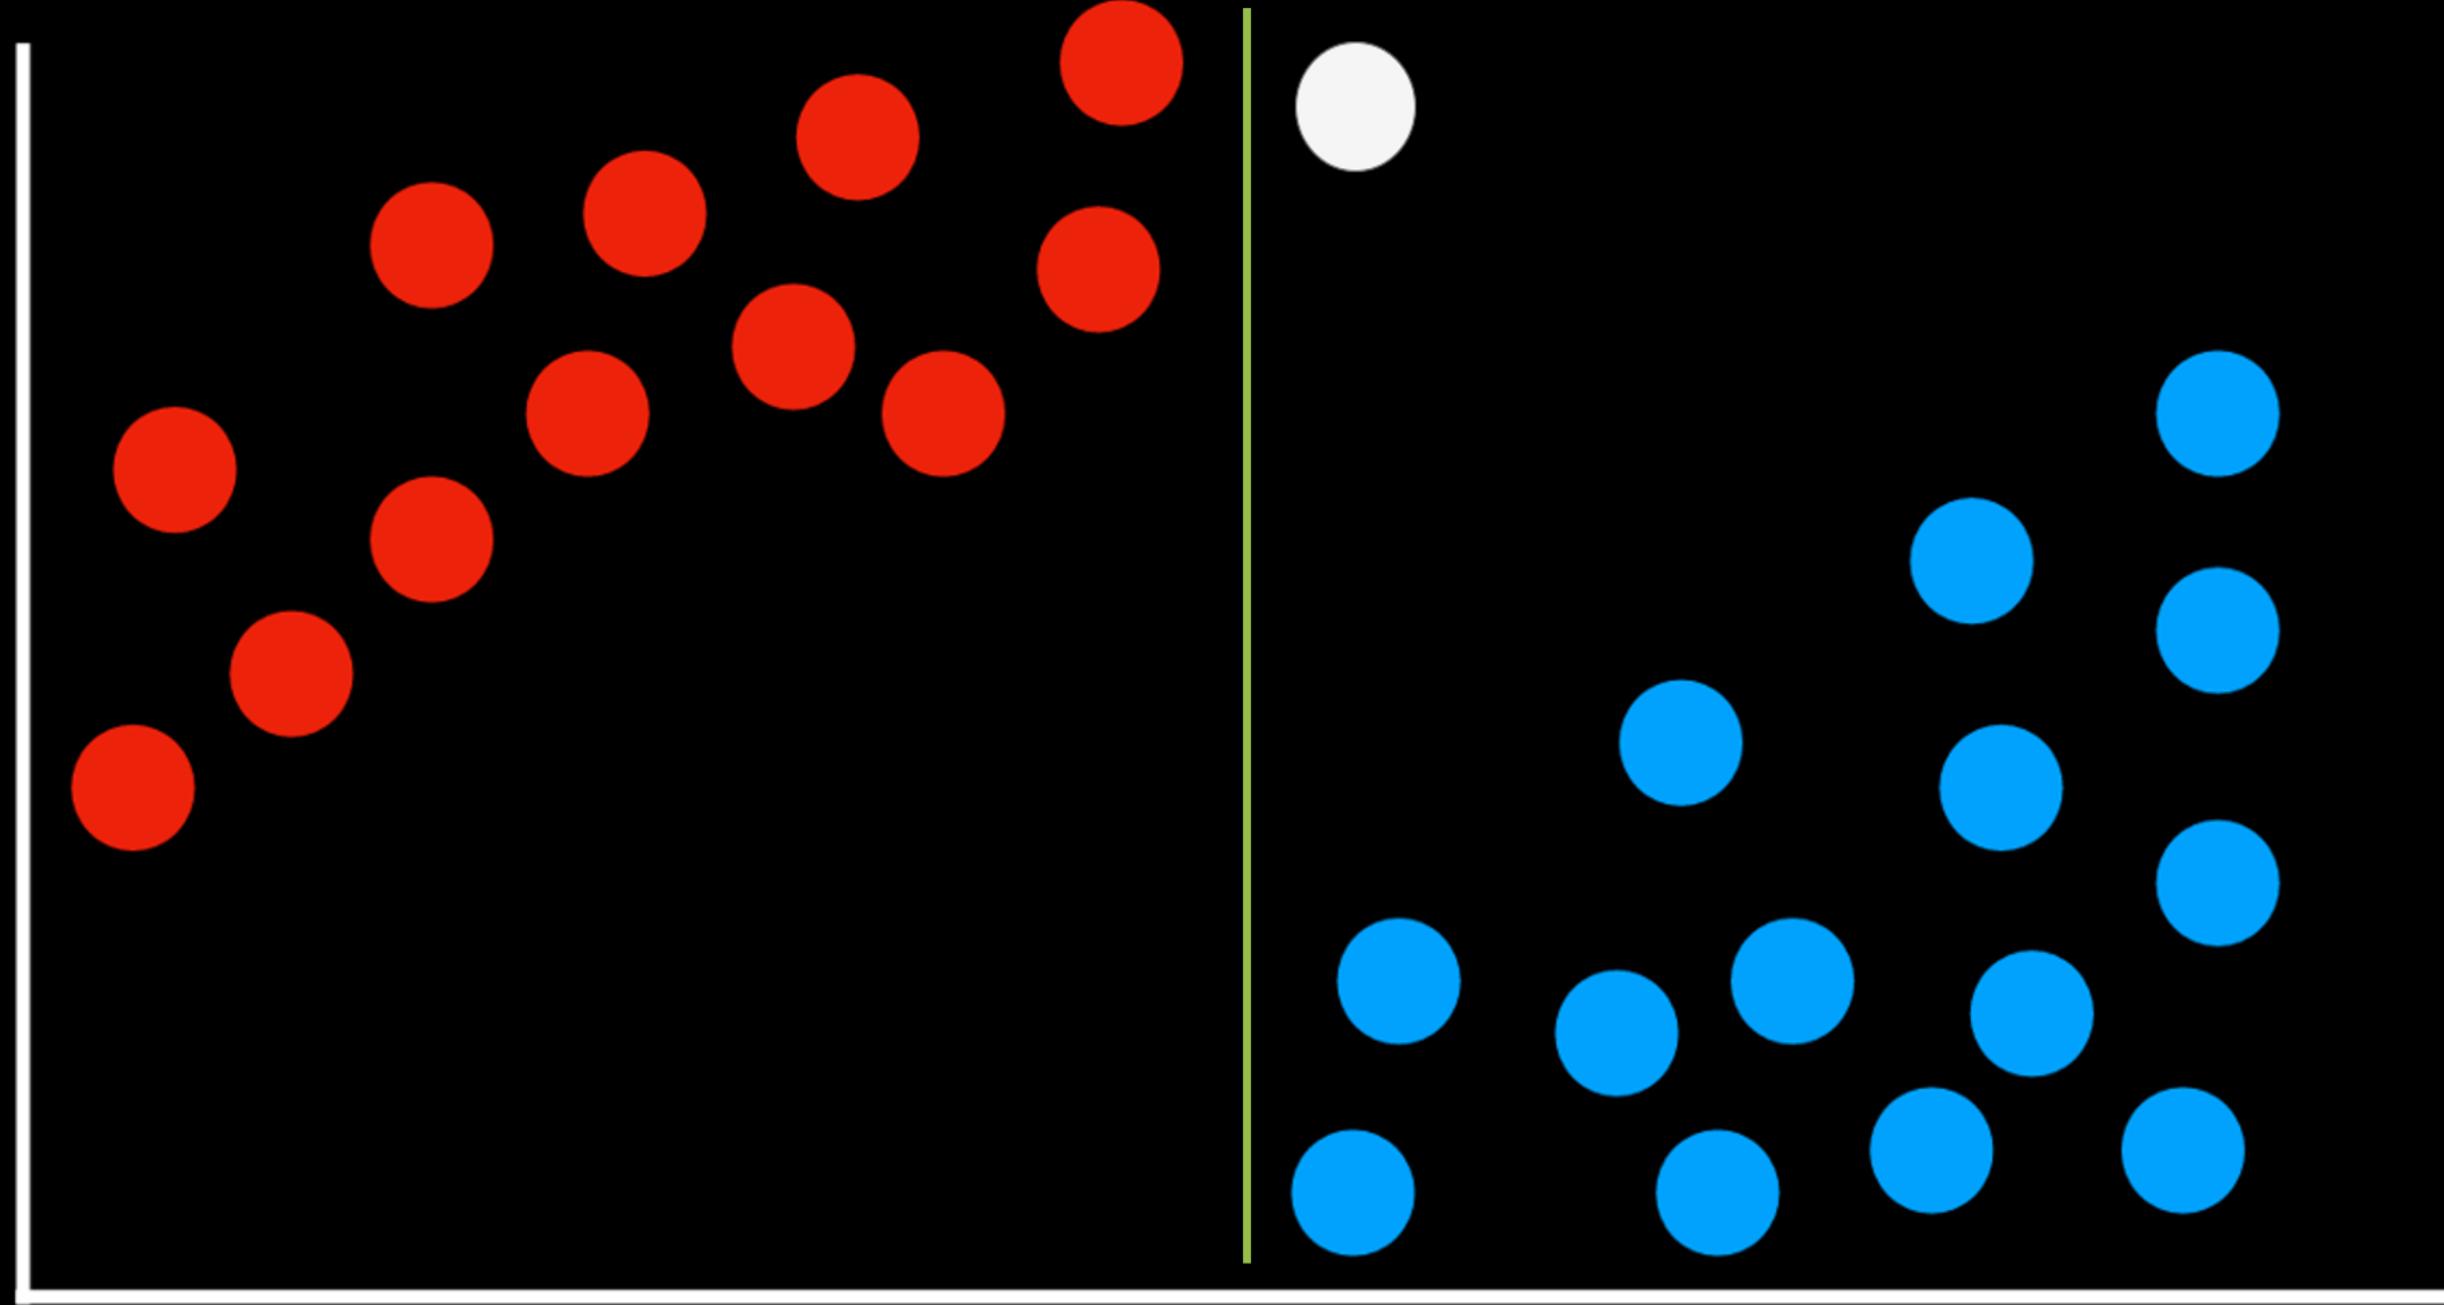
# Support Vector Machines

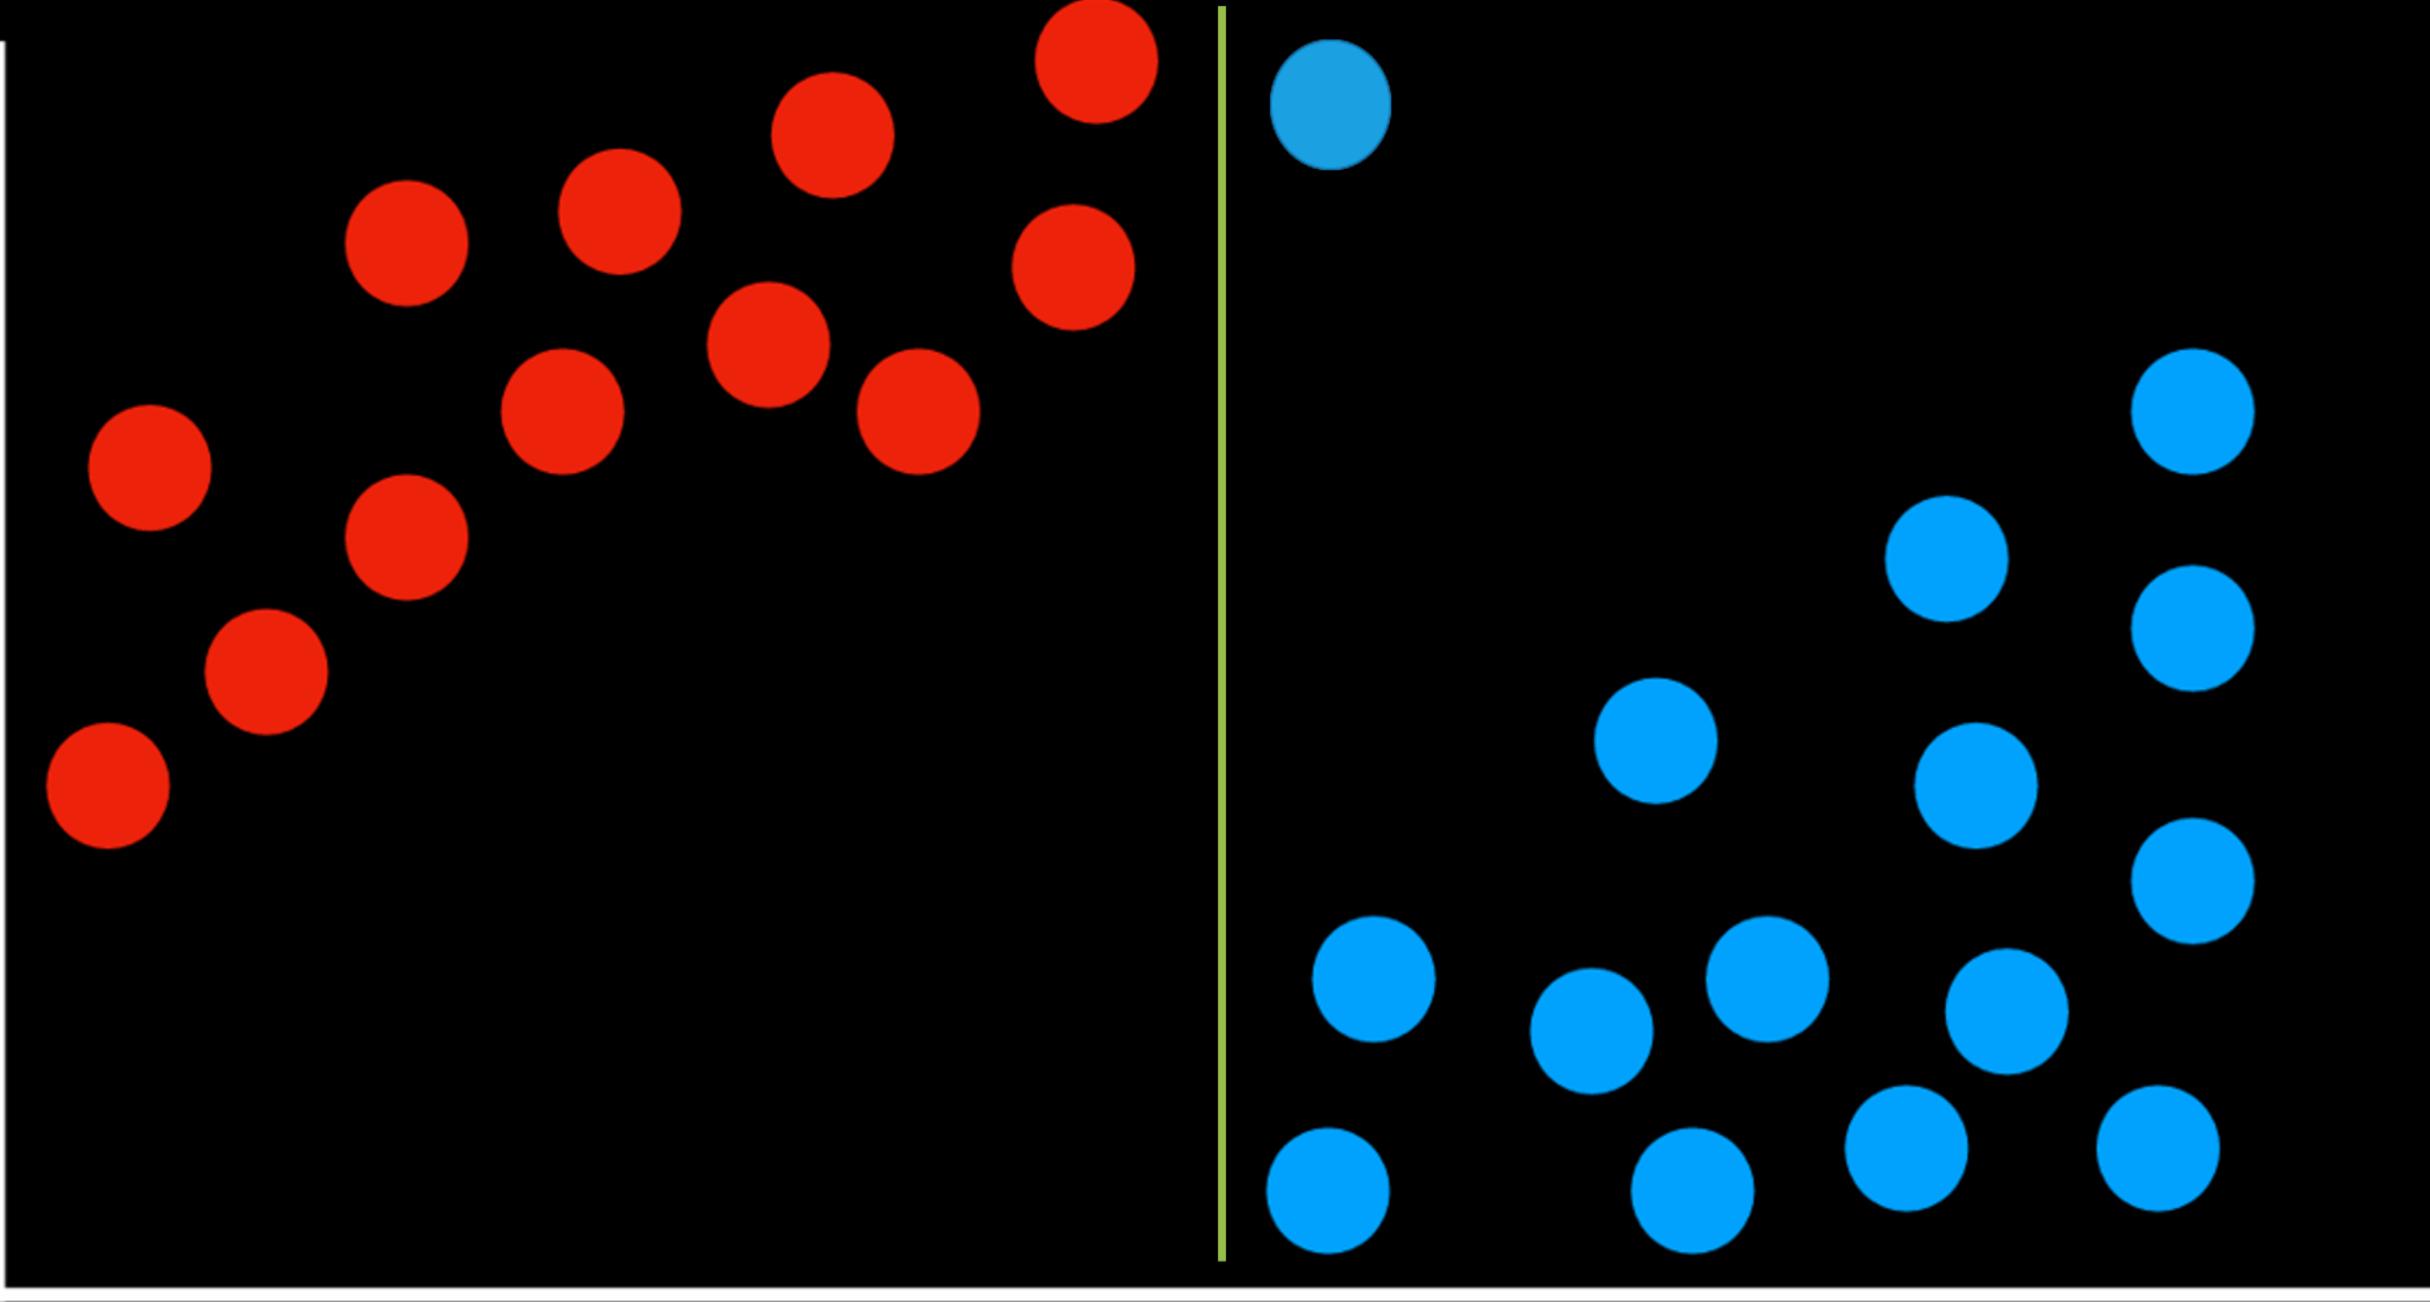


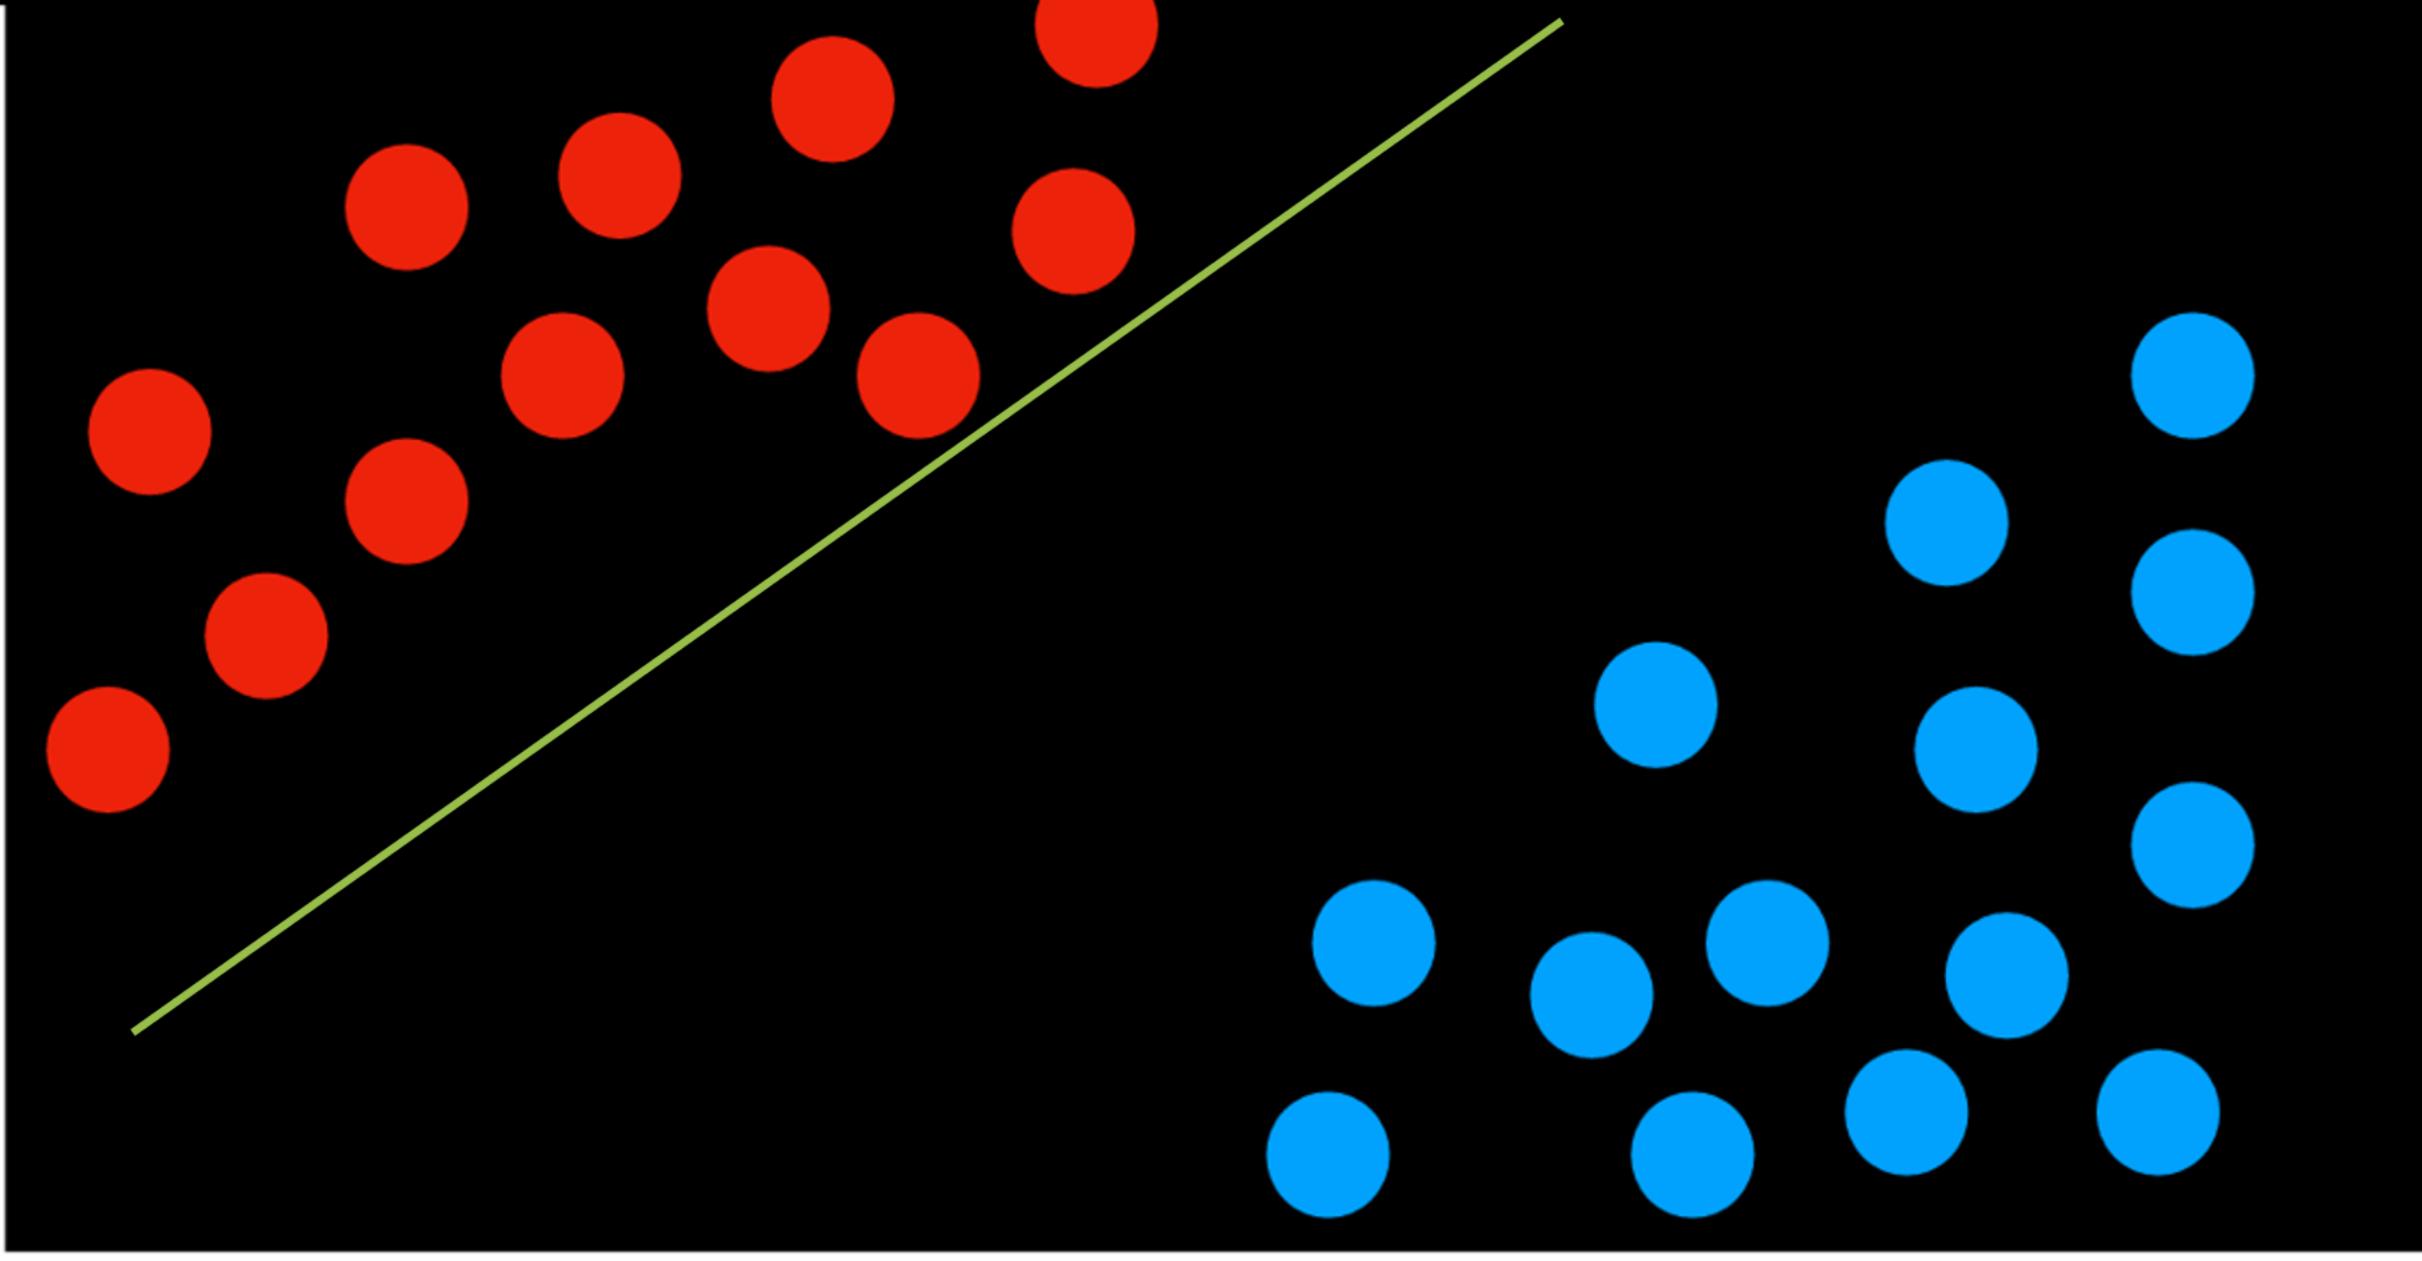


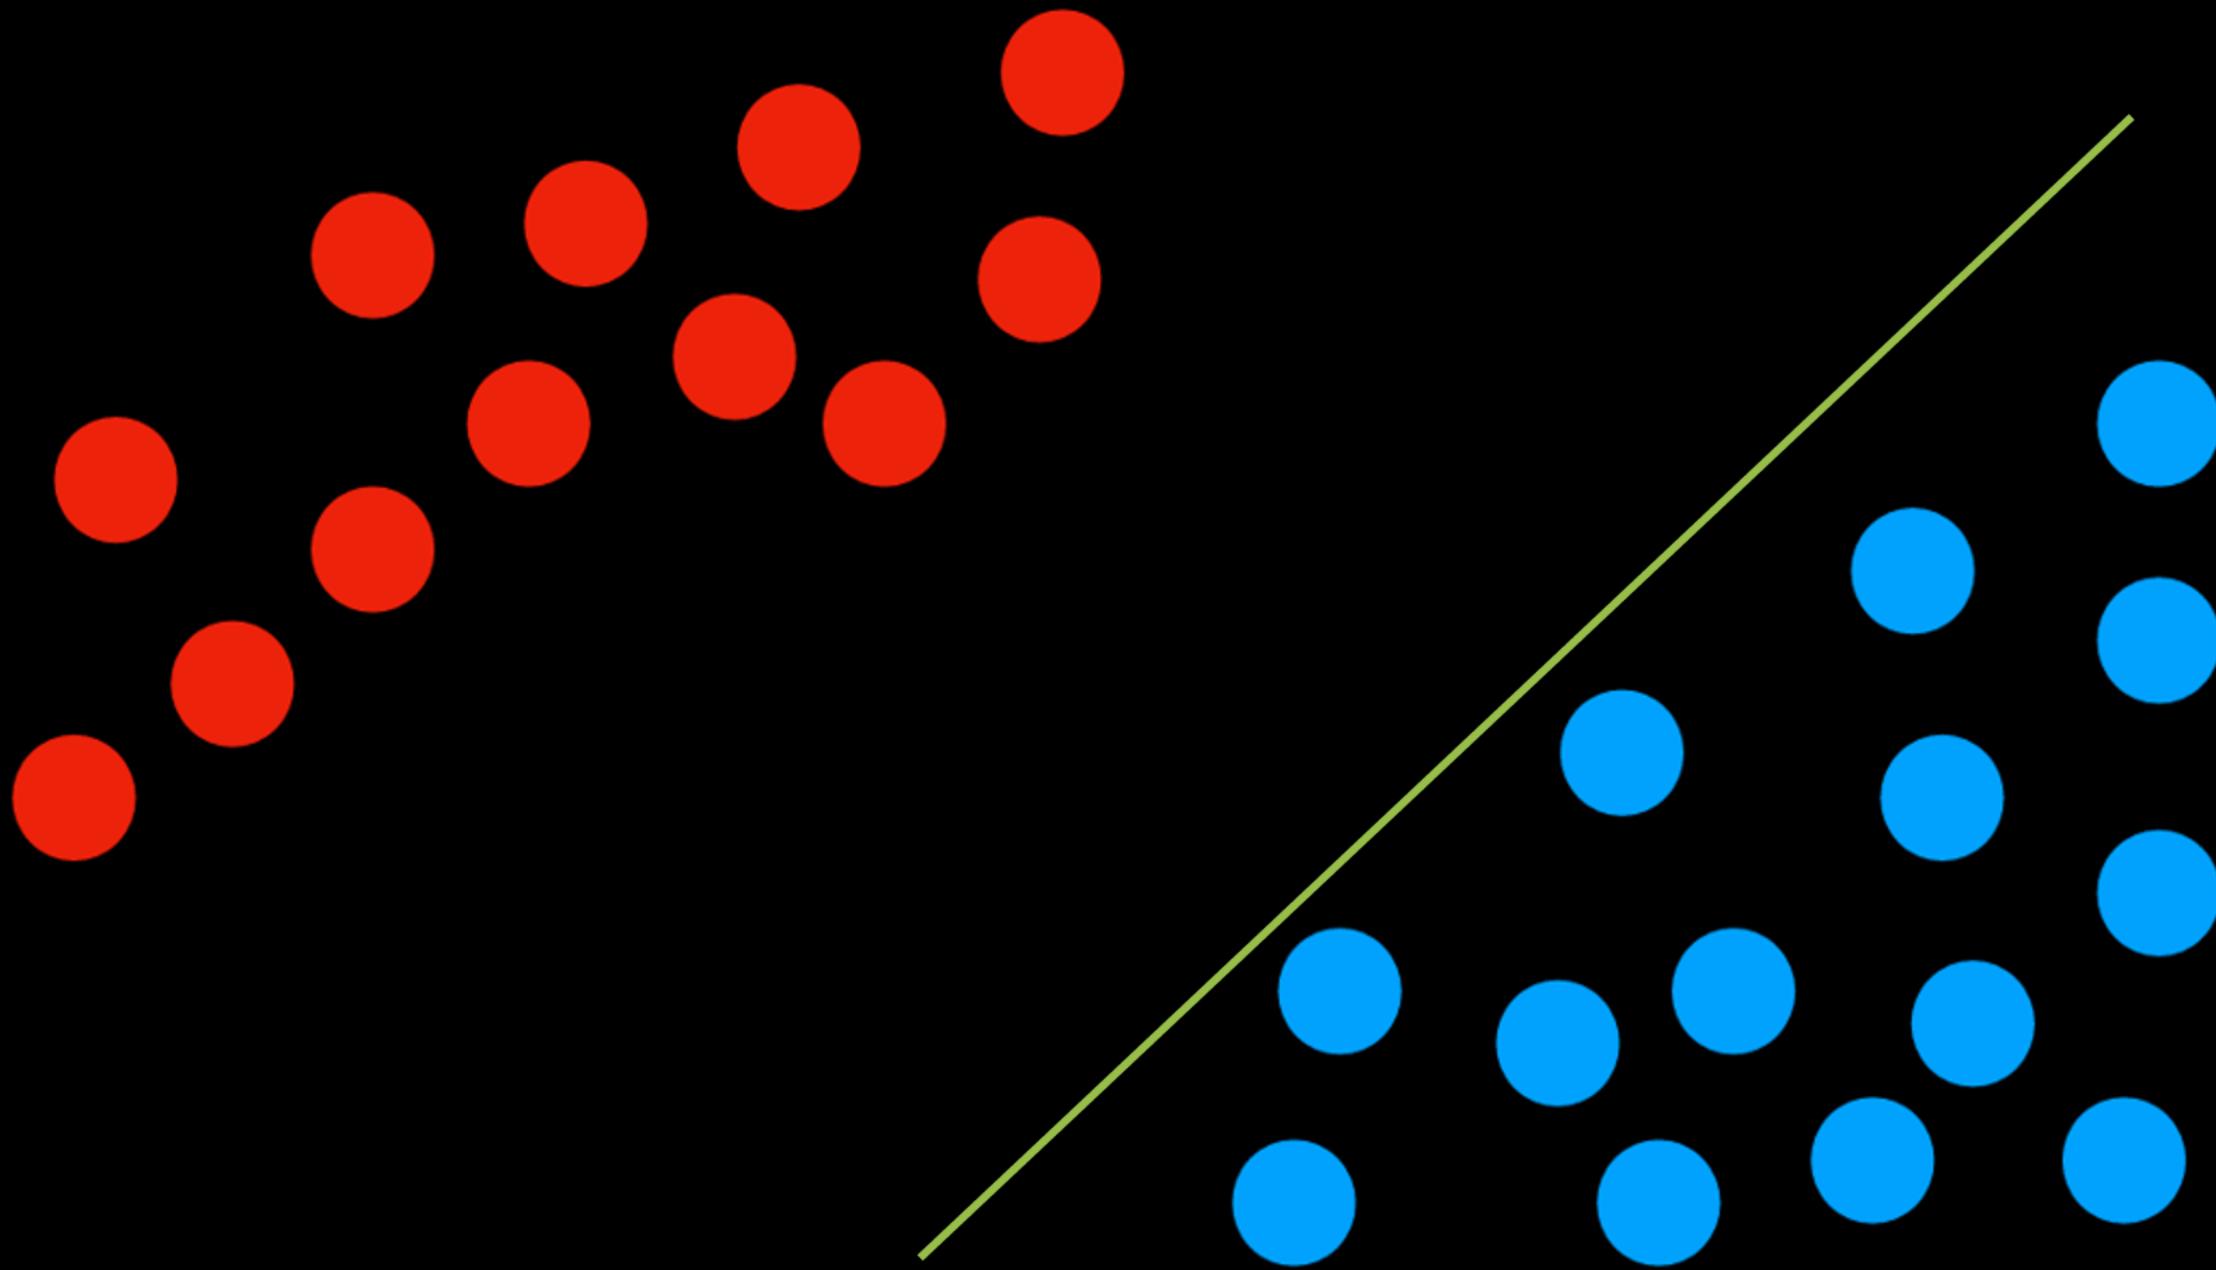


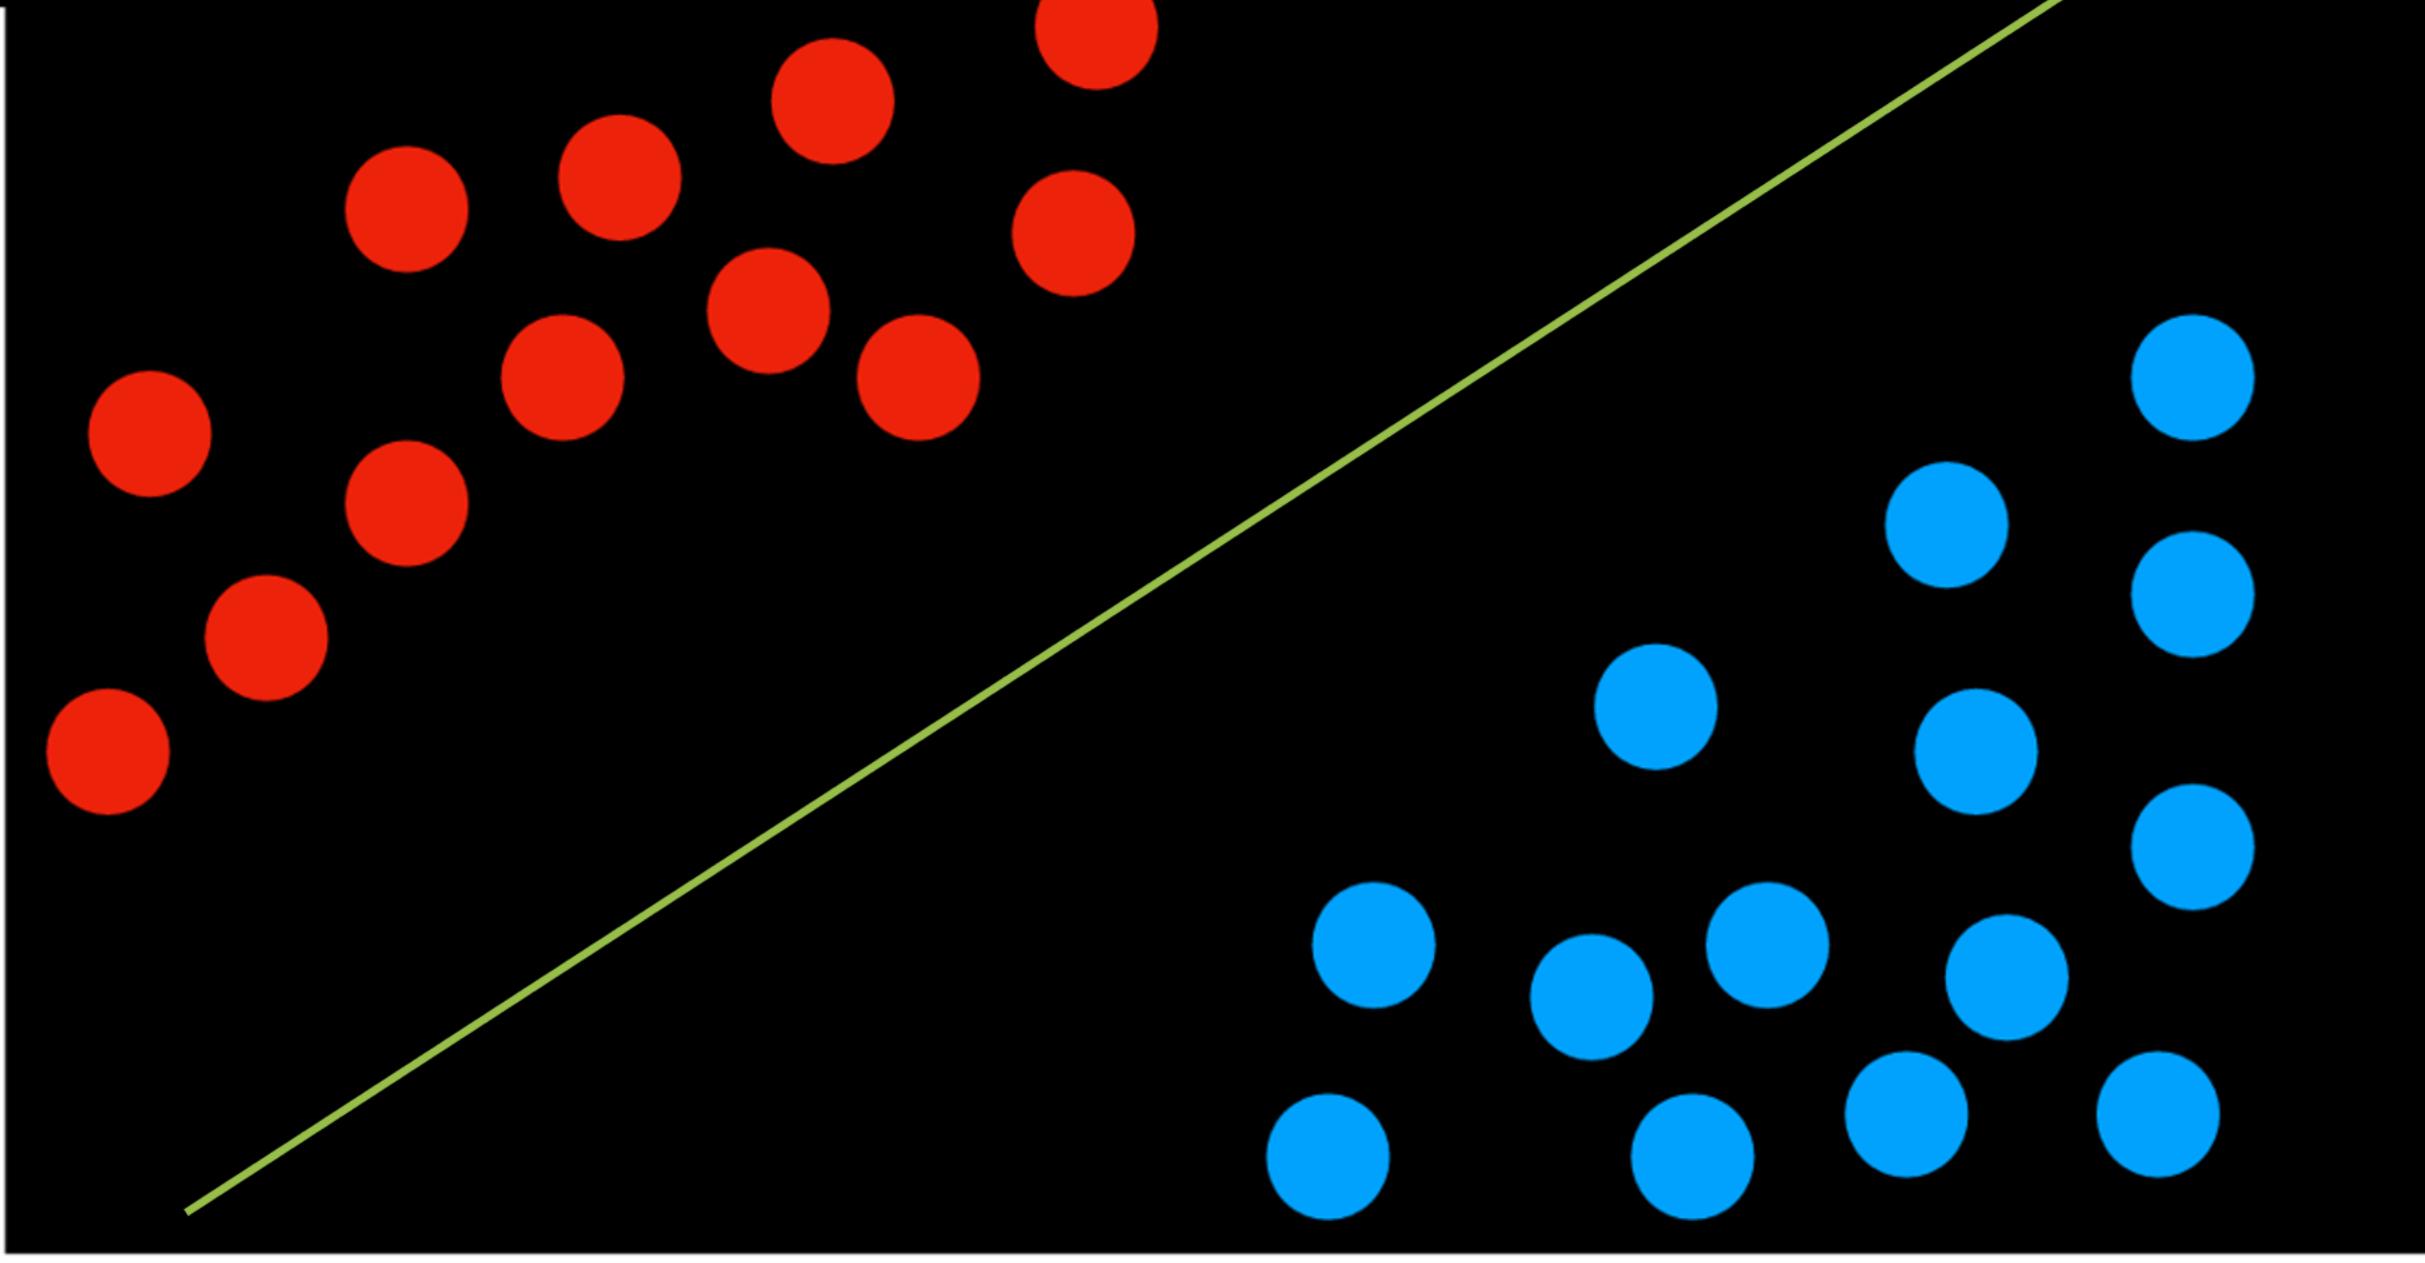






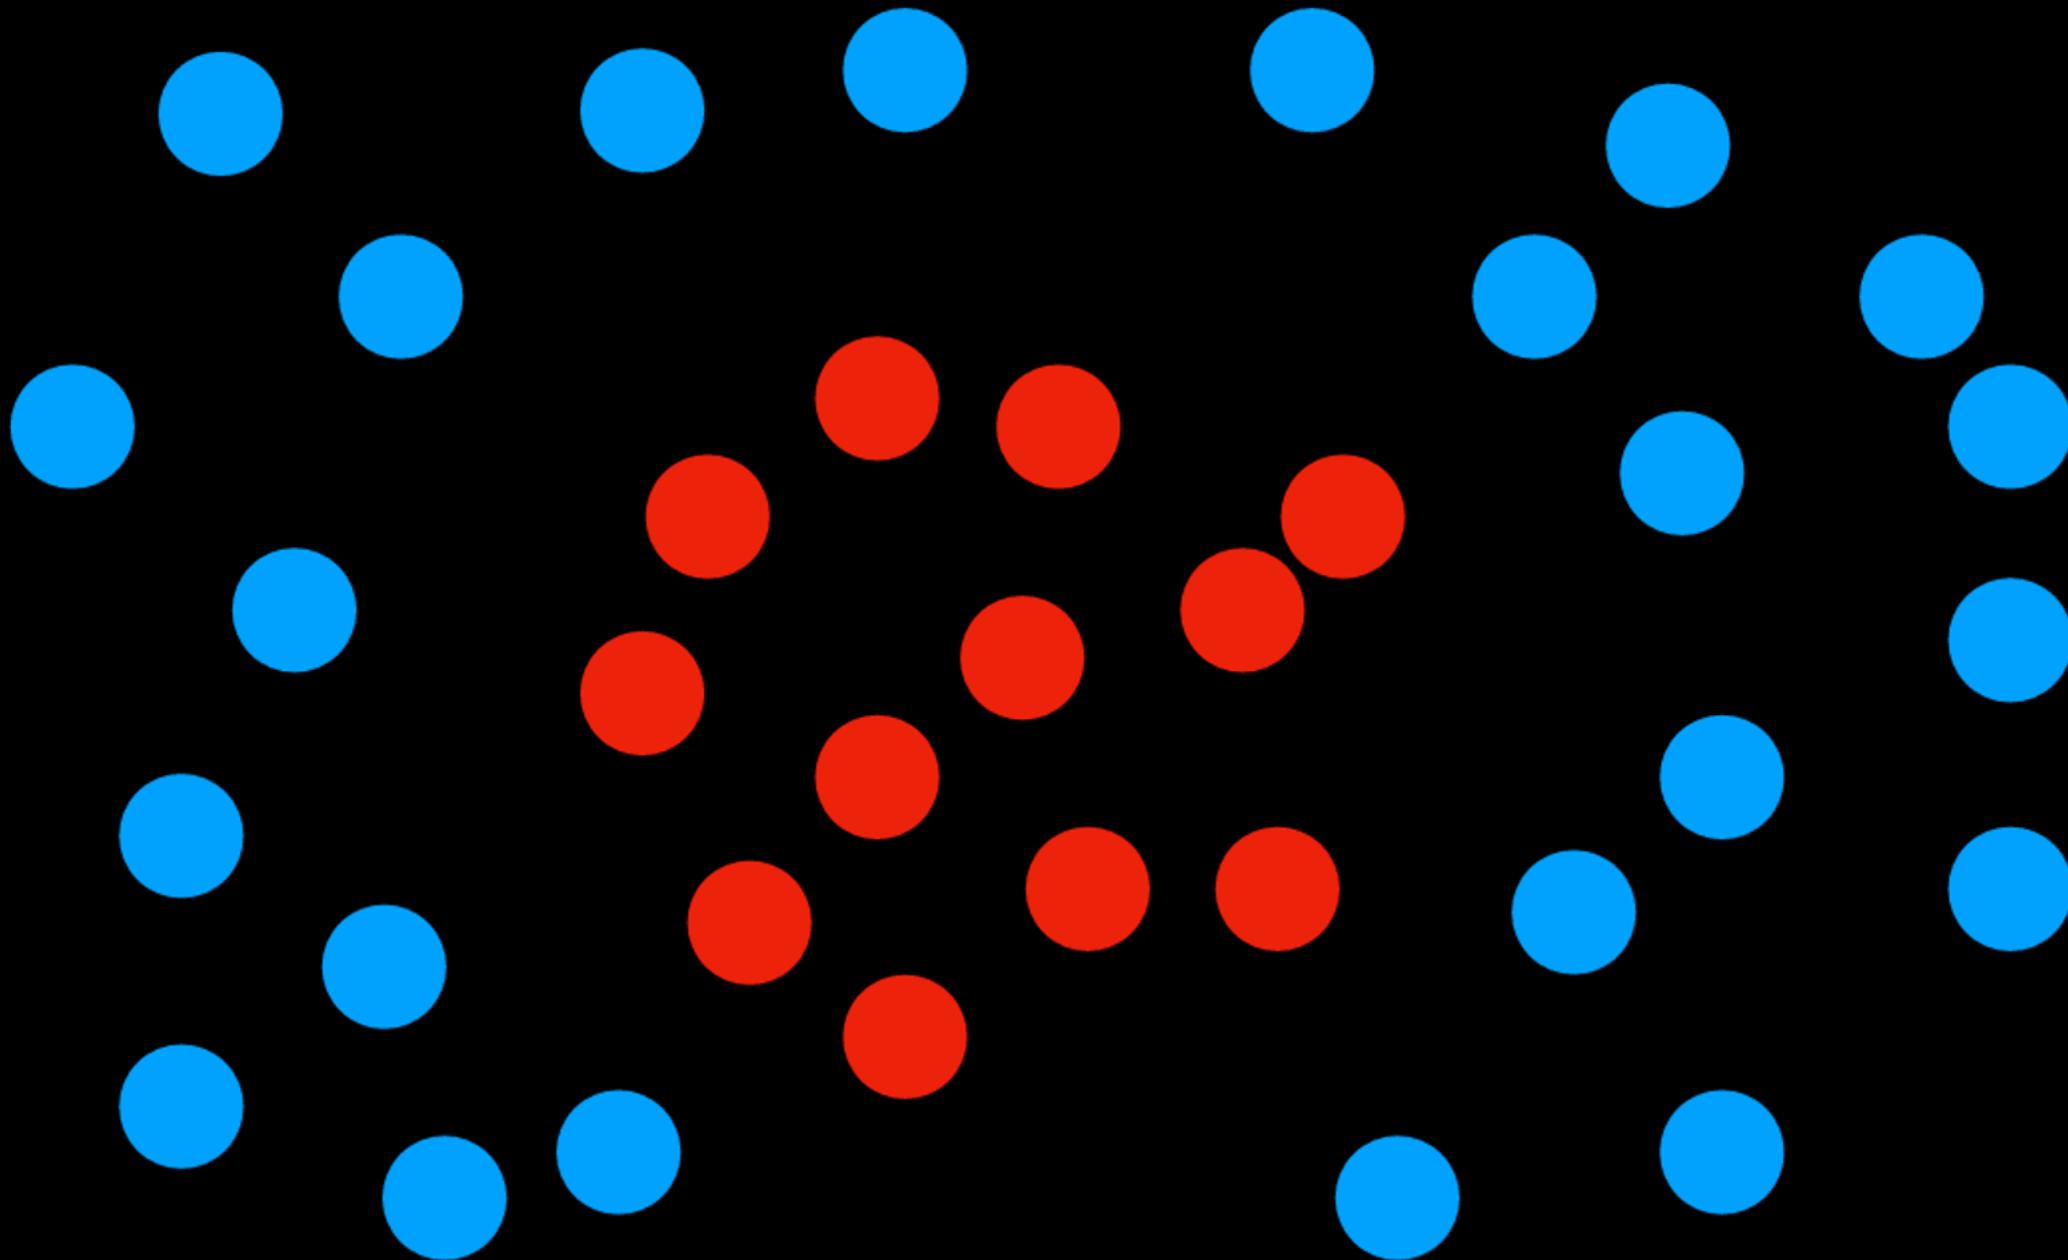


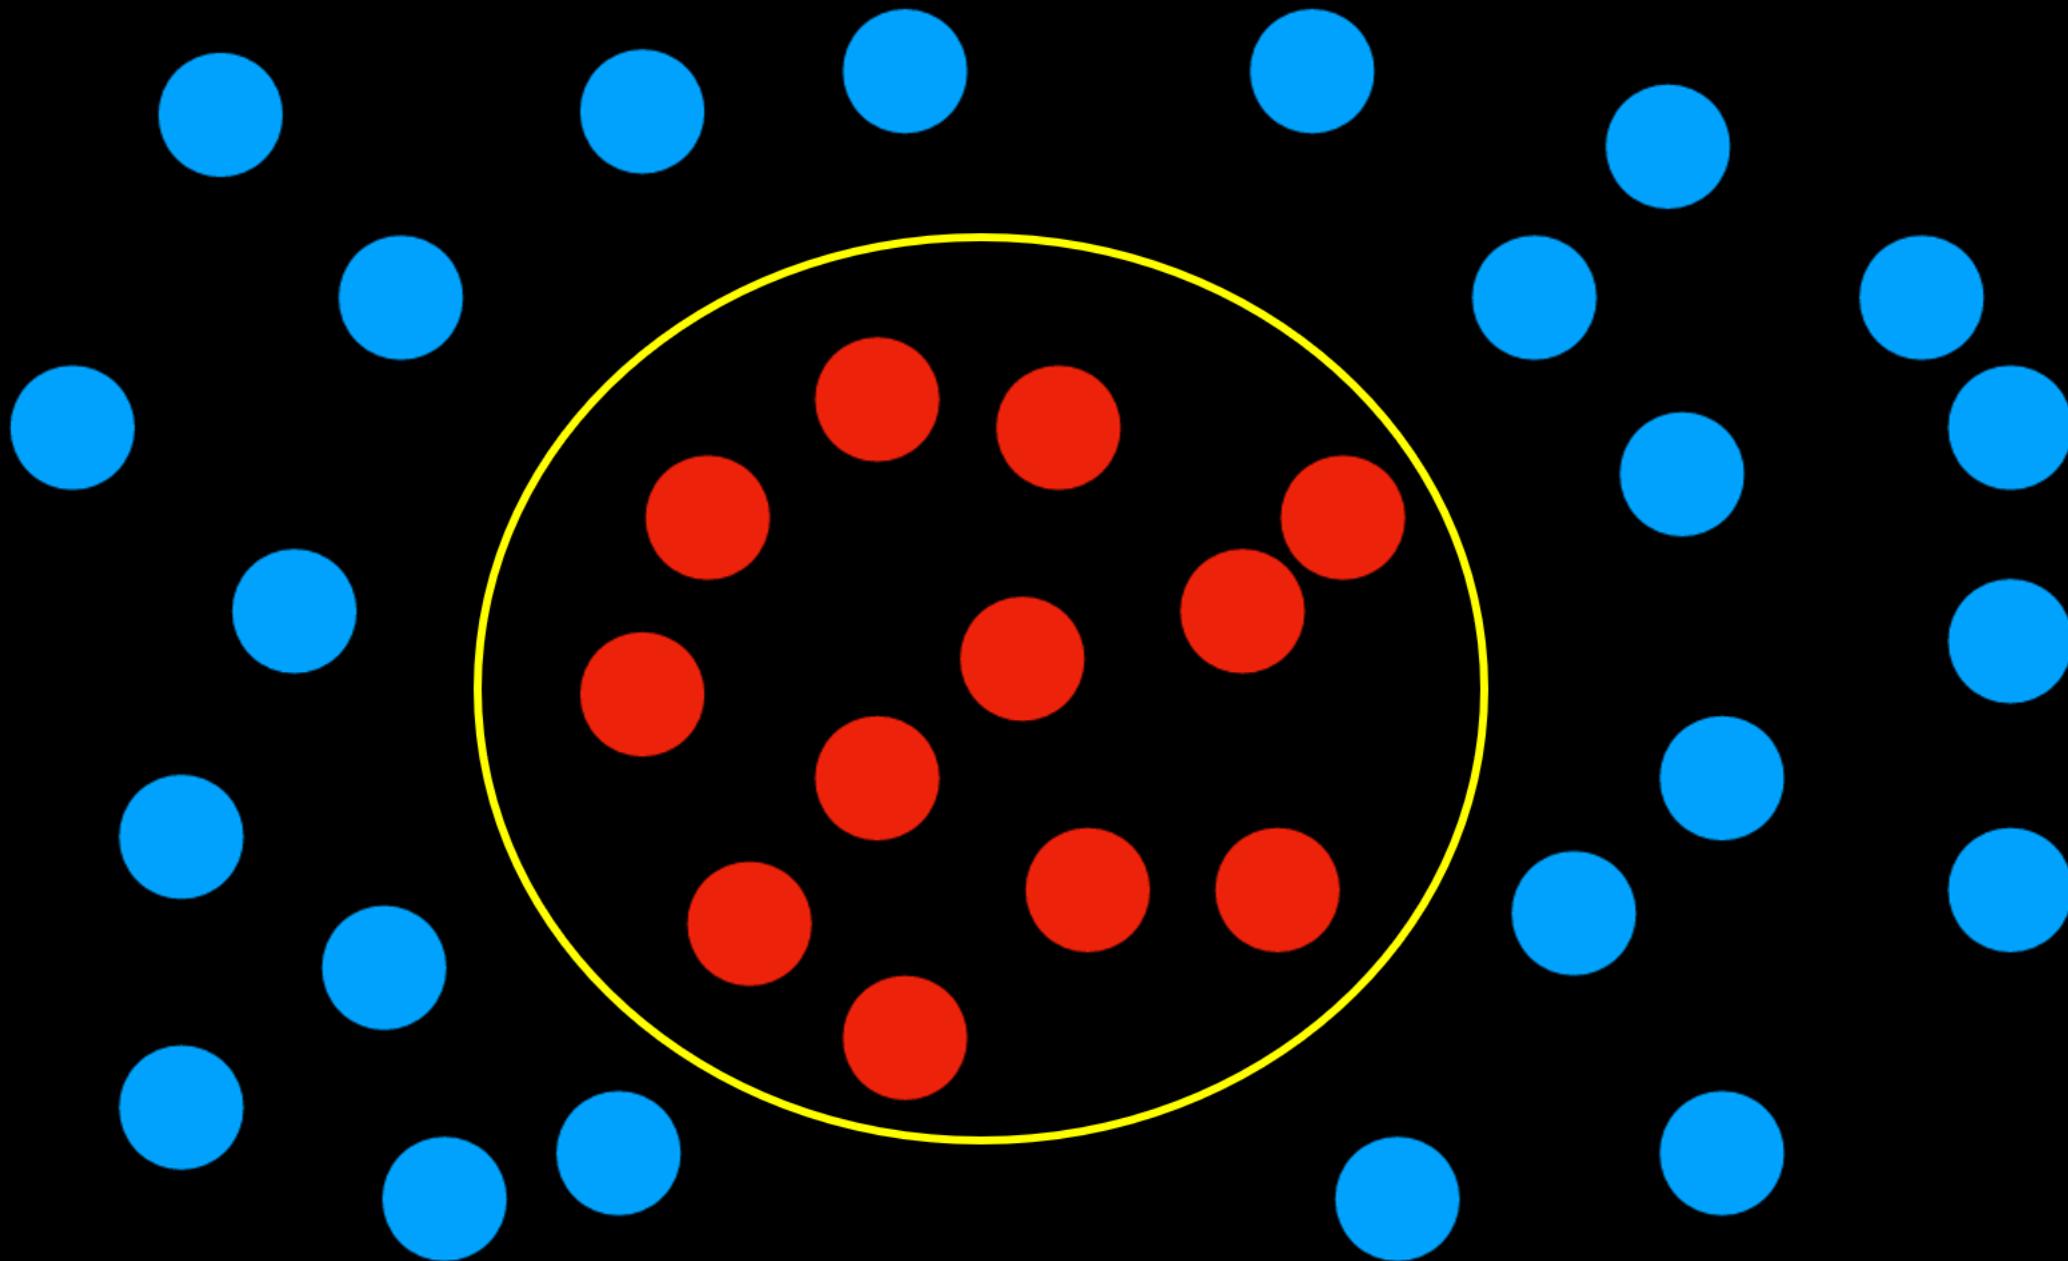




# maximum margin separator

boundary that maximizes the distance  
between any of the data points



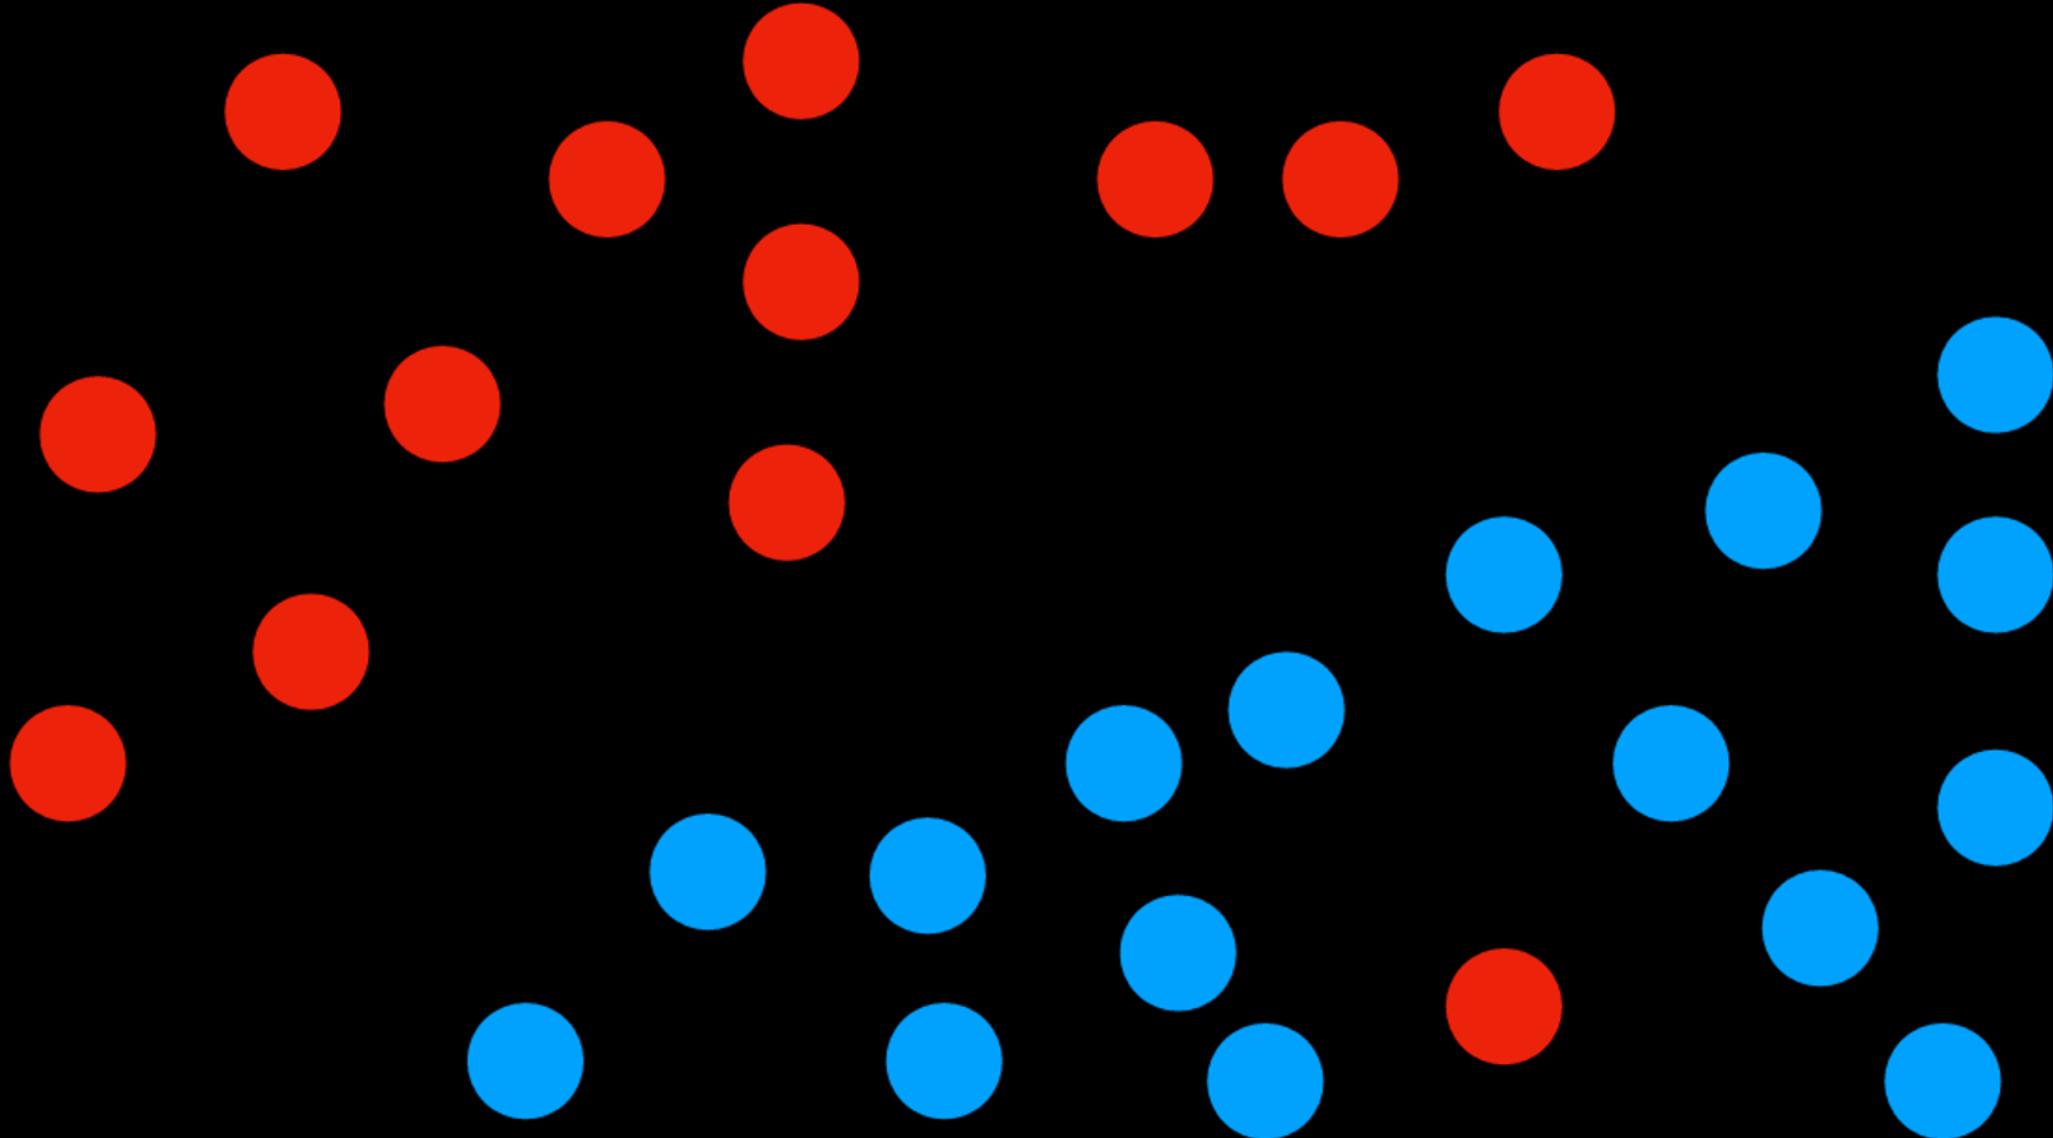


# overfitting

a model that fits too closely to a particular data set and therefore may fail to generalize to future data

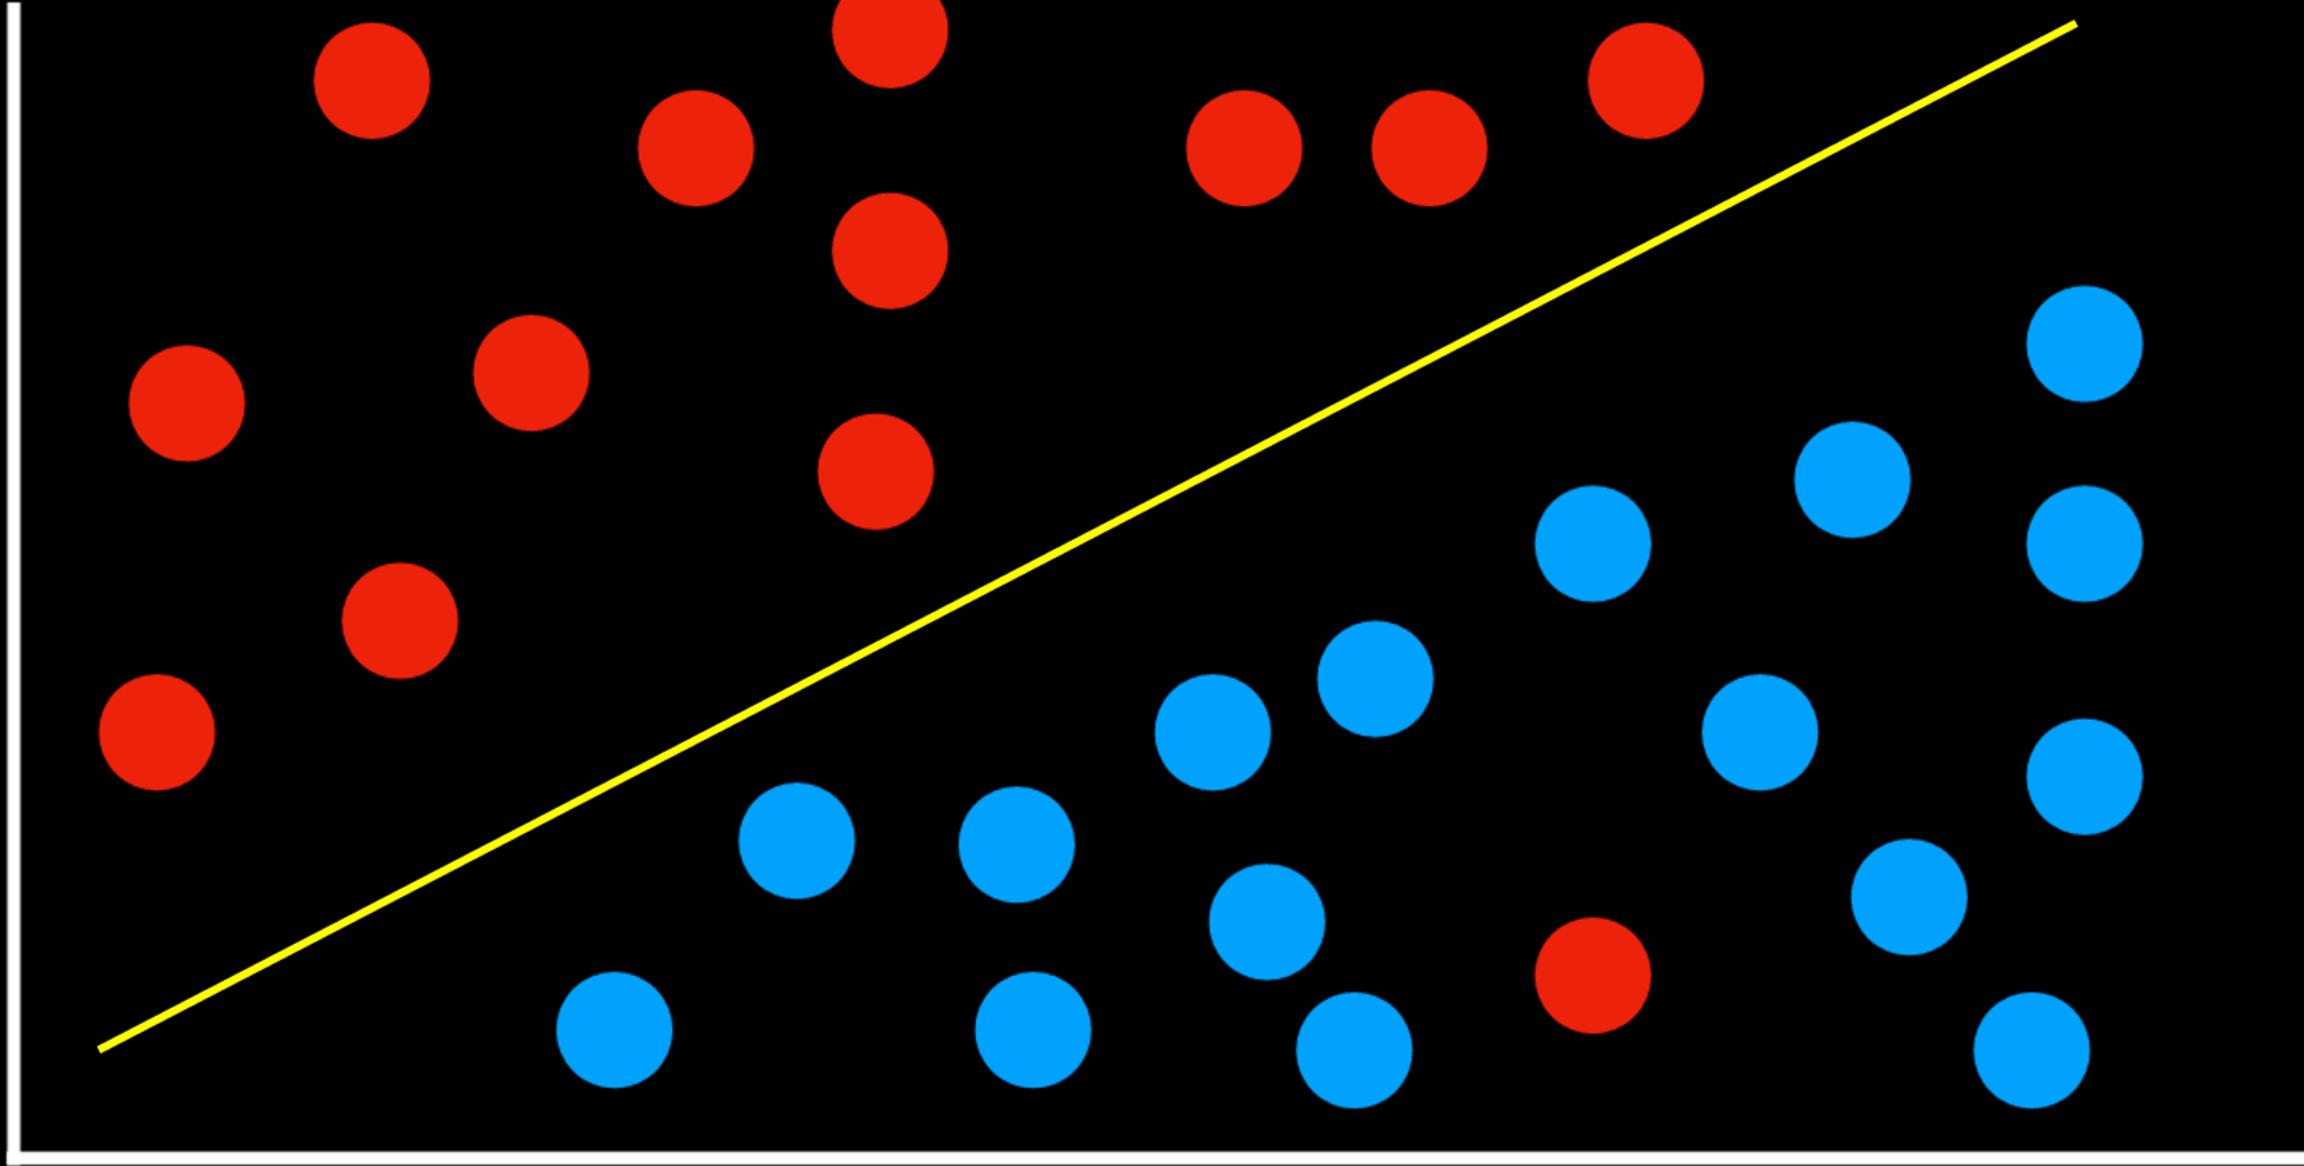
pressure

humidity



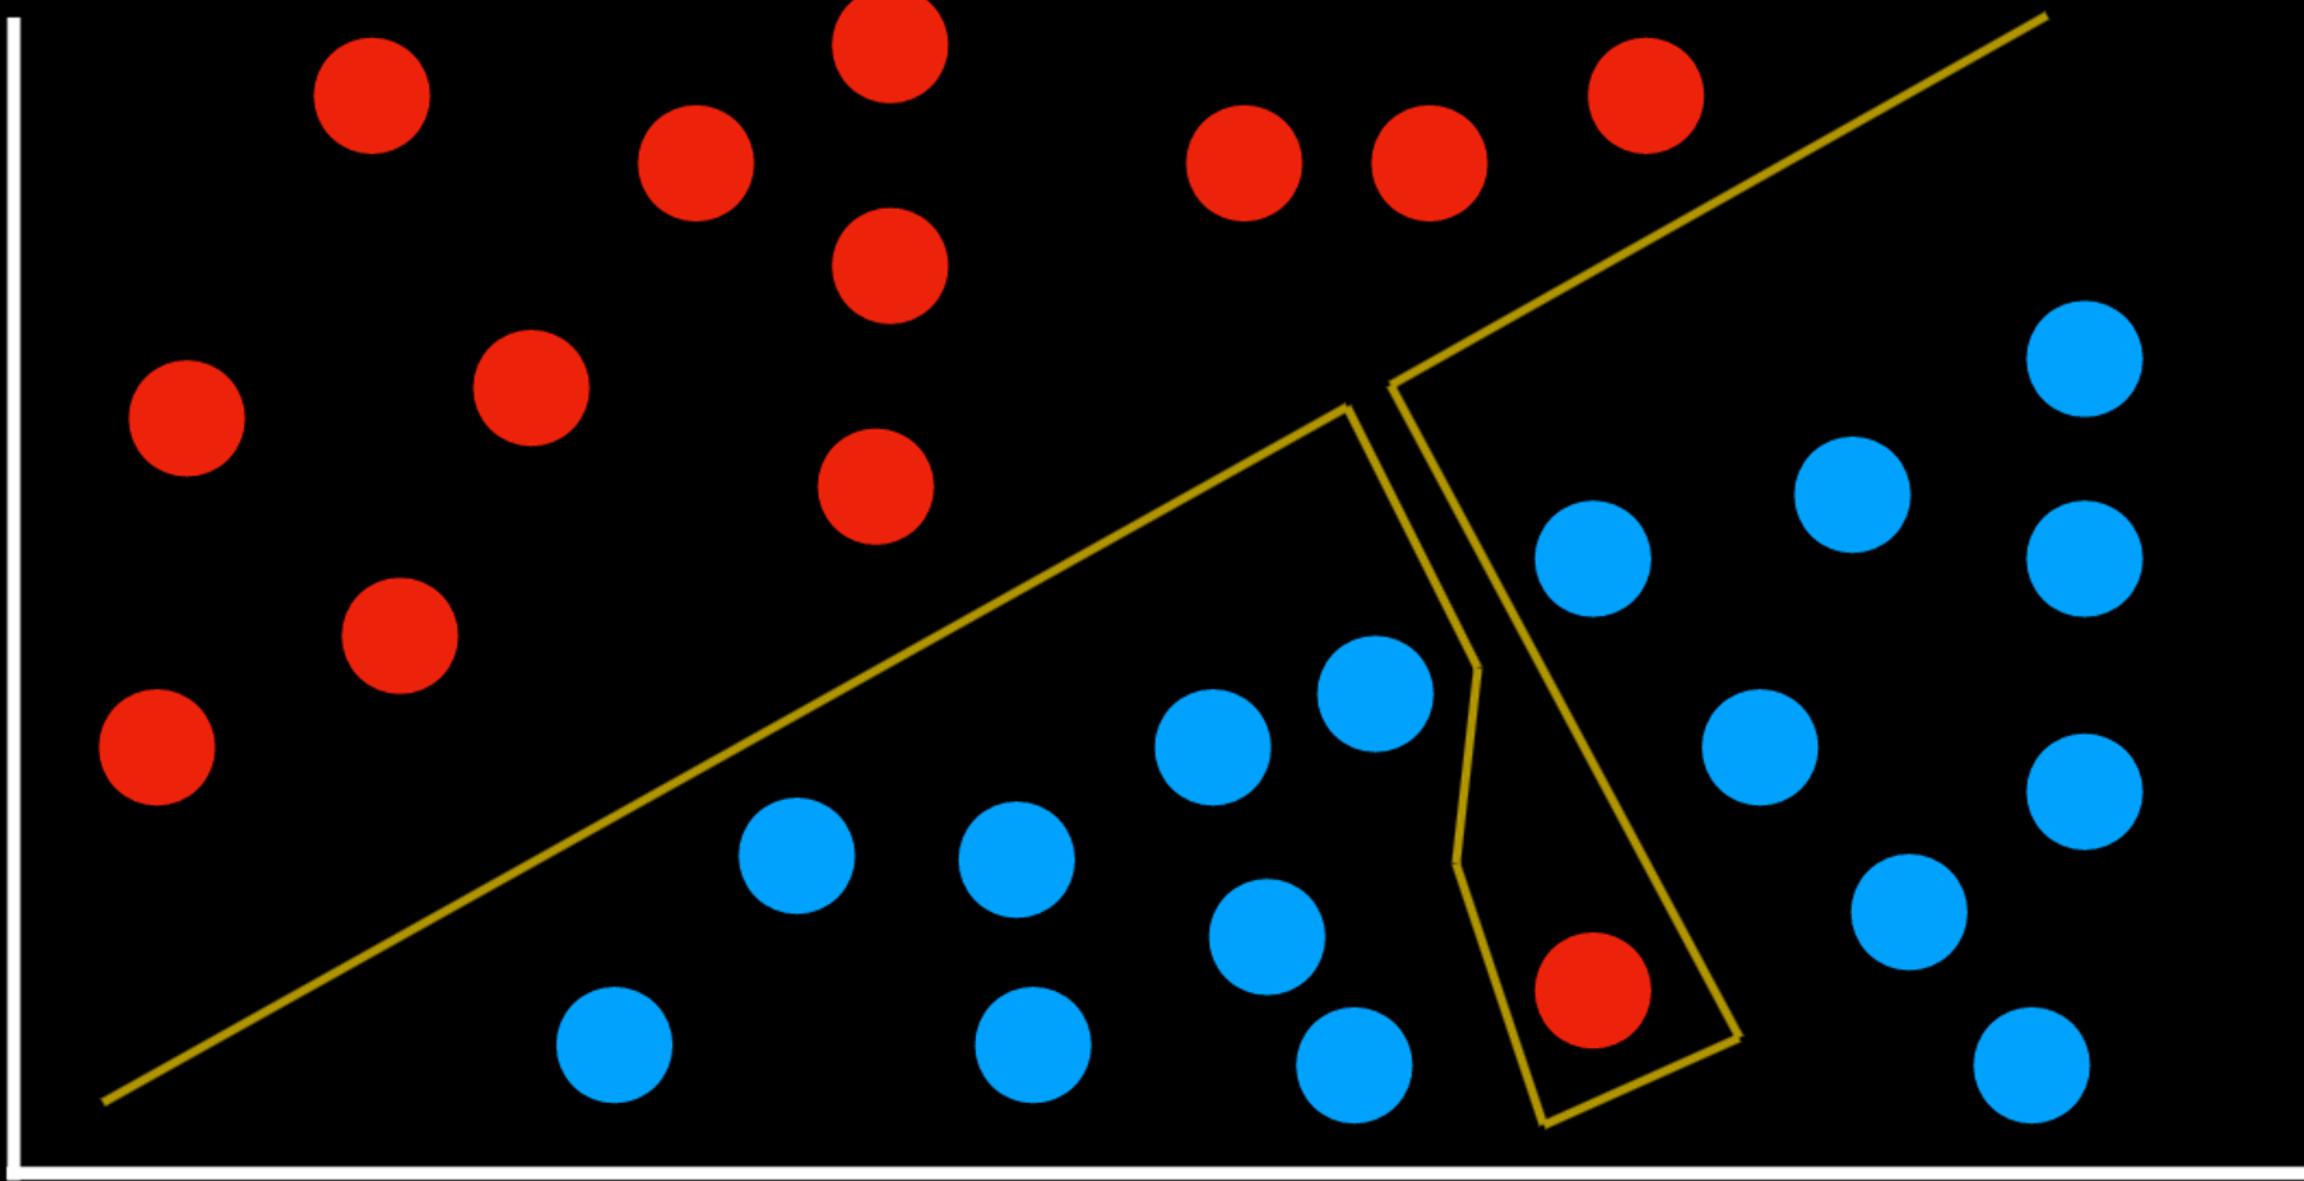
pressure

humidity



pressure

humidity



# regularization

penalizing hypotheses that are more complex  
to favor simpler, more general hypotheses

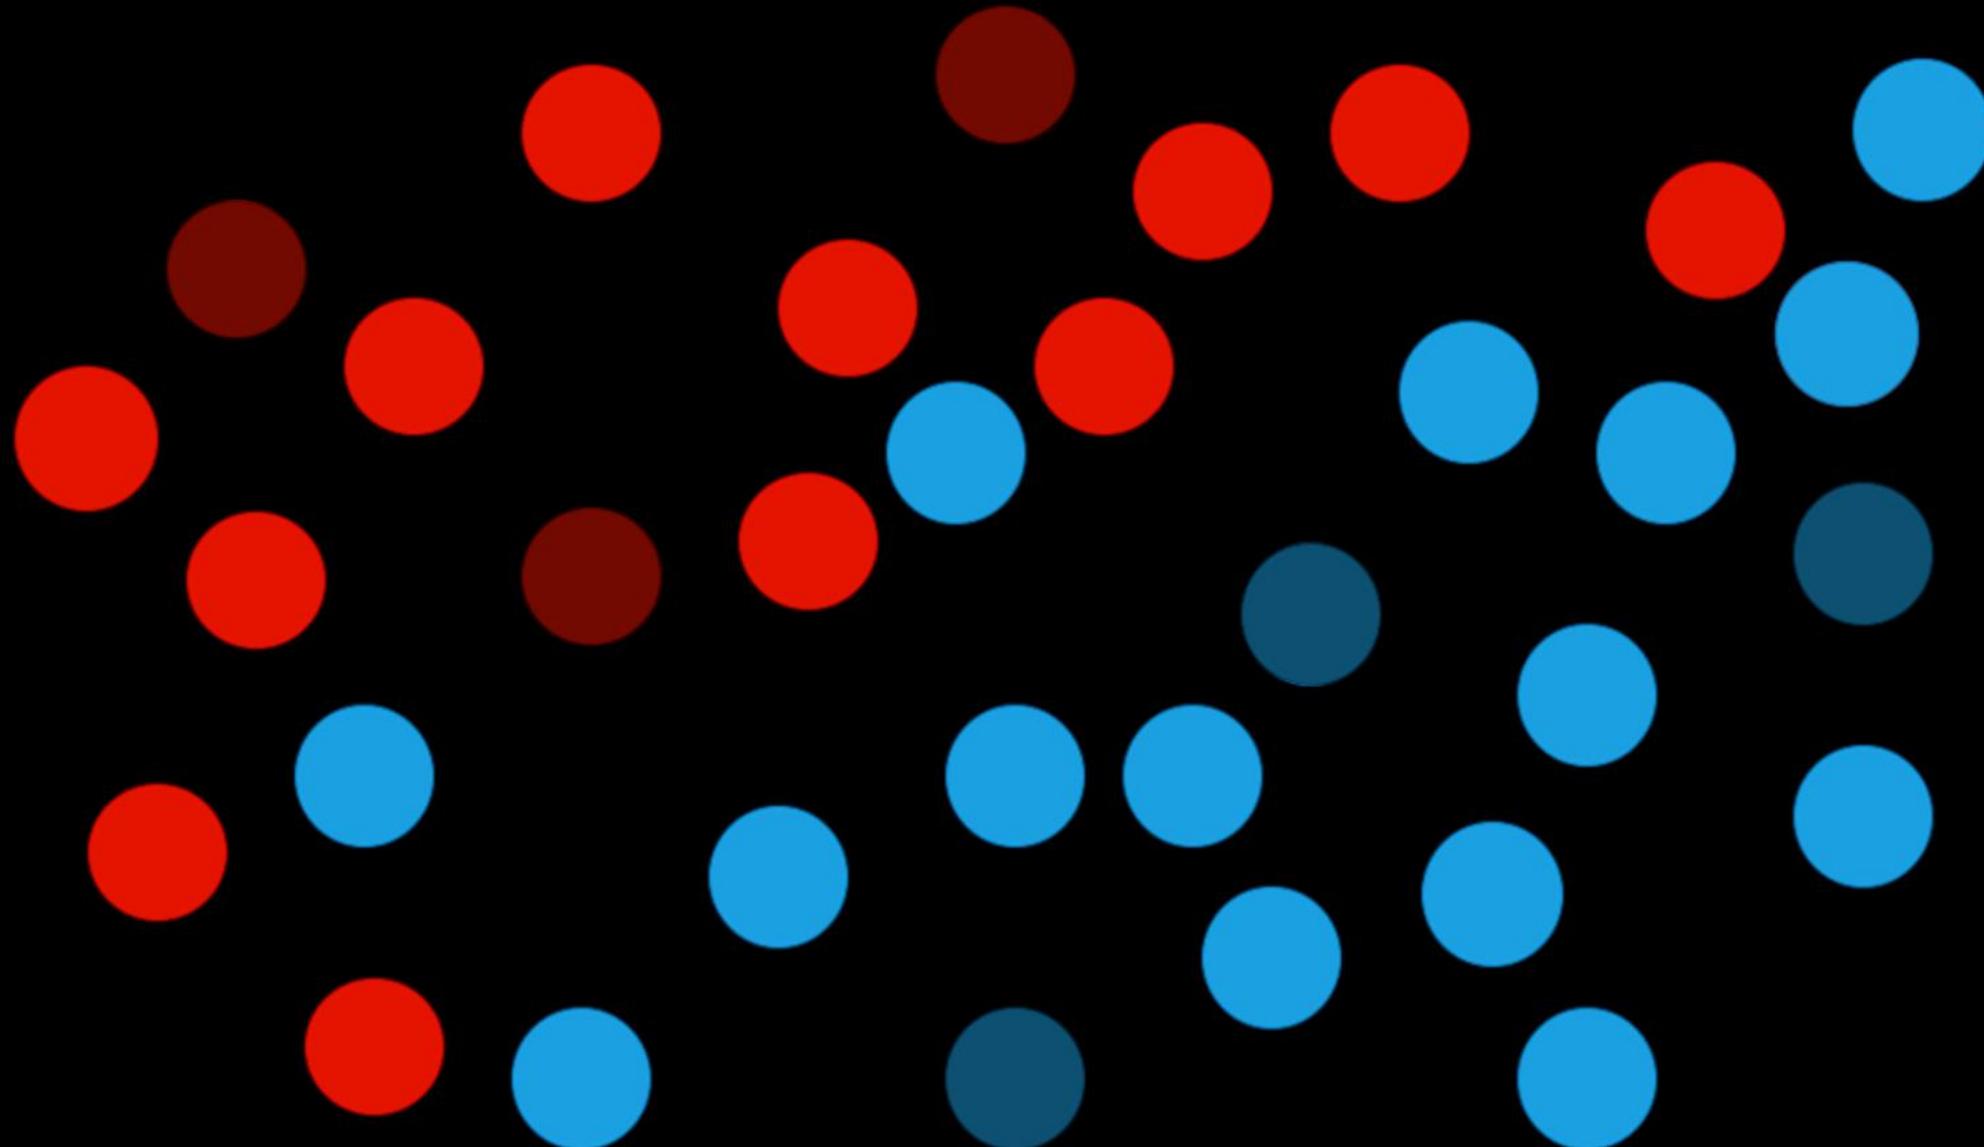
$$cost(h) = loss(h) + \lambda complexity(h)$$

# holdout cross-validation

splitting data into a **training set** and a **test set**, such that learning happens on the training set and is evaluated on the test set

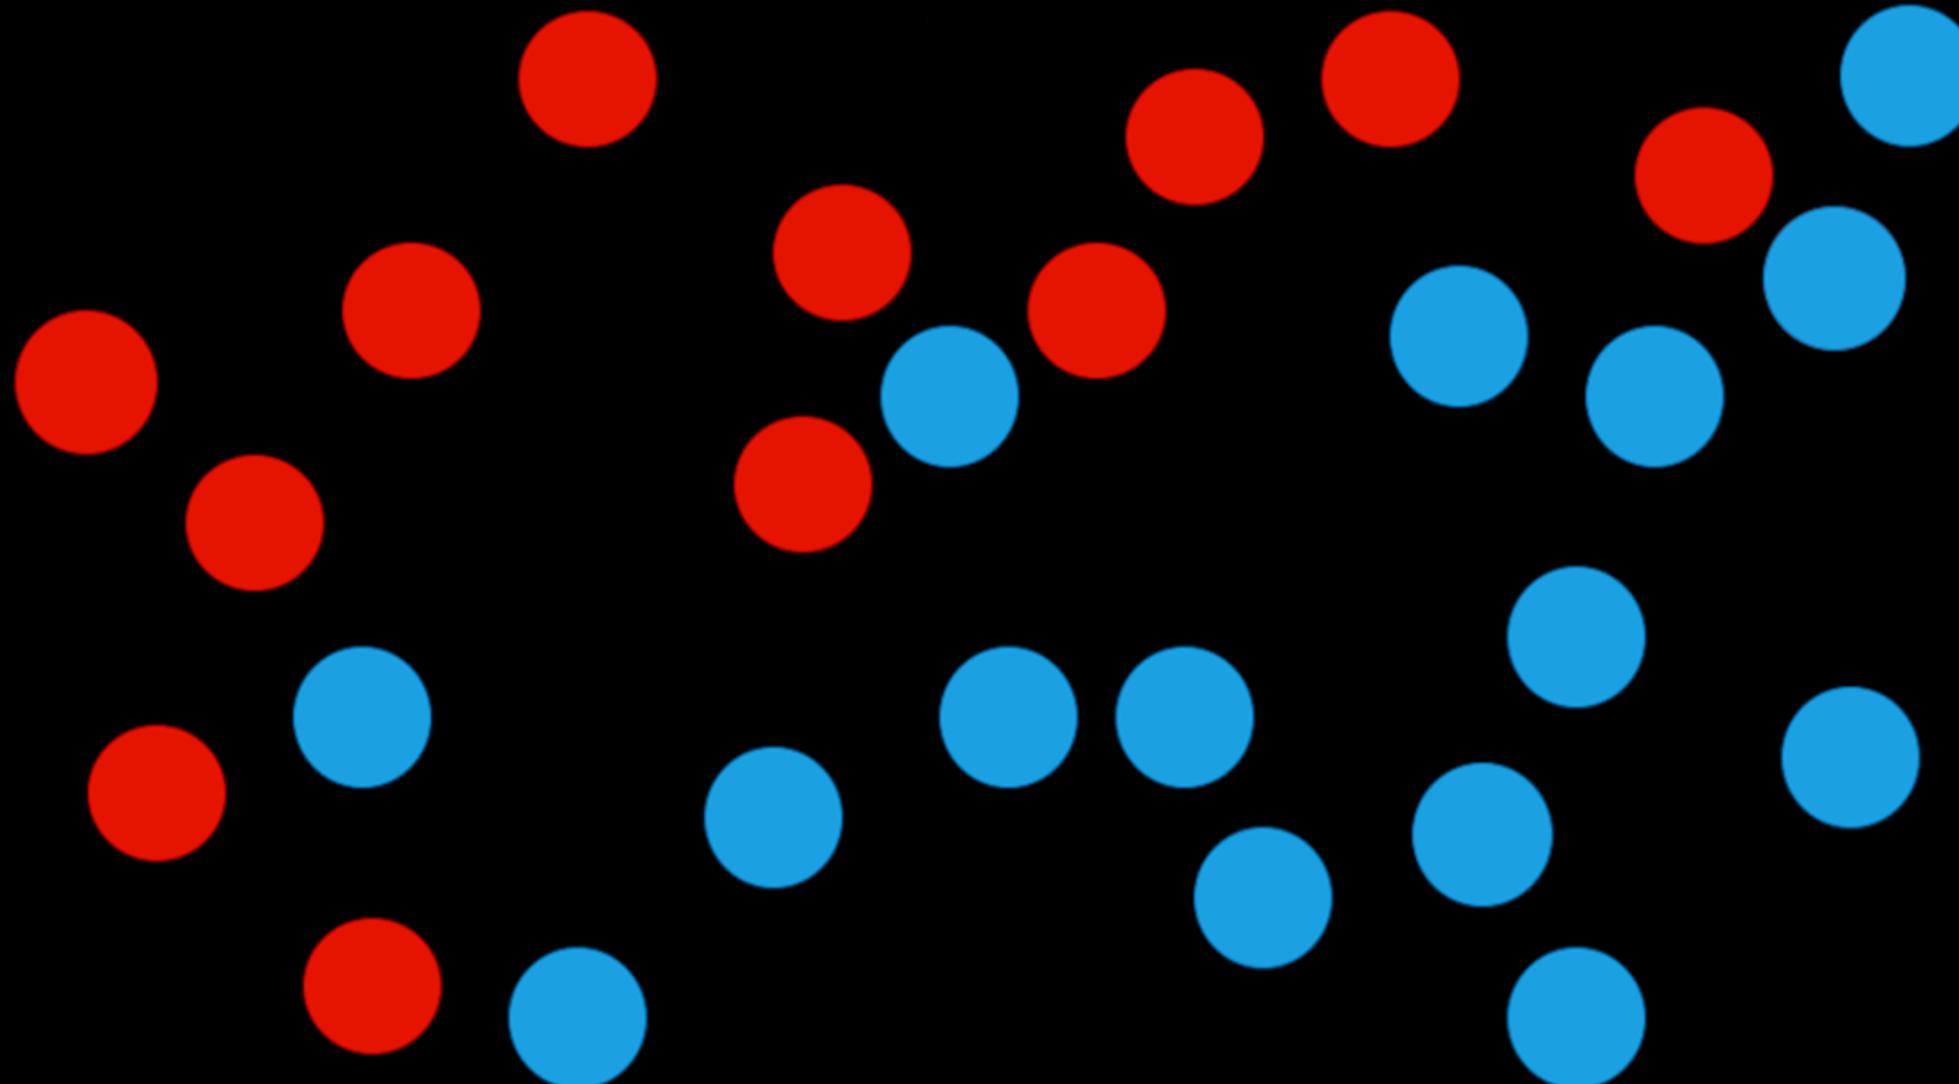
pressure

humidity



pressure

humidity



# Imbalanced Classification

Imbalanced classification refers to classification tasks where the number of examples in each class is unequally distributed.

Typically, imbalanced classification tasks are binary classification tasks where the majority of examples in the training dataset belong to the normal class and a minority of examples belong to the abnormal class.

THE END